

University of
Salford
MANCHESTER

Artificial Intelligent – Week 9

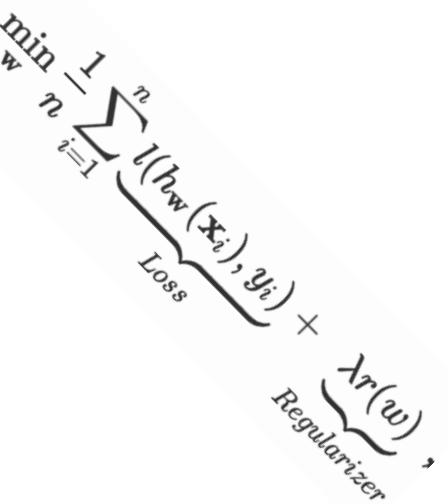
Dr Salem Ameen

S.A.AMEEN1@SALFORD.AC.UK

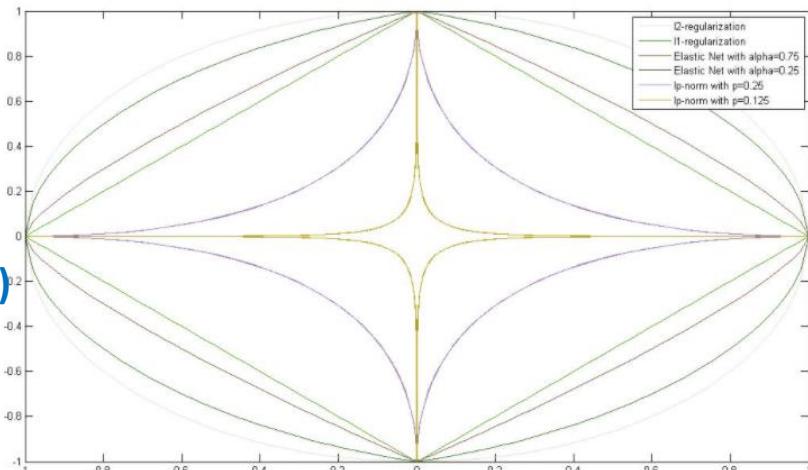
Week 9: Artificial Intelligent – Machine Learning



Loss and Regularizer	Comments
Ordinary Least Squares $\min_{\mathbf{w}} \frac{1}{n} \sum_{i=1}^n (\mathbf{w}^\top \mathbf{x}_i - y_i)^2$	<ul style="list-style-type: none"> Squared Loss No Regularization Closed form solution: $\mathbf{w} = (\mathbf{X}\mathbf{X}^\top)^{-1}\mathbf{X}\mathbf{y}^\top$ $\mathbf{X} = [\mathbf{x}_1, \dots, \mathbf{x}_n]$ $\mathbf{y} = [y_1, \dots, y_n]$
Ridge Regression $\min_{\mathbf{w}} \frac{1}{n} \sum_{i=1}^n (\mathbf{w}^\top \mathbf{x}_i - y_i)^2 + \lambda \ \mathbf{w}\ _2^2$	<ul style="list-style-type: none"> Squared Loss l_2-Regularization $\mathbf{w} = (\mathbf{X}\mathbf{X}^\top + \lambda \mathbb{I})^{-1}\mathbf{X}\mathbf{y}^\top$
Lasso $\min_{\mathbf{w}} \frac{1}{n} \sum_{i=1}^n (\mathbf{w}^\top \mathbf{x}_i - y_i)^2 + \lambda \ \mathbf{w}\ _1$	<ul style="list-style-type: none"> + sparsity inducing (good for feature selection) + Convex - Not strictly convex (no unique solution) - Not differentiable (at 0) Solve with (sub)-gradient descent or SVEN
Elastic Net $\min_{\mathbf{w}} \frac{1}{n} \sum_{i=1}^n (\mathbf{w}^\top \mathbf{x}_i - y_i)^2 + \alpha \ \mathbf{w}\ _1 + (1 - \alpha) \ \mathbf{w}\ _2^2$ $\alpha \in [0, 1]$	<ul style="list-style-type: none"> ADVANTAGE: Strictly convex (i.e. unique solution) + sparsity inducing (good for feature selection) + Dual of squared-loss SVM, see SVEN DISADVANTAGE: - Non-differentiable



Loss and Regularizer	Comments
Logistic Regression $\min_{\mathbf{w}, b} \frac{1}{n} \sum_{i=1}^n \log(1 + e^{-y_i(\mathbf{w}^\top \mathbf{x}_i + b)})$	<ul style="list-style-type: none"> Often l_1 or l_2 Regularized Solve with gradient descent. $\Pr(y x) = \frac{1}{1+e^{-(\mathbf{w}^\top \mathbf{x} + b)}}$
Linear Support Vector Machine $\min_{\mathbf{w}, b} C \sum_{i=1}^n \max[1 - y_i(\mathbf{w}^\top \mathbf{x}_i + b), 0] + \ \mathbf{w}\ _2^2$	<ul style="list-style-type: none"> Typically l_2 regularized (sometimes l_1). Quadratic program. When kernelized leads to sparse solutions. Kernelized version can be solved very efficiently with specialized algorithms (e.g. SMO)

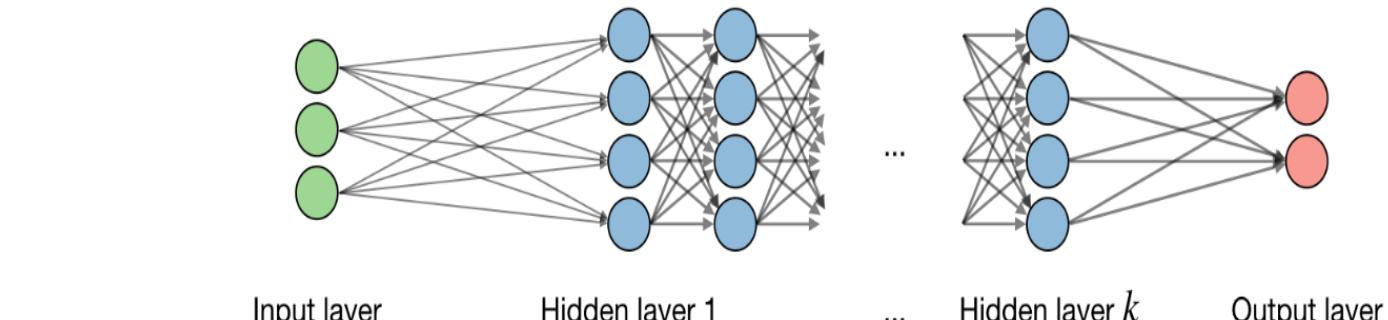


[tf.keras.regularizers.L1\(0.3\)](#)
[tf.keras.regularizers.L2\(0.1\)](#)
[tf.keras.regularizers.L1L2\(l1=0.01, l2=0.01\)](#)

Week 9: Artificial Intelligent – Deep Learning



Features									output
x1	x2	x3	x4	x5	x6	x7	x8	y	



By noting i the i^{th} layer of the network and j the j^{th} hidden unit of the layer, we have:

$$z_j^{[i]} = w_j^{[i]T} x + b_j^{[i]}$$

Week 9: Artificial Intelligent – Deep Learning



- **Architecture of a CNN (Computer Vision):**
- **Architecture of a RNN (Natural Language Processing):**

Week 9: Artificial Intelligent – Deep Learning – CNN - Vision

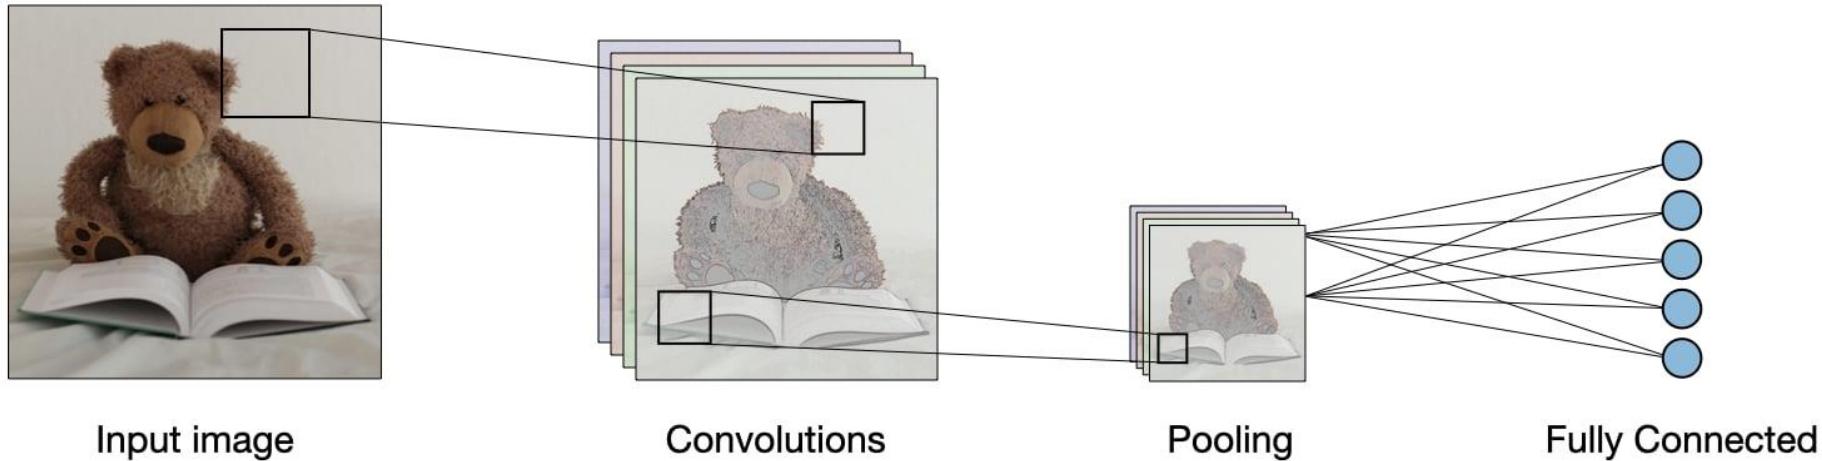


- **Architecture of a traditional CNN:** Convolutional neural networks, also known as CNNs, are a specific type of neural networks that are generally composed of the following layers:

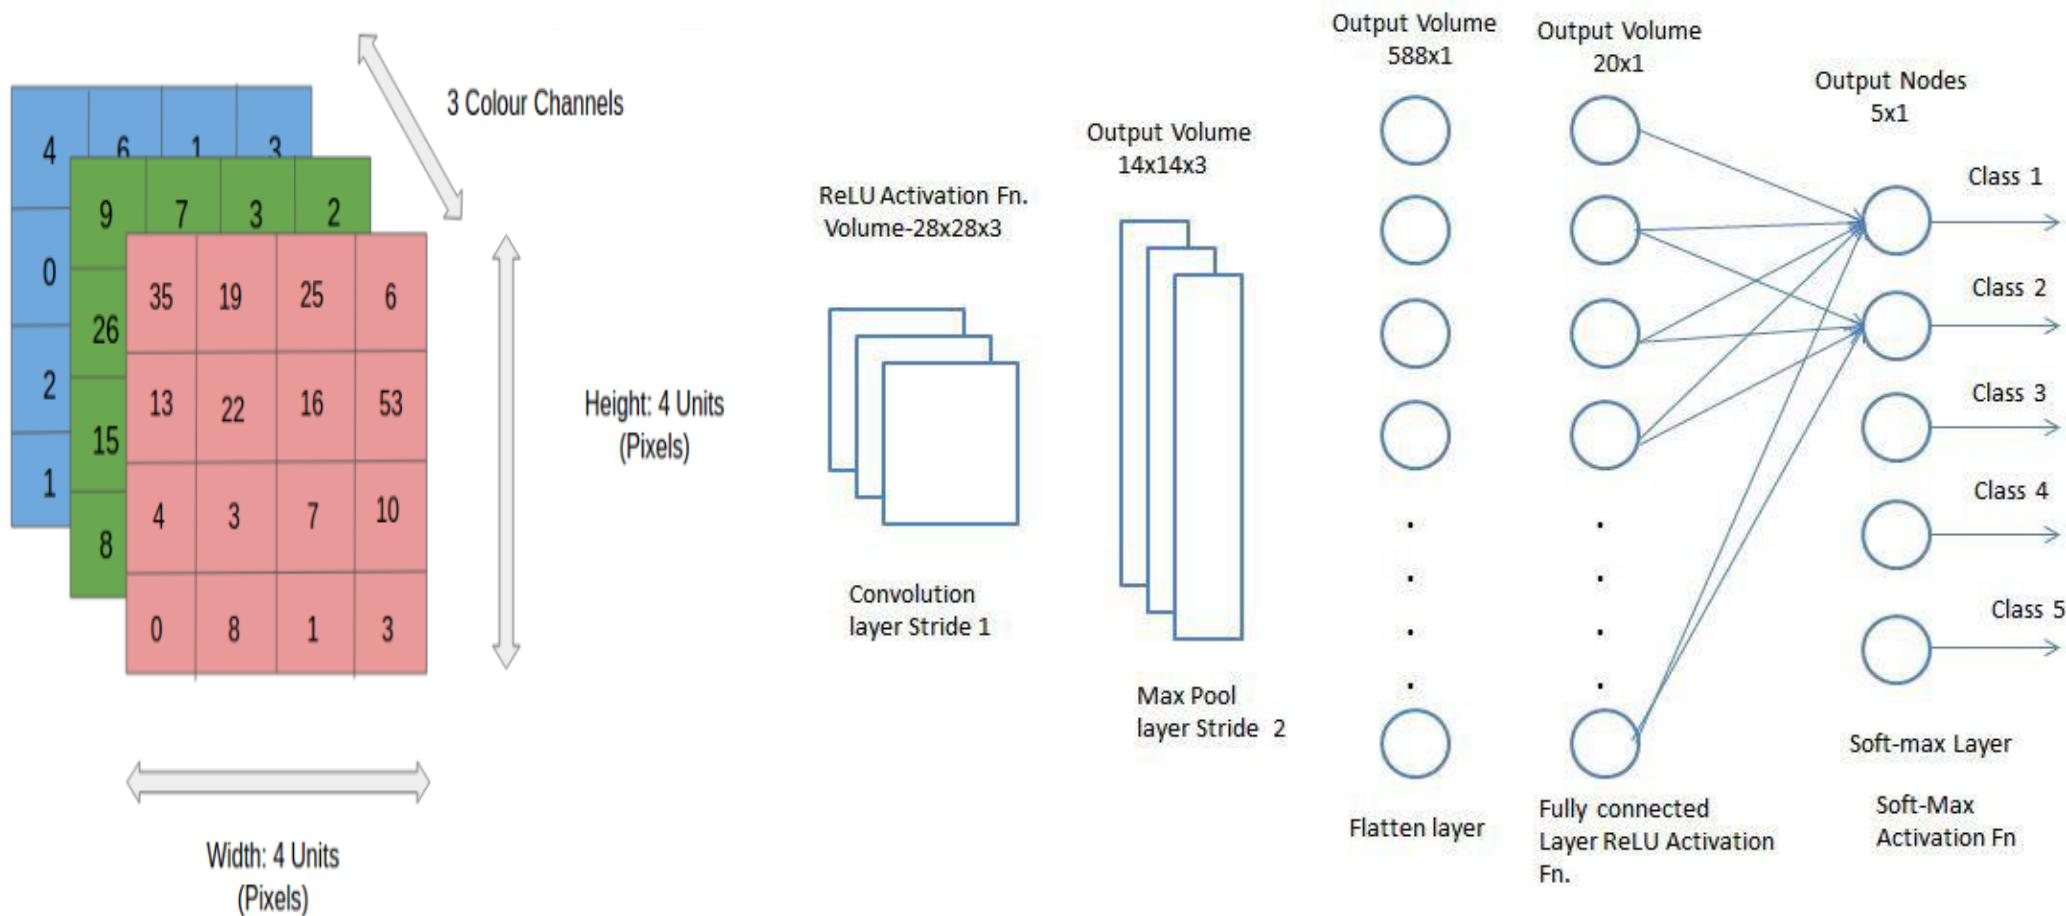
Week 9: Artificial Intelligent – Deep Learning – CNN - Vision



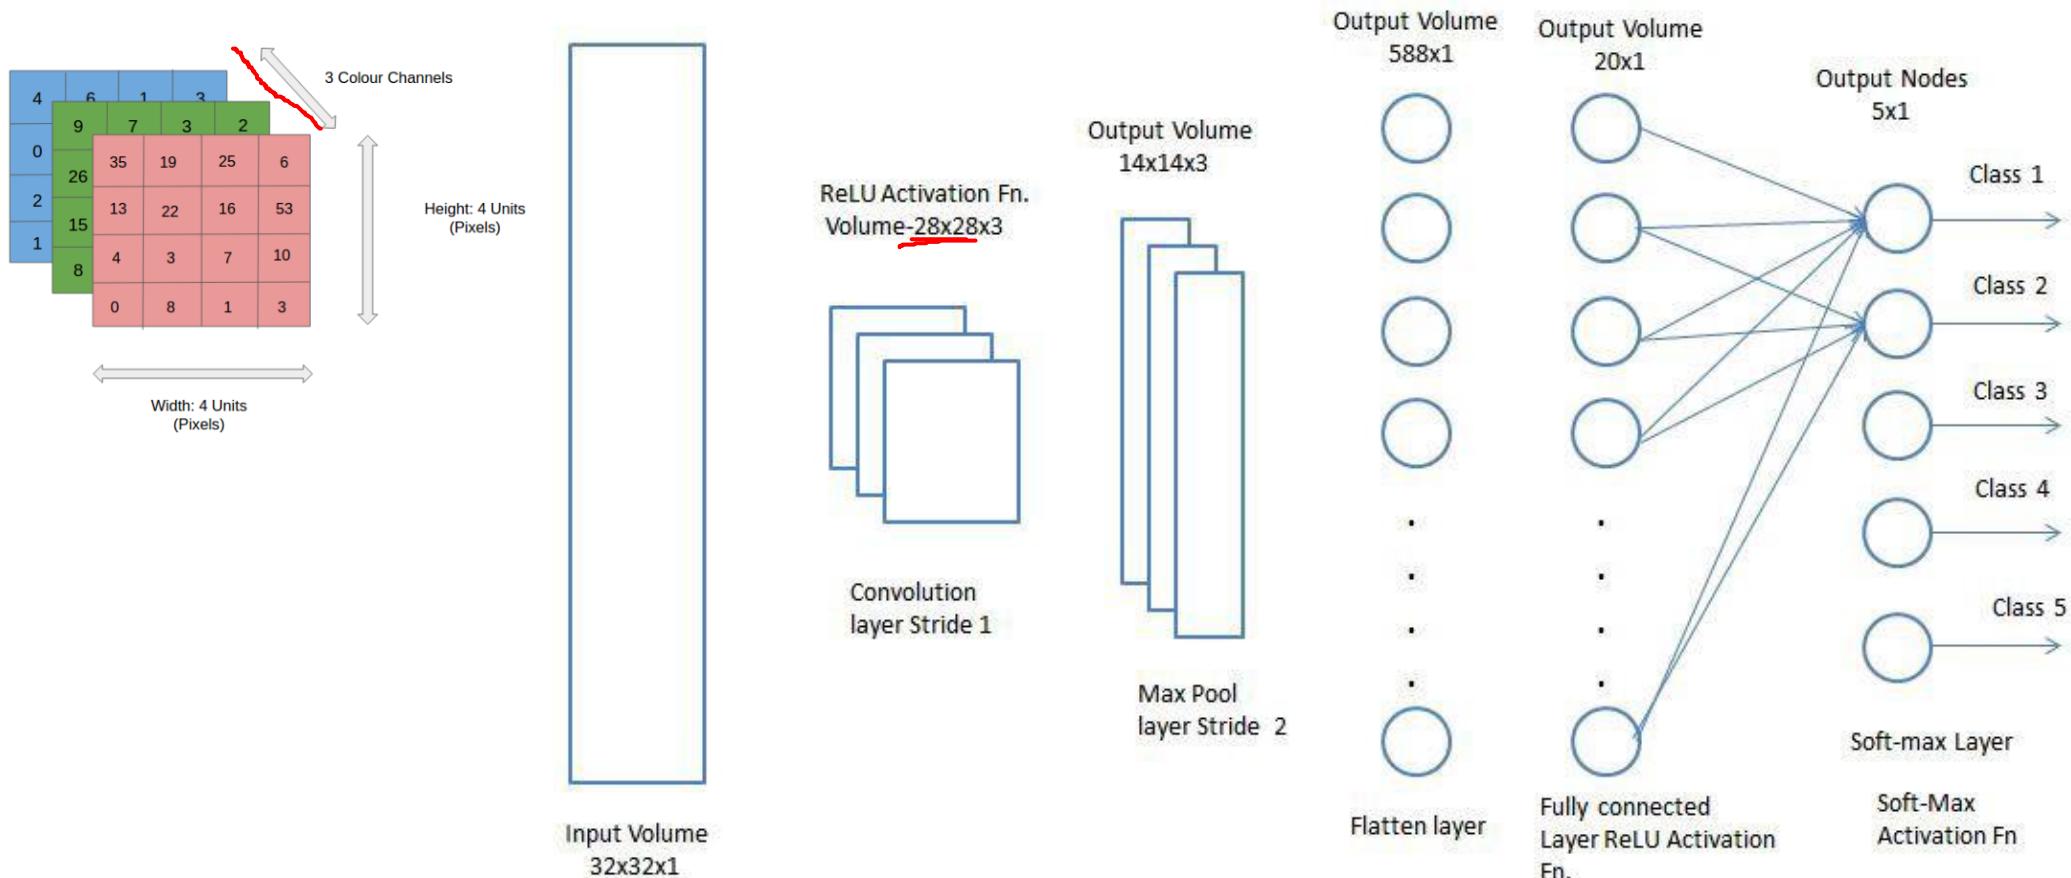
- **Architecture of a traditional CNN:** Convolutional neural networks, also known as CNNs, are a specific type of neural networks that are generally composed of the following layers:



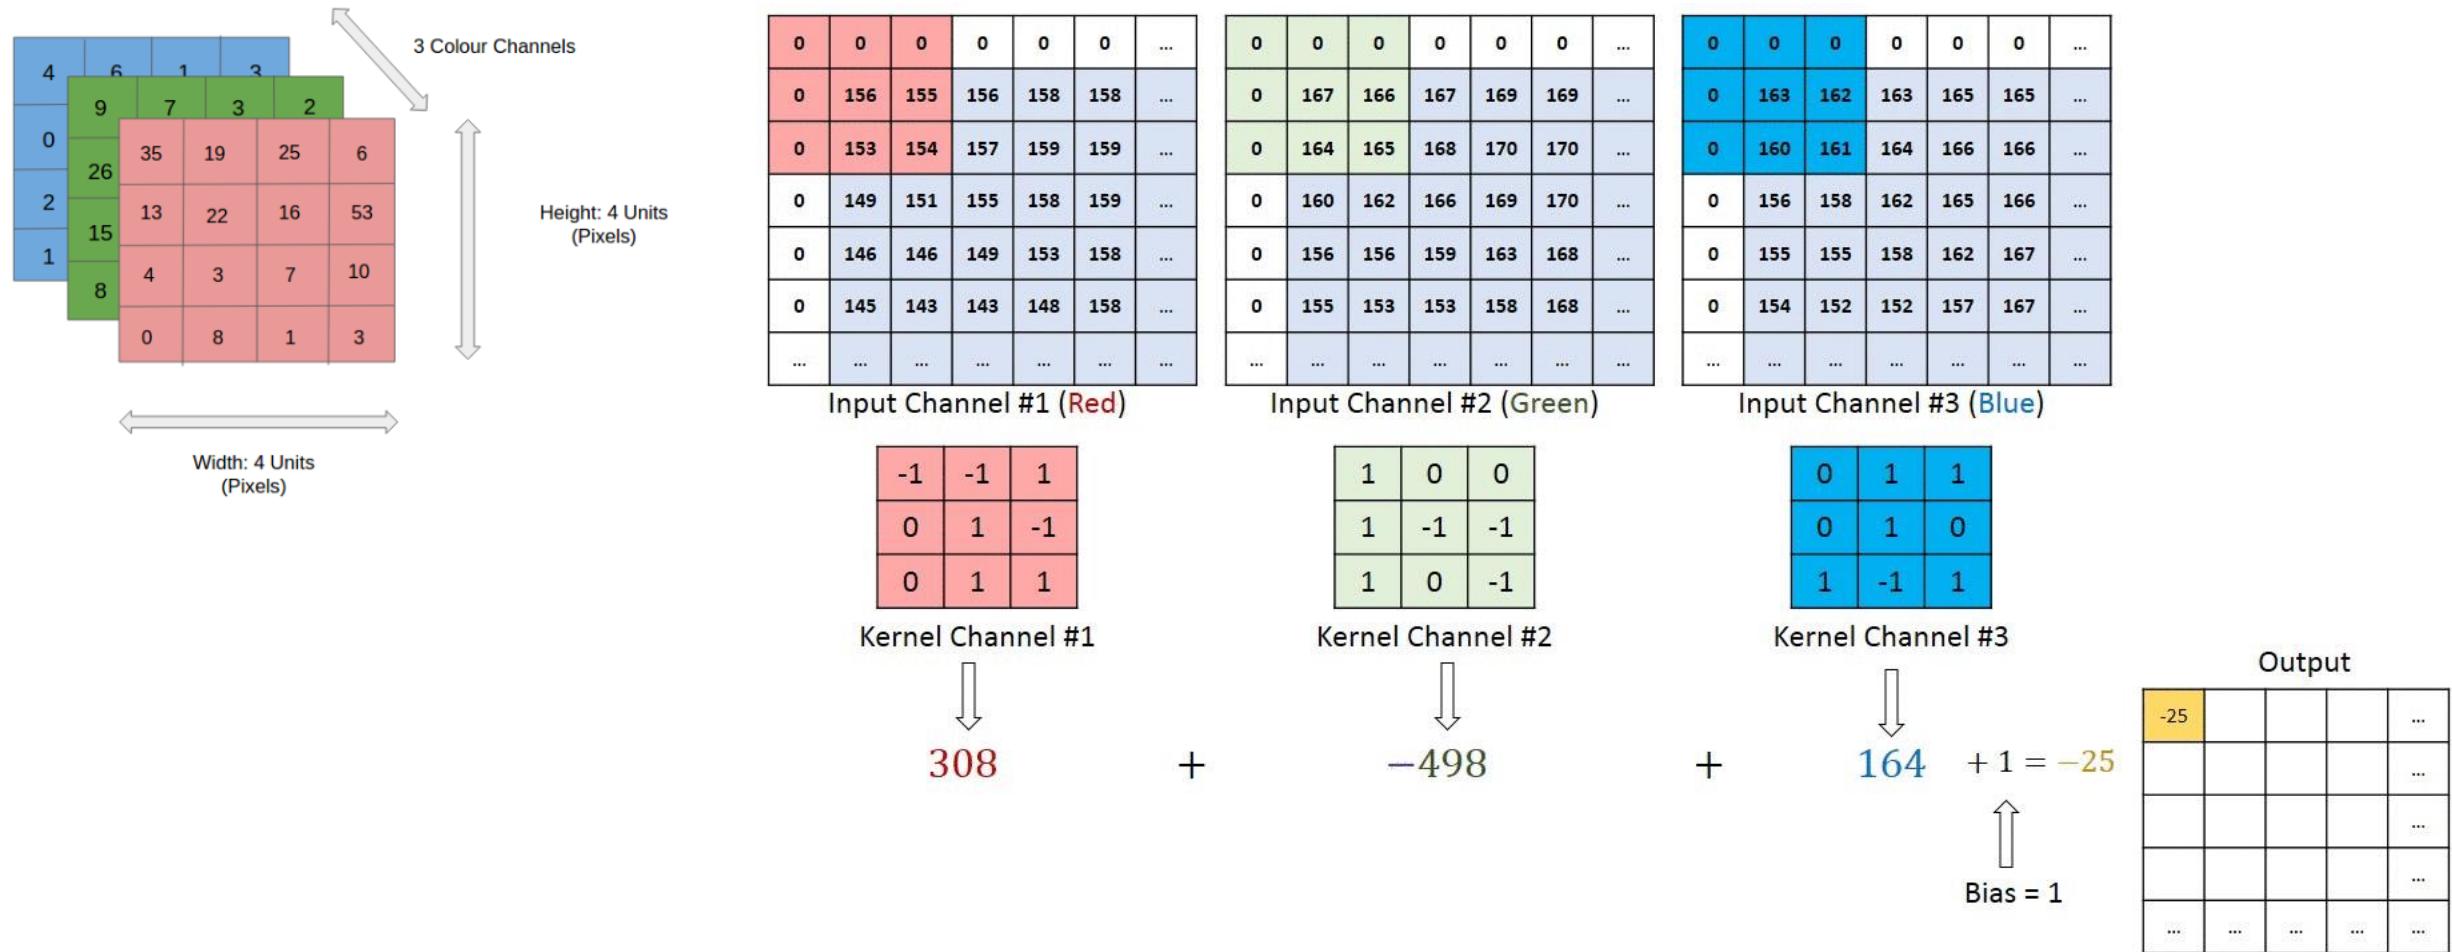
Week 9: Artificial Intelligent – Deep Learning – CNN - Vision



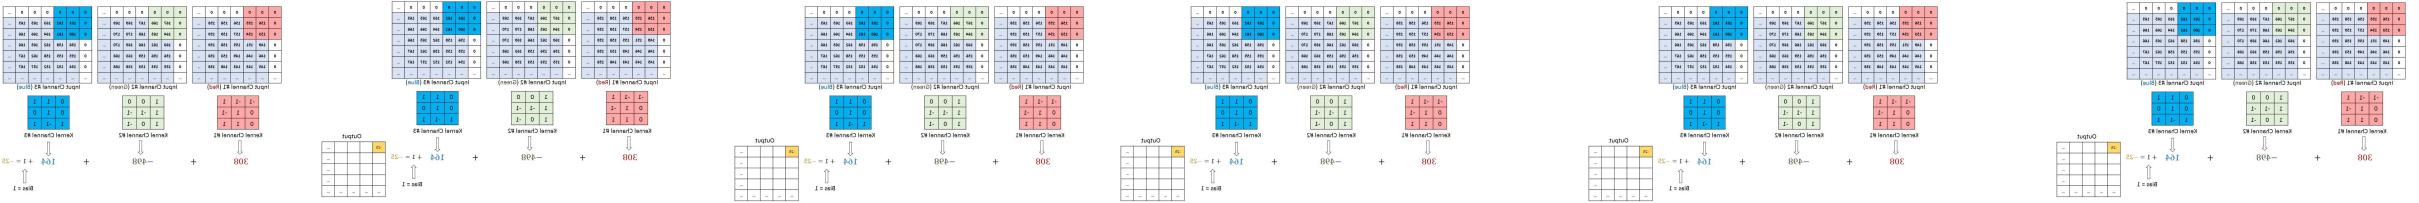
Week 9: Artificial Intelligent – Deep Learning – CNN - Vision



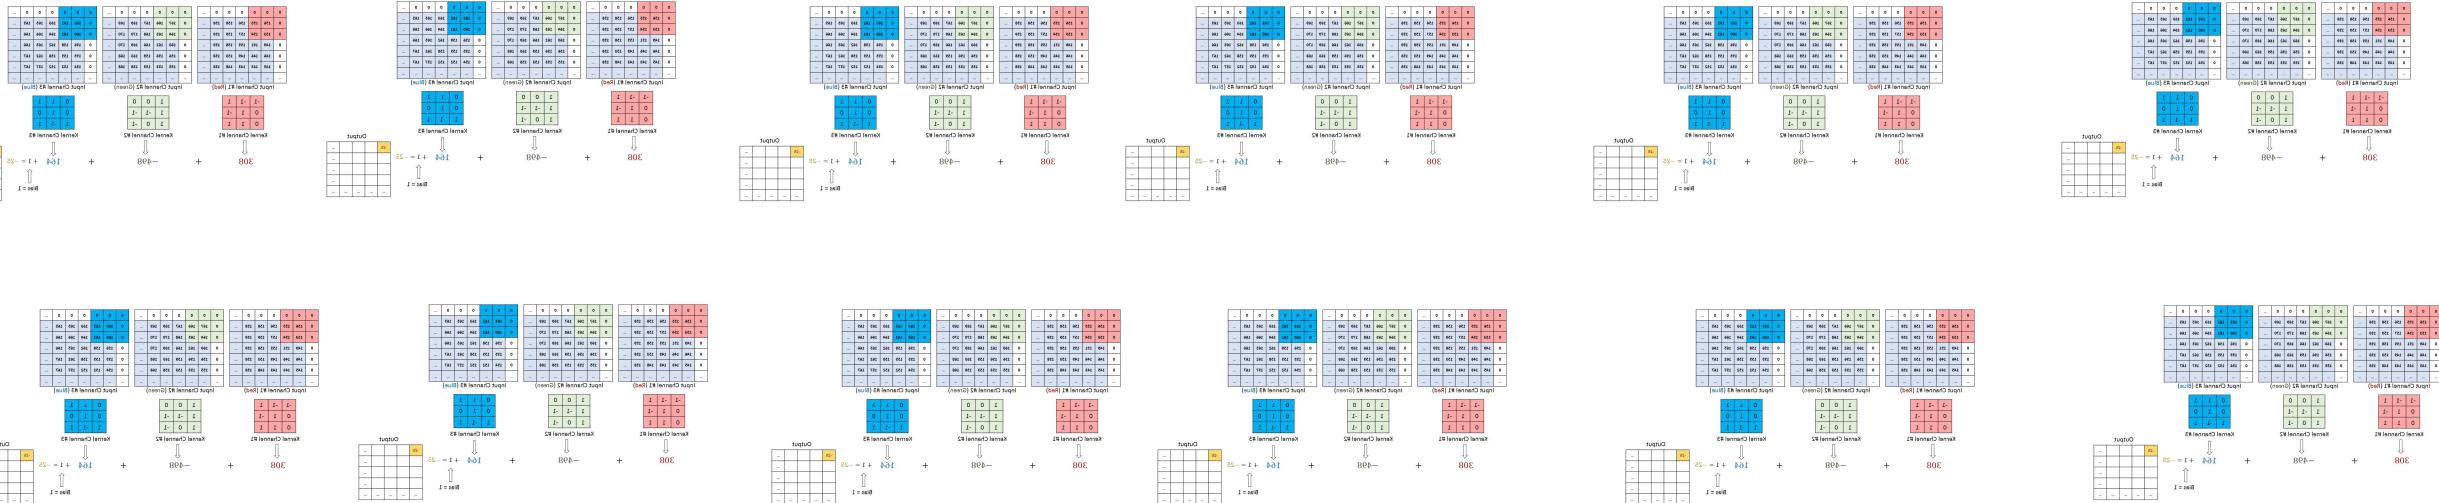
Week 9: Artificial Intelligent – CNN – Convolution layer



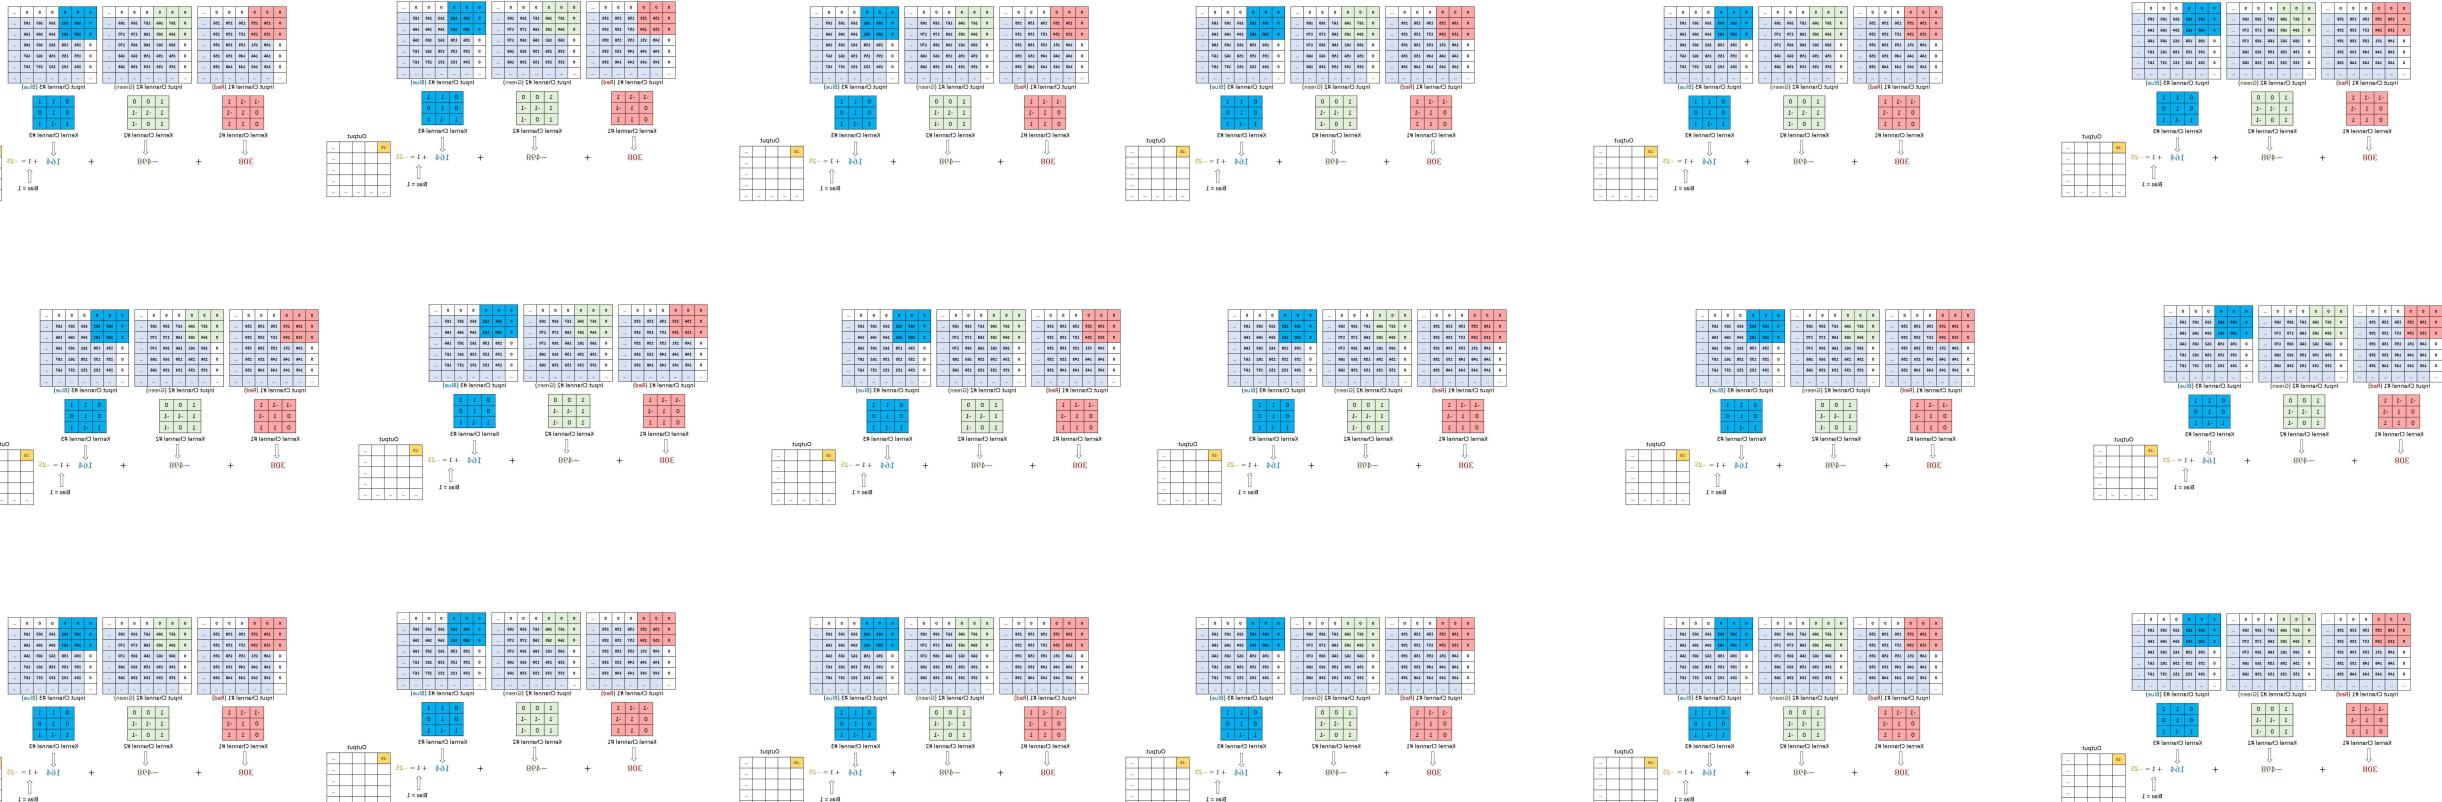
Week 9: Artificial Intelligent – CNN – Convolution layer



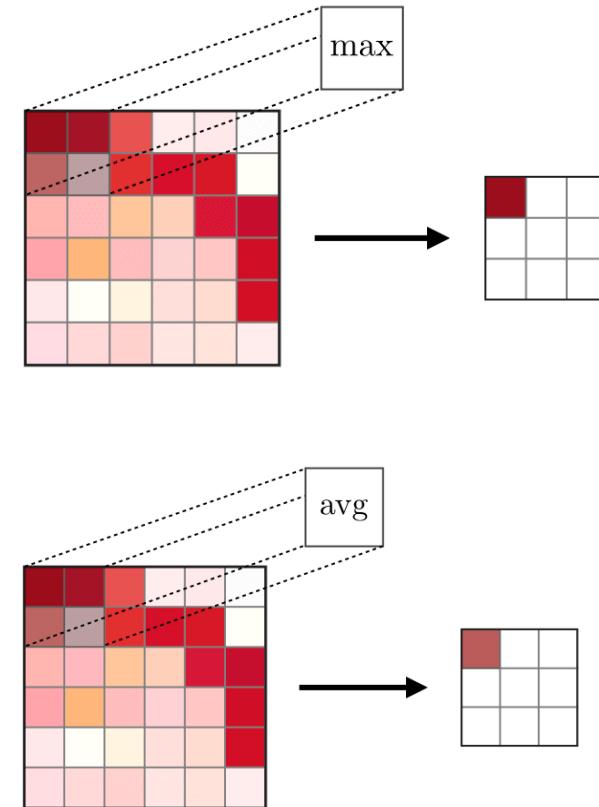
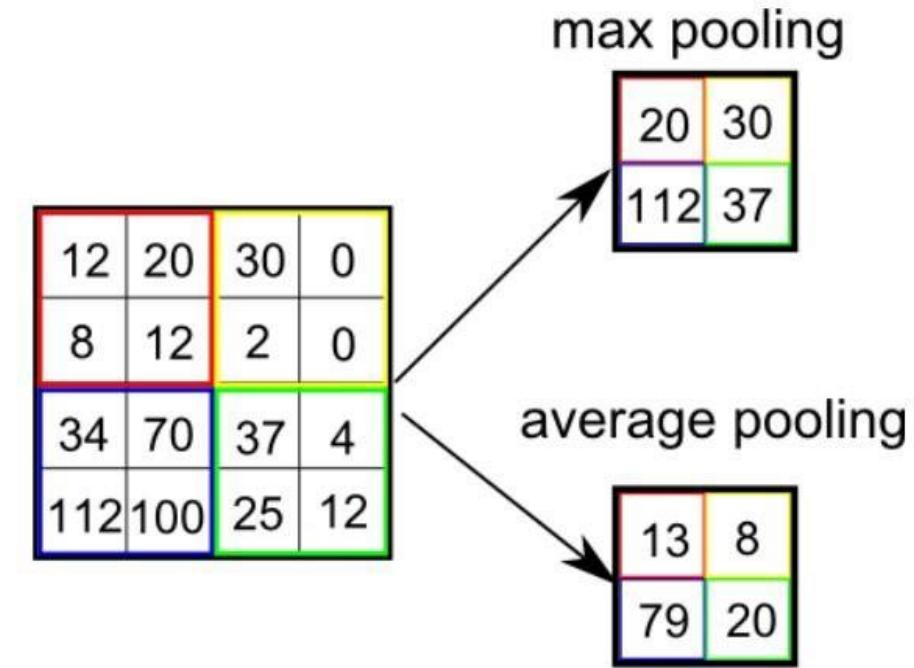
Week 9: Artificial Intelligent – CNN – Convolution layer



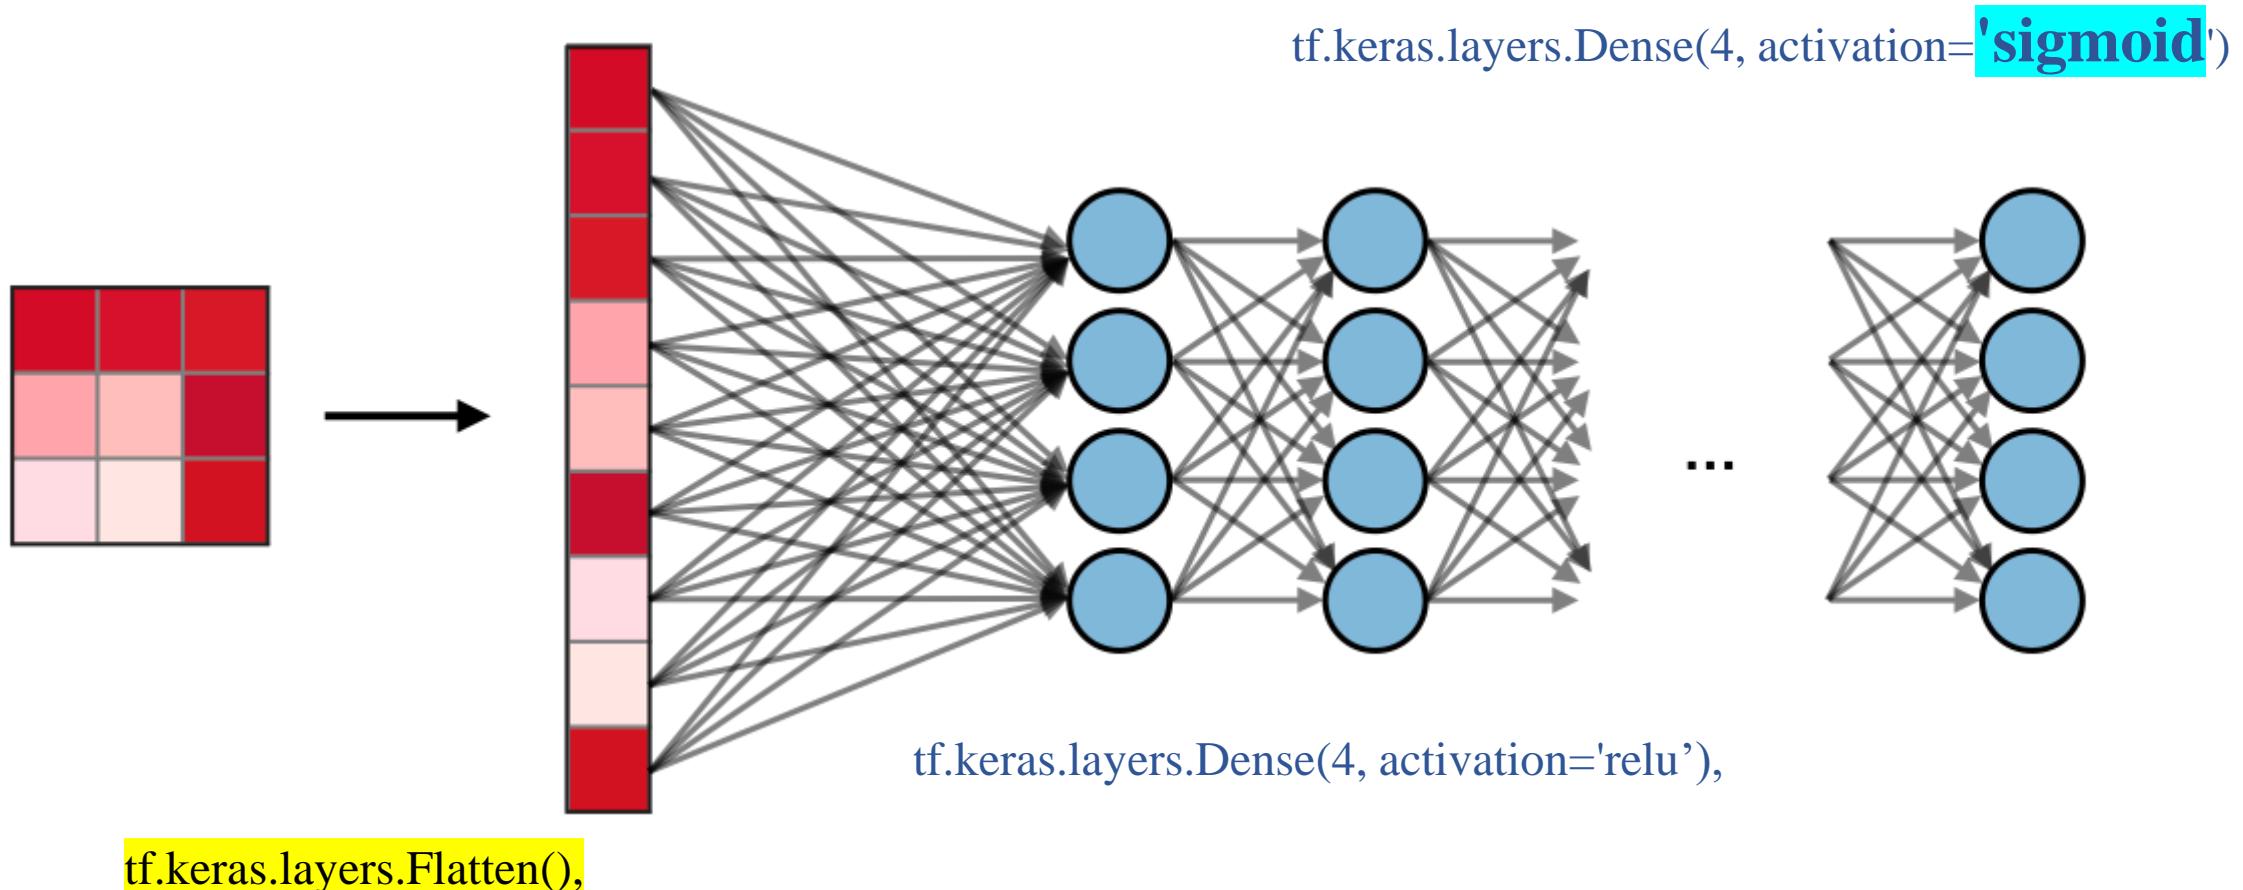
Week 9: Artificial Intelligent – CNN – Convolution layer



Week 9: Artificial Intelligent – CNN – Pooling Layer



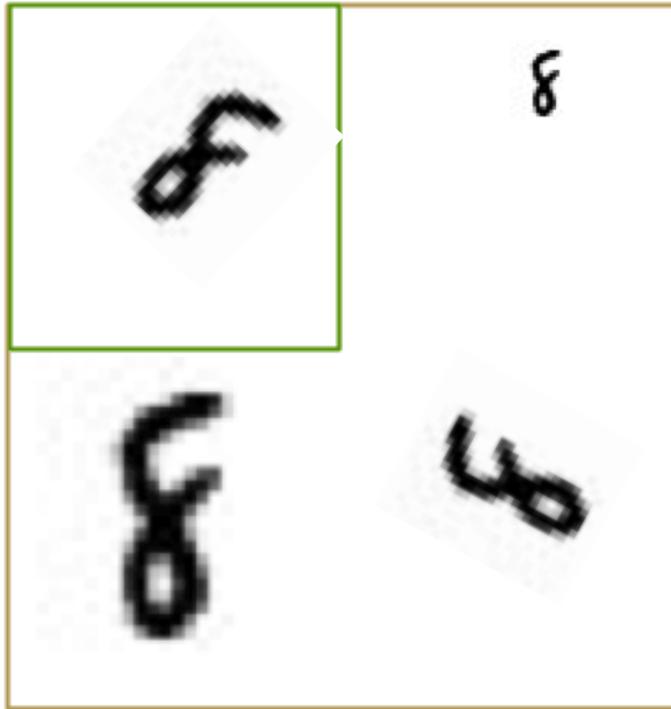
Week 9: Artificial Intelligent – CNN – Fully Connected Layer



Week 9: Artificial Intelligent – CNN – Data processing



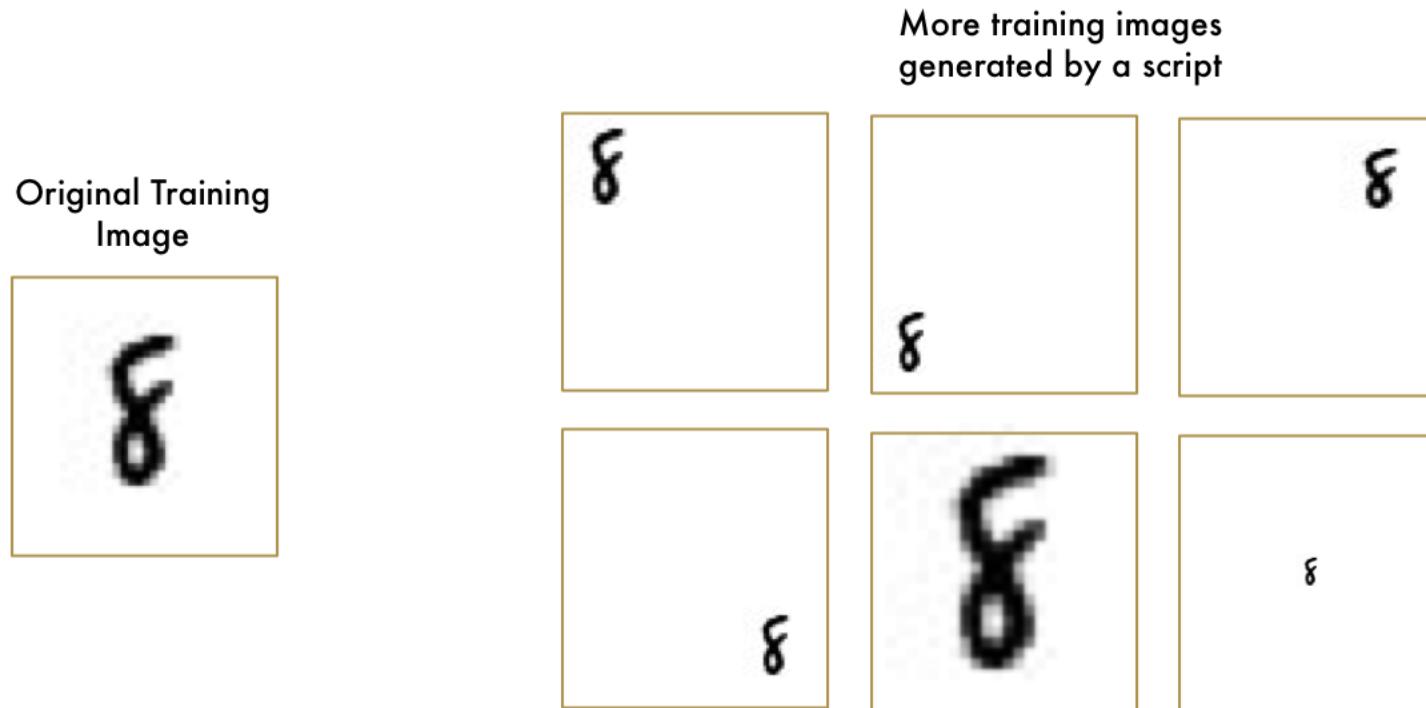
Test Image



Prediction from
our network

No idea?!

Week 9: Artificial Intelligent – CNN – Data processing



Week 9: Artificial Intelligent – CNN – Data processing

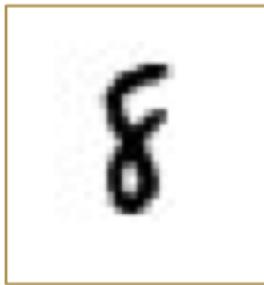


```
from tensorflow.keras.preprocessing.image import  
ImageDataGenerator
```

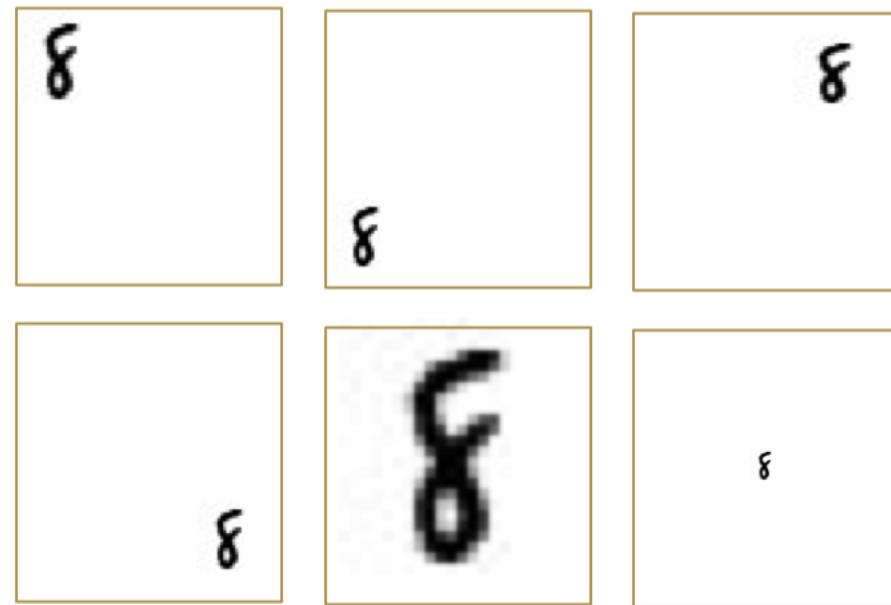
```
# All images will be rescaled by 1./255  
train_datagen = ImageDataGenerator(  
    rescale=1./255,  
    rotation_range=40,  
    width_shift_range=0.2,  
    height_shift_range=0.2,  
    shear_range=0.2,  
    zoom_range=0.2,  
    horizontal_flip=True,  
    fill_mode='nearest')
```

```
validation_datagen = ImageDataGenerator(rescale=1/255)
```

Original Training
Image



More training images
generated by a script



Week 9: Artificial Intelligent – CNN – Data processing



- **Data augmentation:** Deep learning models usually need a lot of data to be properly trained. It is often useful to get more data from the existing ones using data augmentation techniques.

Session Code:
2675
9432

Week 9: Artificial Intelligent – CNN – Data processing



- **Data augmentation:** Deep learning models usually need a lot of data to be properly trained. It is often useful to get more data from the existing ones using data augmentation techniques.
- The main ones are summed up in the table below. More precisely, given the following input image,



• Image without any modification

Week 9: Artificial Intelligent – CNN – Data processing



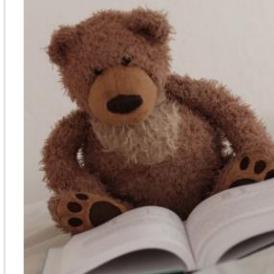
- **Data augmentation:** Deep learning models usually need a lot of data to be properly trained. It is often useful to get more data from the existing ones using data augmentation techniques.
- The main ones are summed up in the table below. More precisely, given the following input image, here are the techniques that we can apply:



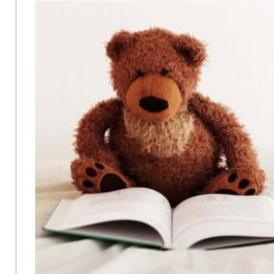
• Image without any modification

Week 9: Artificial Intelligent – CNN – Data processing



Original	Flip	Rotation	Random crop
			

• Image without any modification	• Flipped with respect to an axis for which the meaning of the image is preserved	• Rotation with a slight angle • Simulates incorrect horizon calibration	• Random focus on one part of the image • Several random crops can be done in a row
----------------------------------	---	---	--

Color shift	Noise addition	Information loss	Contrast change
			

• Nuances of RGB is slightly changed • Captures noise that can occur with light exposure	• Addition of noise • More tolerance to quality variation of inputs	• Parts of image ignored • Mimics potential loss of parts of image	• Luminosity changes • Controls difference in exposition due to time of day
---	--	---	--

Week 9: Artificial Intelligent – CNN – Modelling



```
model = tf.keras.models.Sequential([
    tf.keras.layers.Conv2D(16, (3,3), activation='relu', input_shape=(300, 300, 3)),
    tf.keras.layers.MaxPooling2D(2, 2),
    tf.keras.layers.Conv2D(32, (3,3), activation='relu'),
    tf.keras.layers.MaxPooling2D(2,2),
    tf.keras.layers.Conv2D(64, (3,3), activation='relu'),
    tf.keras.layers.MaxPooling2D(2,2),
    tf.keras.layers.Conv2D(64, (3,3), activation='relu'),
    tf.keras.layers.MaxPooling2D(2,2),
    tf.keras.layers.Conv2D(64, (3,3), activation='relu'),
    tf.keras.layers.MaxPooling2D(2,2),
    tf.keras.layers.Flatten(),
    tf.keras.layers.Dense(512, activation='relu'),
    tf.keras.layers.Dense(1, activation='sigmoid')
])
```

Week 9: Artificial Intelligent – CNN – Training

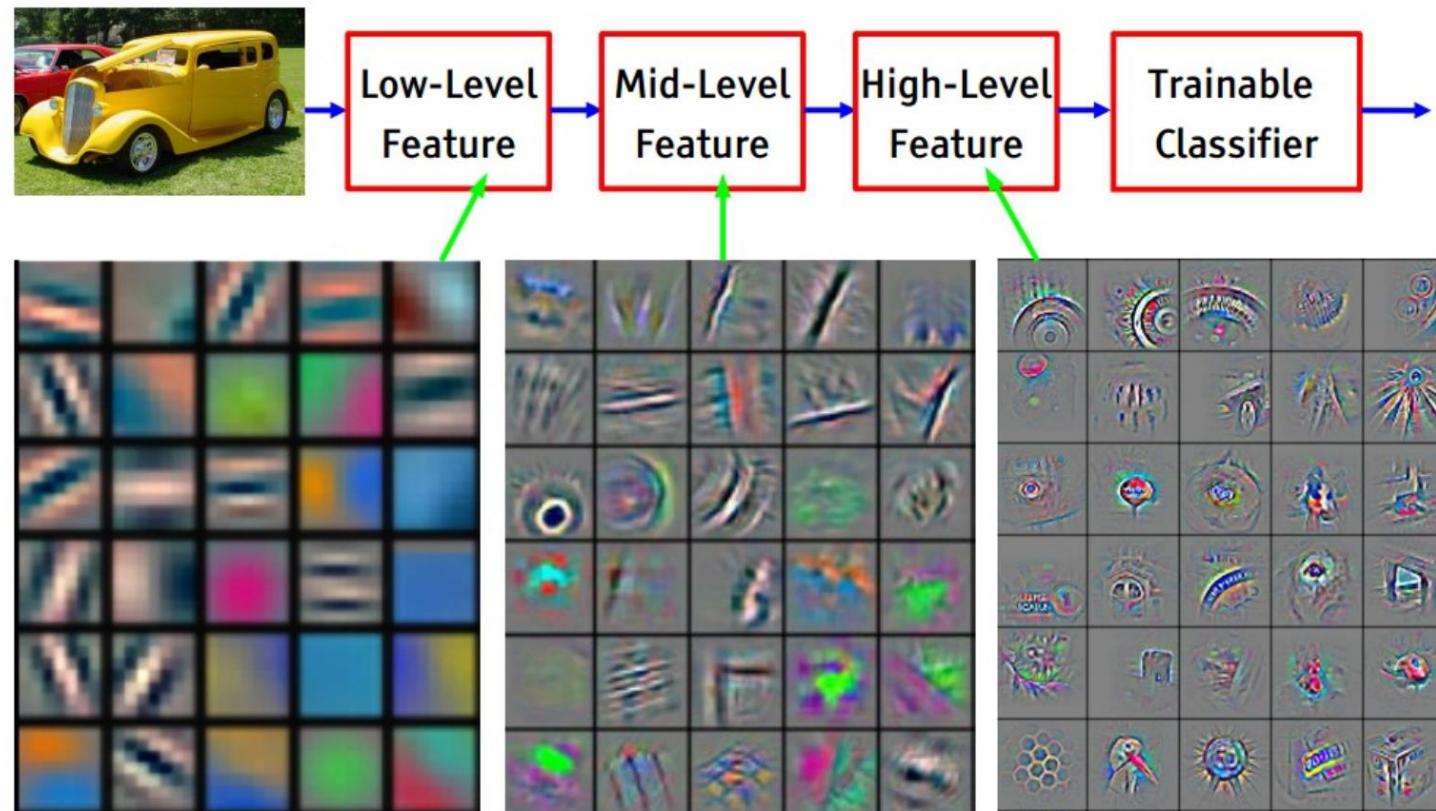


```
from tensorflow.keras.optimizers import Adam, SGD, RMSProp, RMSProp, Adagrad
```

```
model.compile(loss='binary_crossentropy', optimizer=Adam, metrics=['accuracy'])
```

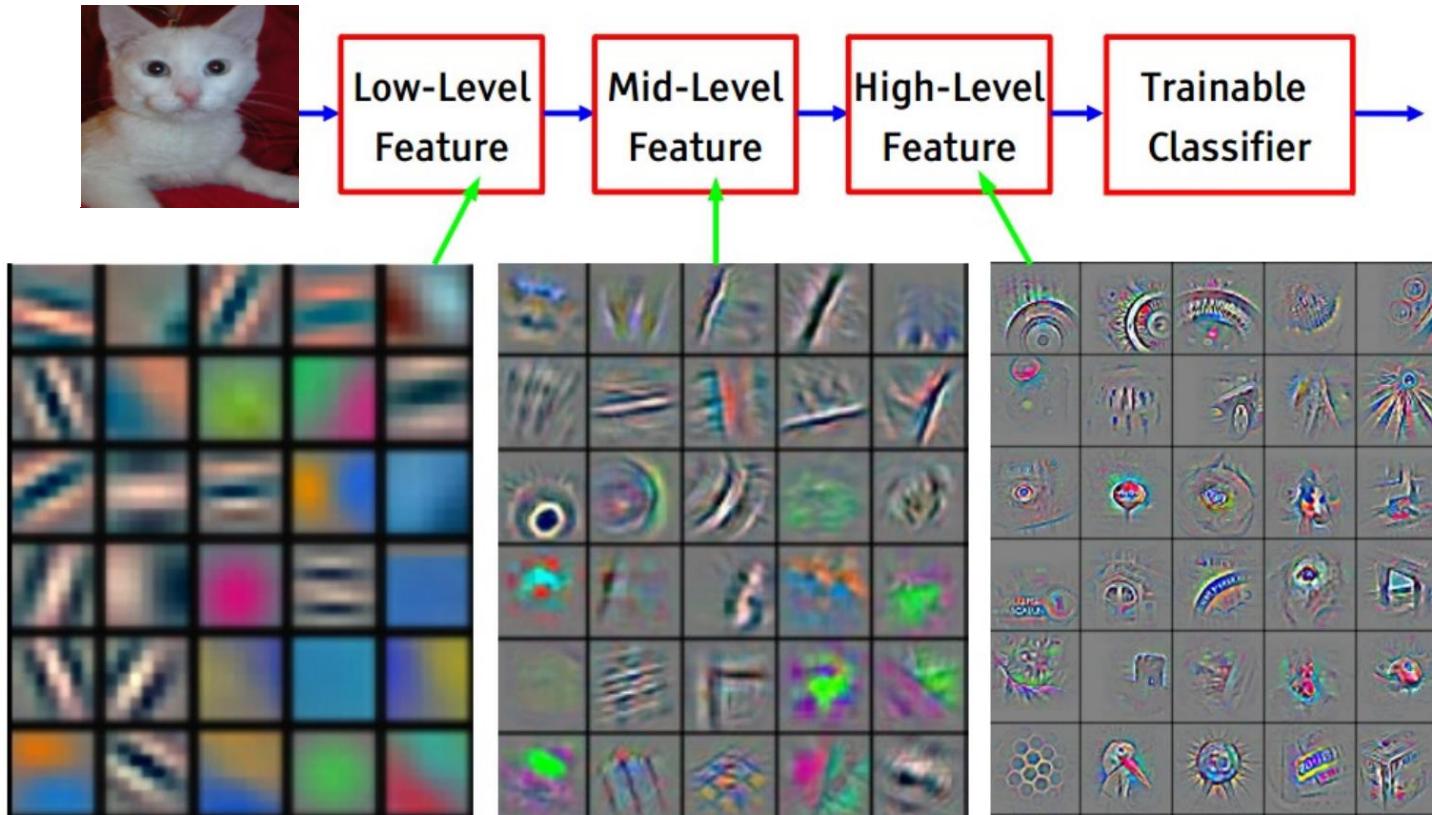
```
history = model.fit( train_generator, steps_per_epoch=8, epochs=30,verbose=1,  
validation_data = validation_generator, validation_steps=8)
```

Week 9: Artificial Intelligent – CNN – Feature Detector



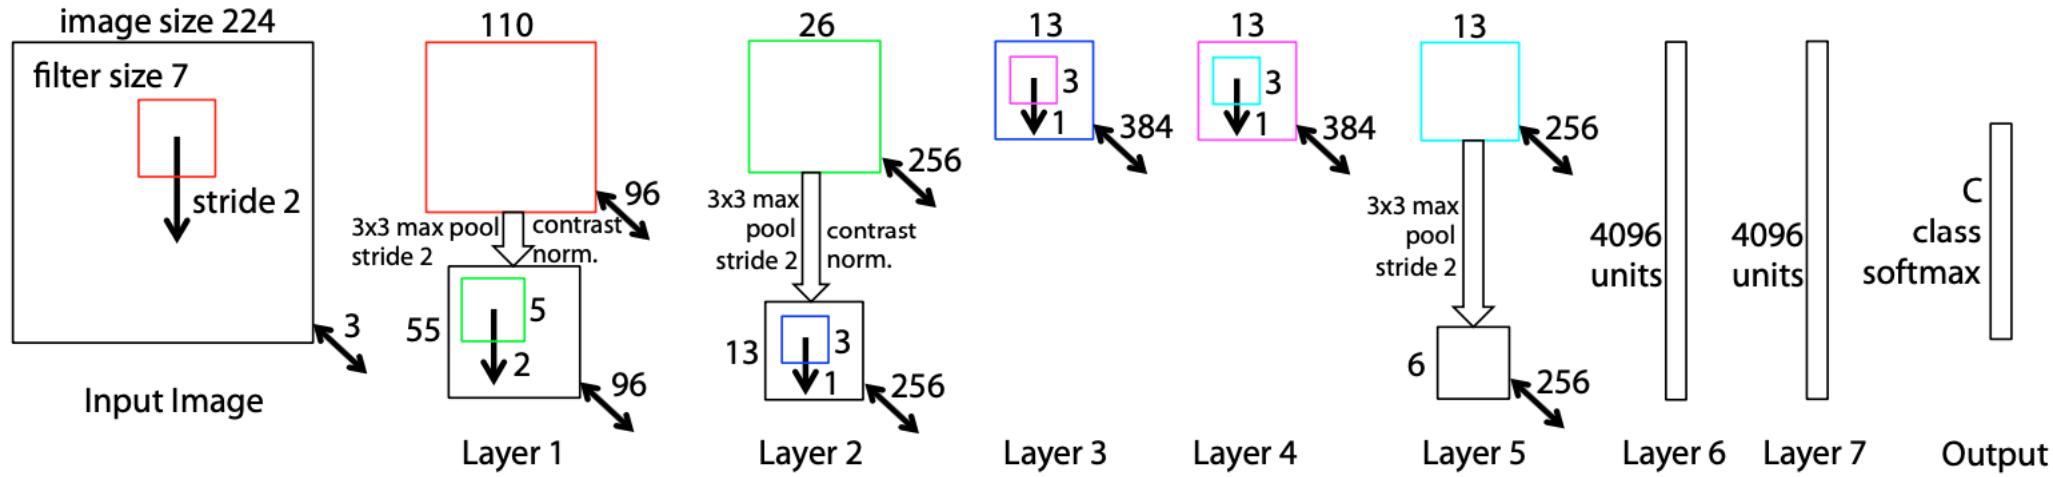
Feature visualization of convolutional net trained on ImageNet from [Zeiler & Fergus 2013]

Week 9: Artificial Intelligent – CNN – Feature Detector



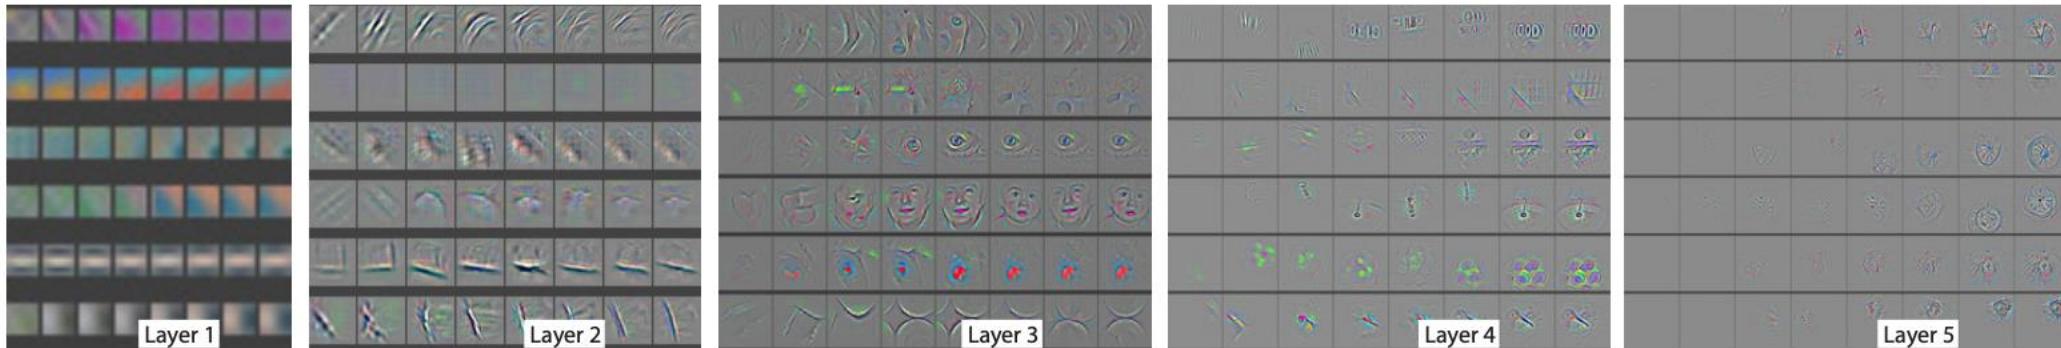
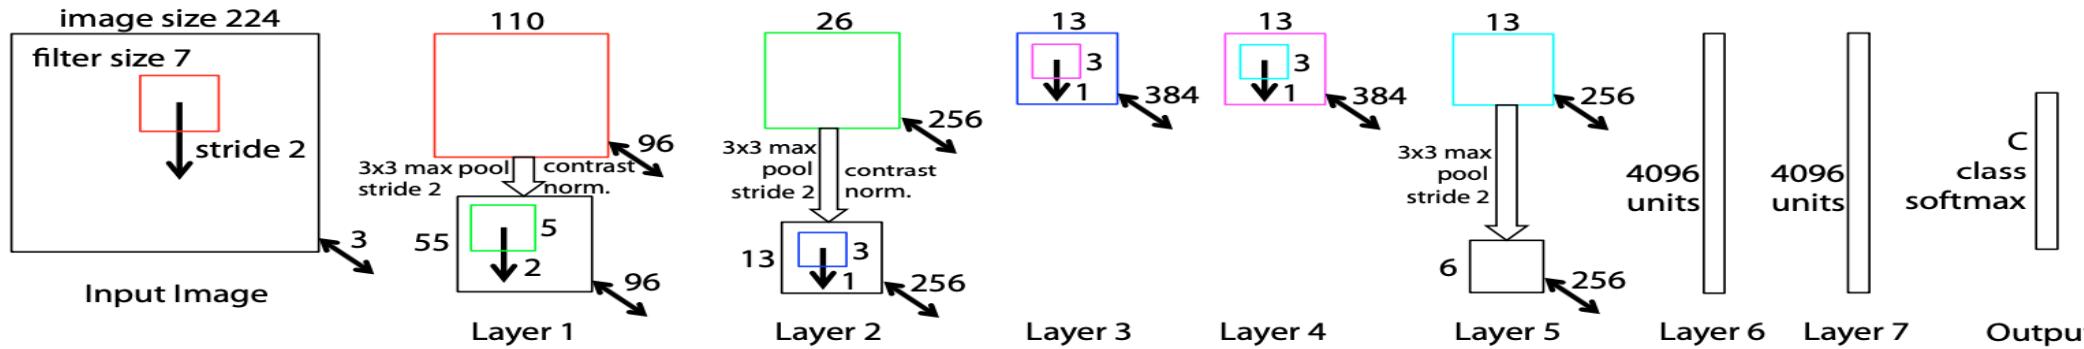
Feature visualization of convolutional net trained on ImageNet from [Zeiler & Fergus 2013]

Week 9: Artificial Intelligent – CNN – Feature Detector “ImageNet”



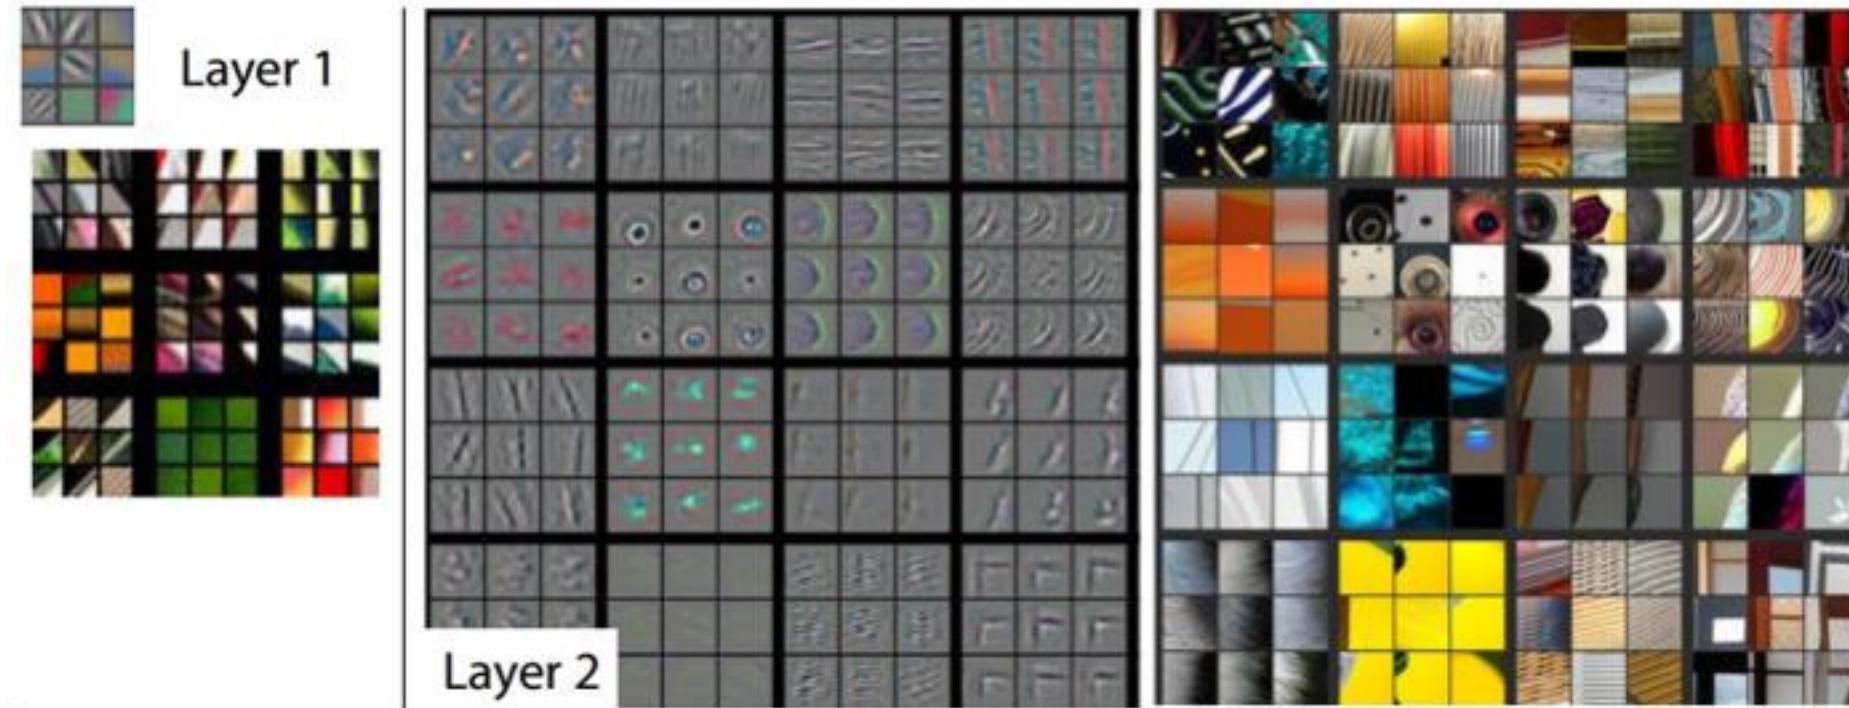
[Visualizing and Understanding Convolutional Networks](#) MD Zeiler · 2013

Week 9: Artificial Intelligent – CNN – Feature Detector “ImageNet”



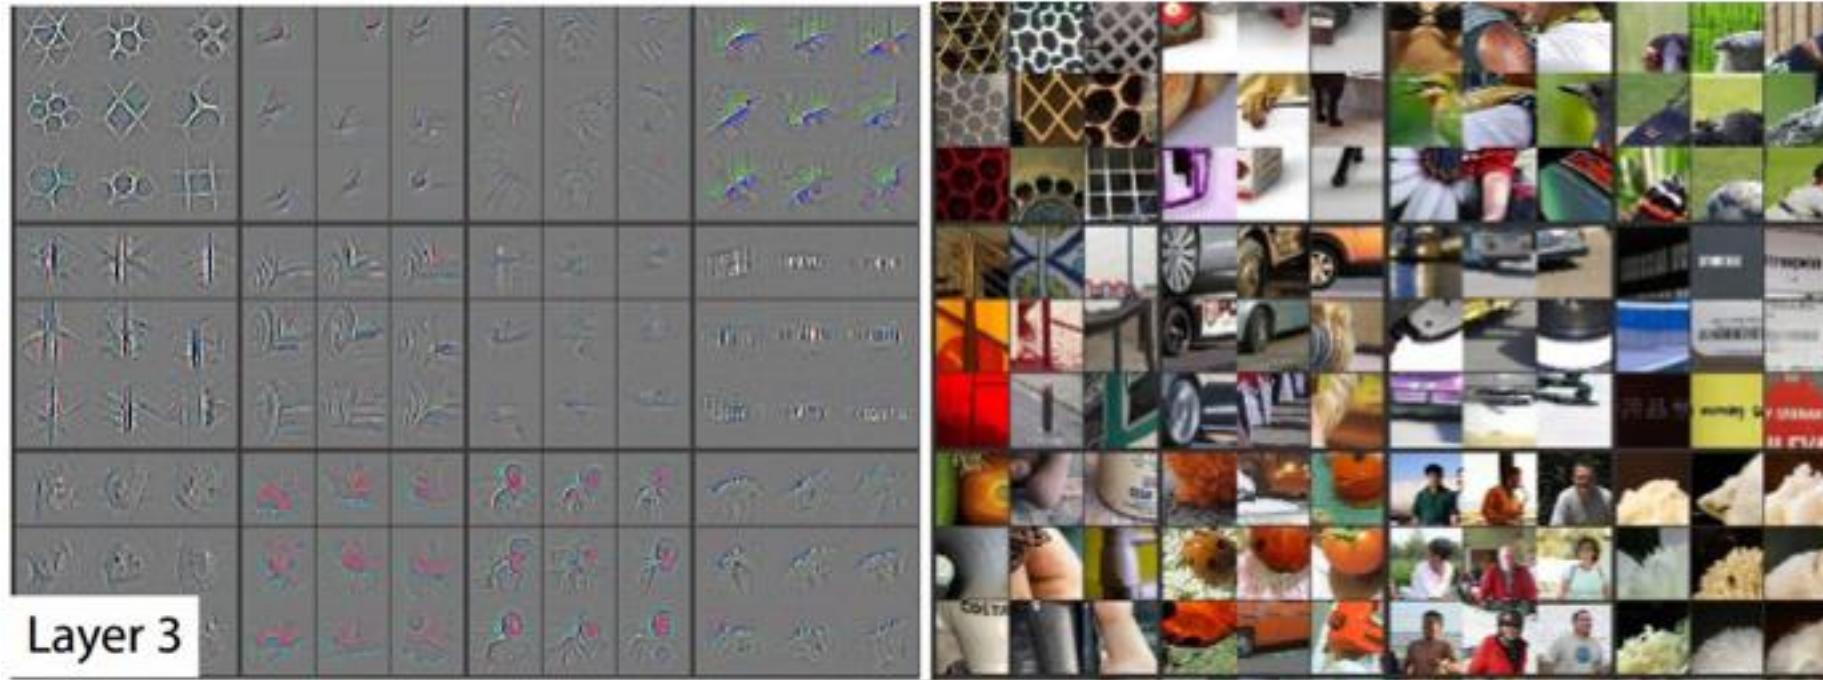
[Visualizing and Understanding Convolutional Networks](#) MD Zeiler · 2013

Week 9: Artificial Intelligent – CNN – Feature Detector “ImageNet”



[Visualizing and Understanding Convolutional Networks](#) MD Zeiler · 2013

Week 9: Artificial Intelligent – CNN – Feature Detector “ImageNet”



Visualizing and Understanding Convolutional Networks MD Zeiler · 2013

Week 9: Artificial Intelligent – CNN – Feature Detector “ImageNet”



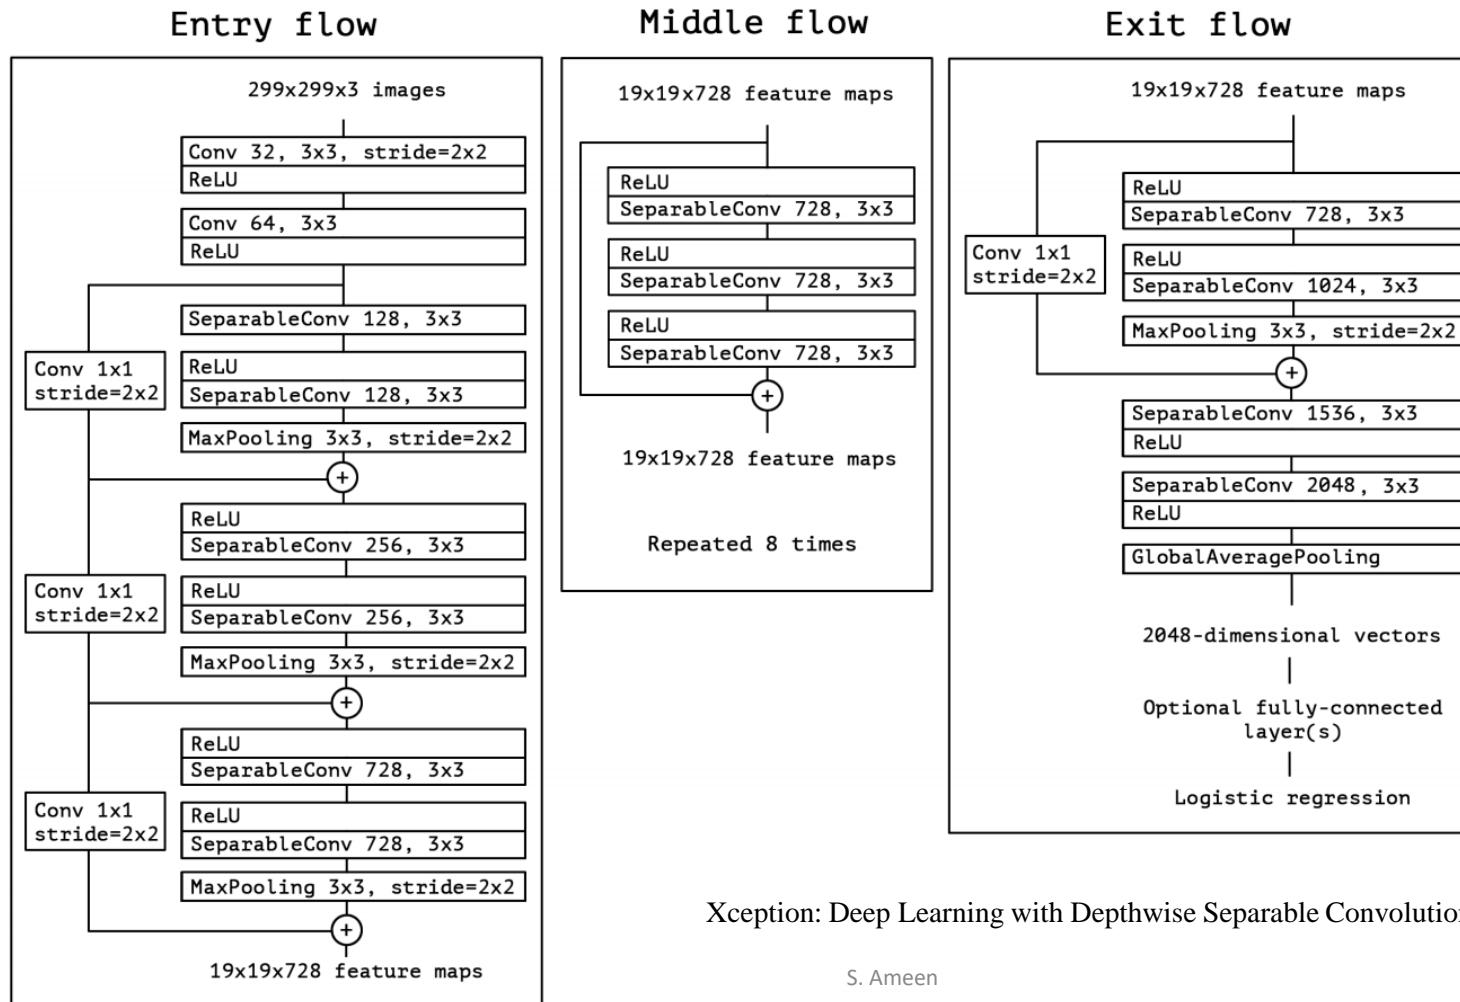
Visualizing and Understanding Convolutional Networks MD Zeiler · 2013

Week 9: Artificial Intelligent – CNN – Transfer Learning



	very similar dataset	very different dataset
very little data	Use Linear Classifier on top layer	You're in trouble... Try linear classifier from different stages
quite a lot of data	Finetune a few layers	Finetune a larger number of layers

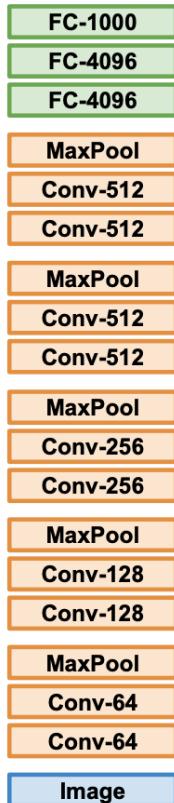
Week 9: Artificial Intelligent – CNN – Transfer Learning - Xception



Week 9: Artificial Intelligent – CNN – Transfer Learning



1. Train on Imagenet



AlexNet, VGG, ResNet, SqueezeNet, DenseNet, Inception v3,
GoogLeNet, ShuffleNet v2, MobileNet v2, ResNeXt, Wide
ResNet and MNASNet

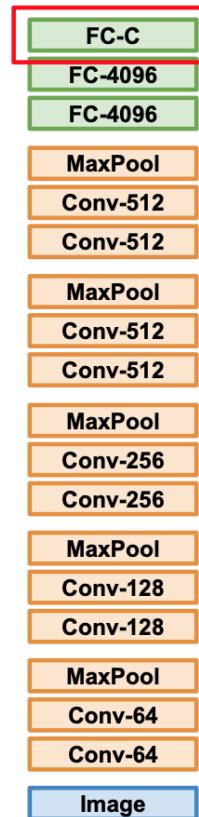
Week 9: Artificial Intelligent – CNN – Transfer Learning



1. Train on Imagenet



2. Small Dataset (C classes)



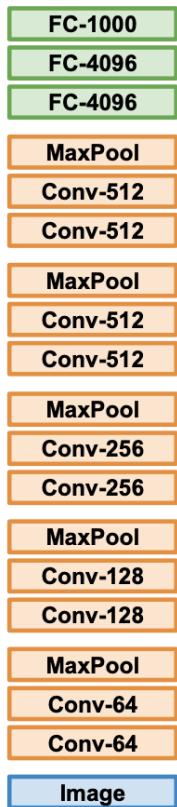
VGG, AlexNet, ...

```
base_model = keras.applications.Xception(  
    weights='imagenet', # Load weights pre-trained on  
    ImageNet.  
    input_shape=(150, 150, 3),  
    include_top=False) # Do not include the ImageNet  
    classifier at the top.
```

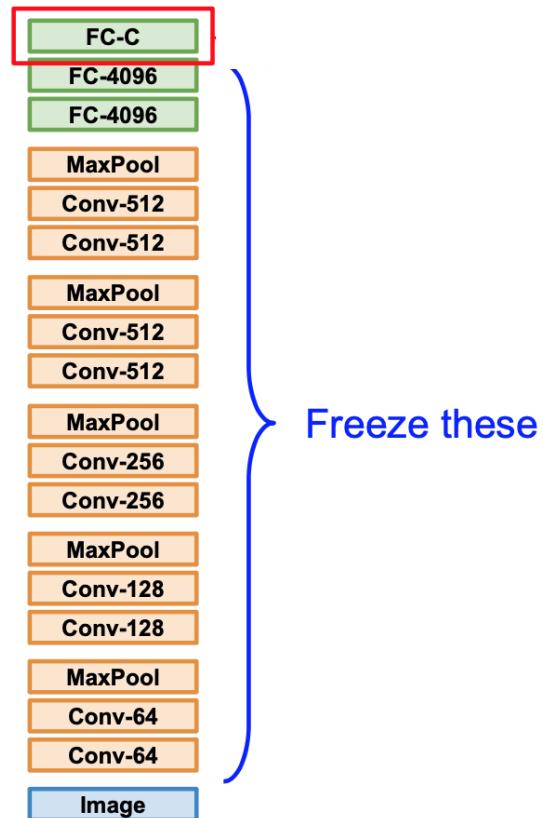
Week 9: Artificial Intelligent – CNN – Transfer Learning



1. Train on Imagenet



2. Small Dataset (C classes)



```
base_model = keras.applications.Xception(  
    weights='imagenet',  
    input_shape=(150, 150, 3),  
    include_top=False)
```

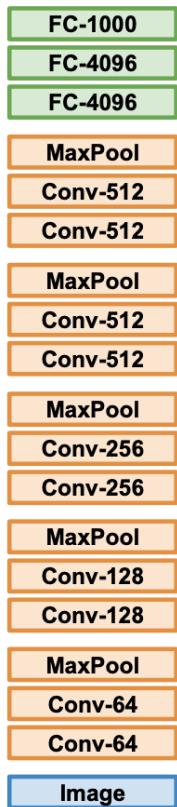
```
base_model.trainable = False
```

Session Code:
2675
9432

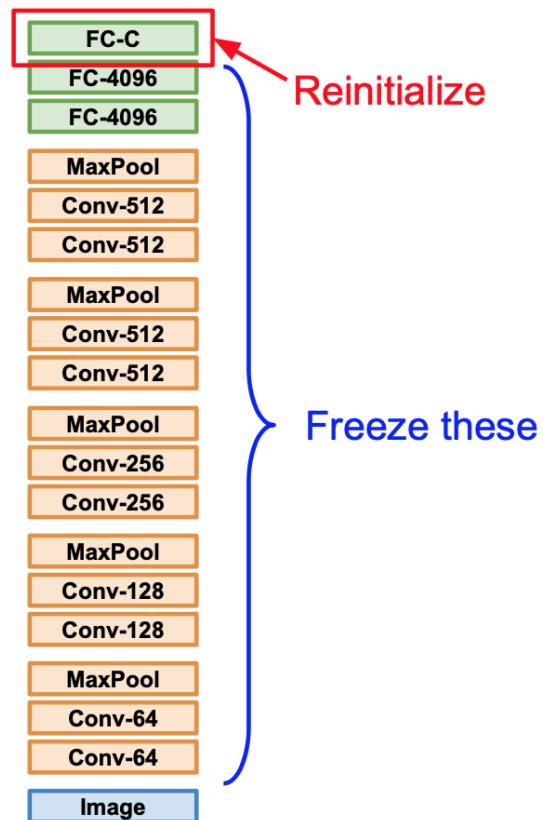
Week 9: Artificial Intelligent – CNN – Transfer Learning



1. Train on Imagenet



2. Small Dataset (C classes)

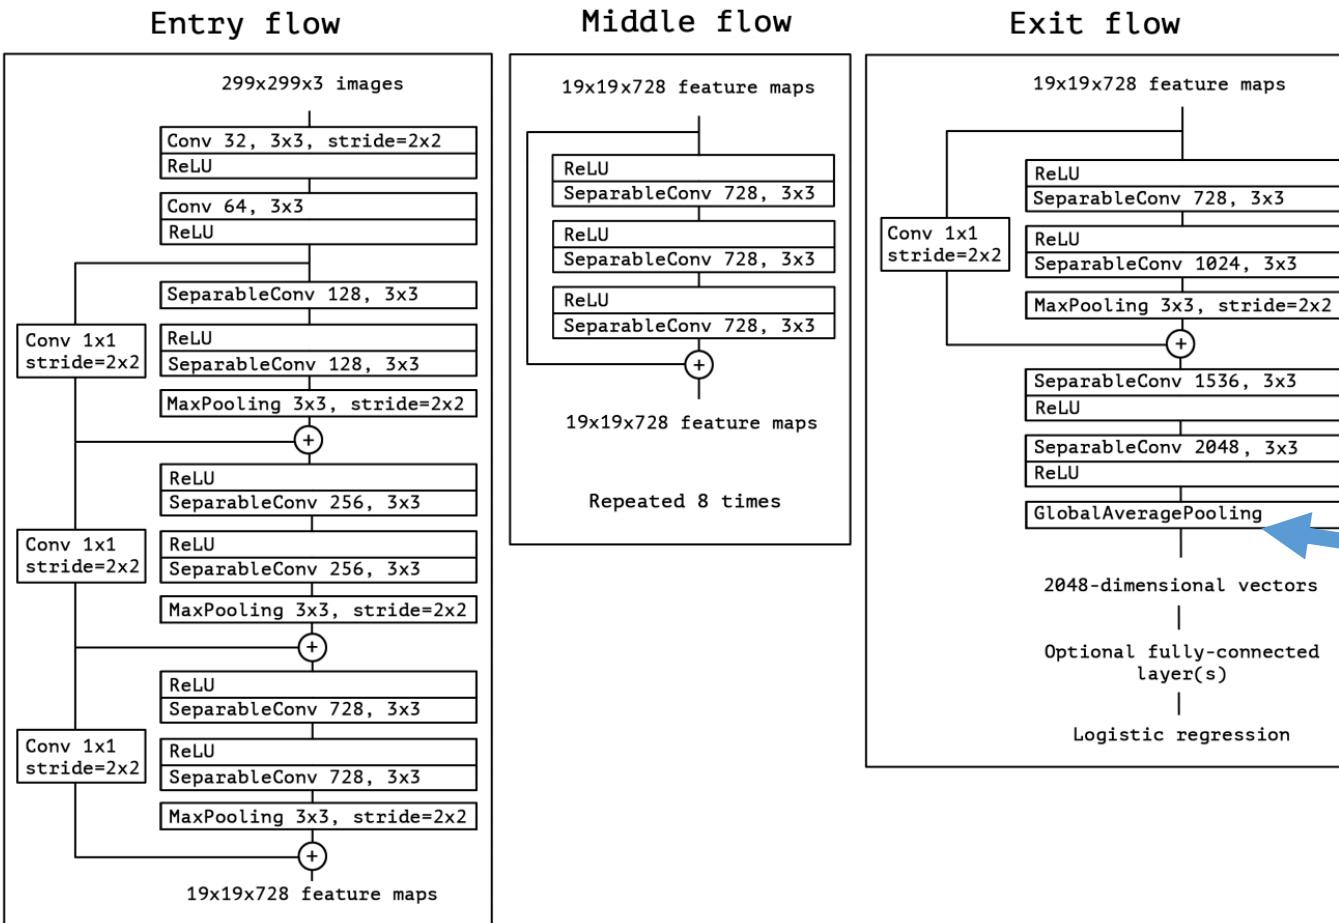


```
base_model = keras.applications.Xception(  
    weights='imagenet',  
    input_shape=(150, 150, 3),  
    include_top=False)
```

```
base_model.trainable = False
```

```
inputs = keras.Input(shape=(150, 150, 3))  
x = base_model(inputs, training=False)  
x = keras.layers.GlobalAveragePooling2D()(x)  
outputs = keras.layers.Dense(1)(x)  
model = keras.Model(inputs, outputs)
```

Week 9: Artificial Intelligent – CNN – Transfer Learning



```
base_model = keras.applications.Xception(
    weights='imagenet',
    input_shape=(150, 150, 3),
    include_top=False)
```

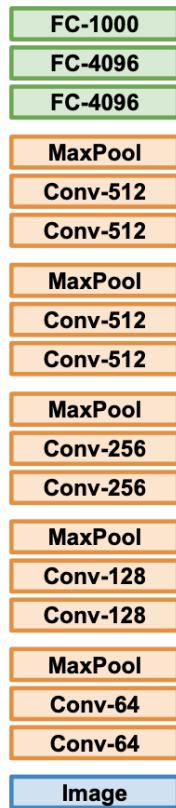
```
base_model.trainable = False
```

```
inputs = keras.Input(shape=(150, 150, 3))
x = base_model(inputs, training=False)
x = keras.layers.GlobalAveragePooling2D()(x)
outputs = keras.layers.Dense(1)(x)
model = keras.Model(inputs, outputs)
```

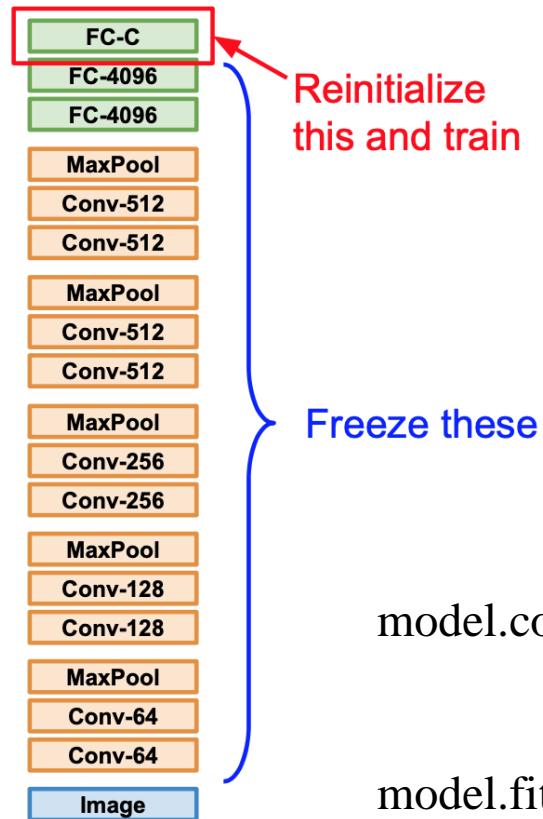
Week 9: Artificial Intelligent – CNN – Transfer Learning



1. Train on Imagenet



2. Small Dataset (C classes)



```
base_model = keras.applications.Xception(
```

```
    weights='imagenet',
```

```
    input_shape=(150, 150, 3),
```

```
    include_top=False)
```

```
base_model.trainable = False
```

```
inputs = keras.Input(shape=(150, 150, 3))
```

```
x = base_model(inputs, training=False)
```

```
x = keras.layers.GlobalAveragePooling2D()(x)
```

```
outputs = keras.layers.Dense(1)(x)
```

```
model = keras.Model(inputs, outputs)
```

```
model.compile(optimizer=keras.optimizers.Adam(),
```

```
    loss=keras.losses.BinaryCrossentropy(from_logits=True),
```

```
    metrics=[keras.metrics.BinaryAccuracy()])
```

```
model.fit(new_dataset, epochs=20)
```

Week 9: Artificial Intelligent – CNN – Transfer Learning



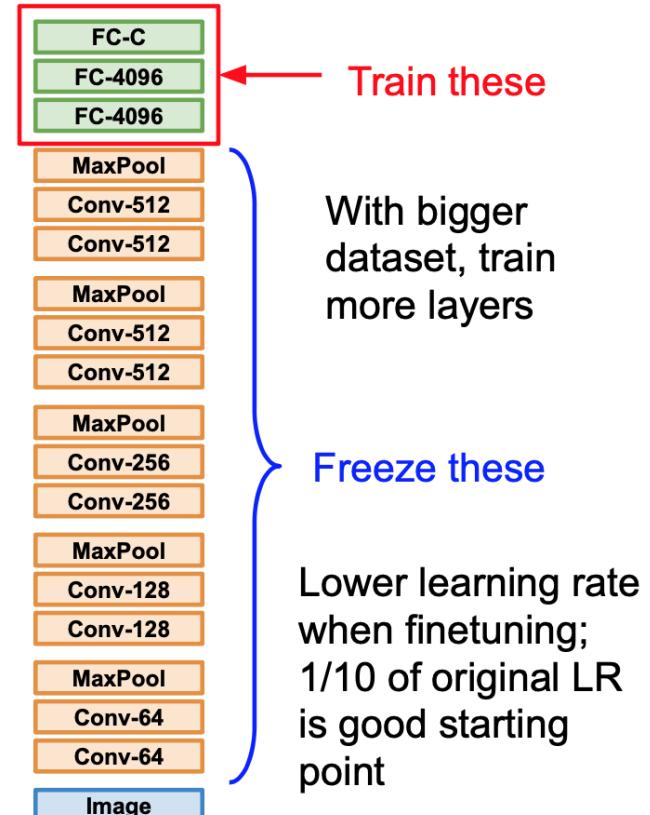
	very similar dataset	very different dataset
very little data	Use Linear Classifier on top layer	You're in trouble... Try linear classifier from different stages
quite a lot of data	Finetune a few layers	Finetune a larger number of layers

Week 9: Artificial Intelligent – CNN – Transfer Learning



base_model.trainable = True

3. Bigger dataset



Week 9: Artificial Intelligent – CNN – Transfer Learning



```
base_model.trainable = True
```

```
# Let's take a look to see how many layers are in the base model
print("Number of layers in the base model: ",
len(base_model.layers))
```

```
# Fine-tune from this layer onwards
fine_tune_at = 100
```

```
# Freeze all the layers before the `fine_tune_at` layer
for layer in base_model.layers[:fine_tune_at]:
    layer.trainable = False
```

3. Bigger dataset



Train these

With bigger dataset, train more layers

Freeze these

Lower learning rate when finetuning;
1/10 of original LR is good starting point

Week 9: Artificial Intelligent – CNN – Transfer Learning



```
base_model.trainable = True
```

```
len(base_model.layers))
```

```
fine_tune_at = 100
```

```
for layer in base_model.layers[:fine_tune_at]:  
    layer.trainable = False
```

```
model.compile(loss=tf.keras.losses.BinaryCrossentropy(from_logits=True),  
              optimizer = tf.keras.optimizers.RMSprop(lr=base_learning_rate/10),  
              metrics=['accuracy'])
```

```
history_fine = model.fit(train_dataset,  
                         epochs=total_epochs,  
                         initial_epoch=history.epoch[-1],  
                         validation_data=validation_dataset)
```

3. Bigger dataset



Train these

With bigger dataset, train more layers

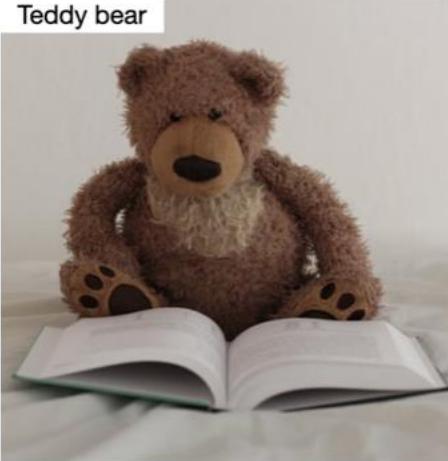
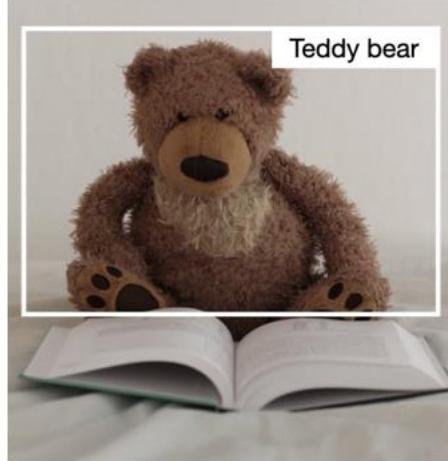
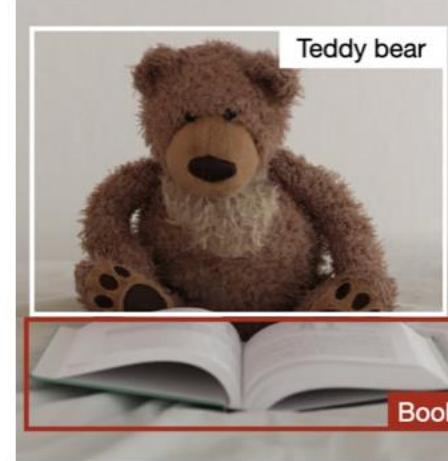
Freeze these

Lower learning rate when finetuning;
1/10 of original LR is good starting point

Week 9: CNN – Object detection



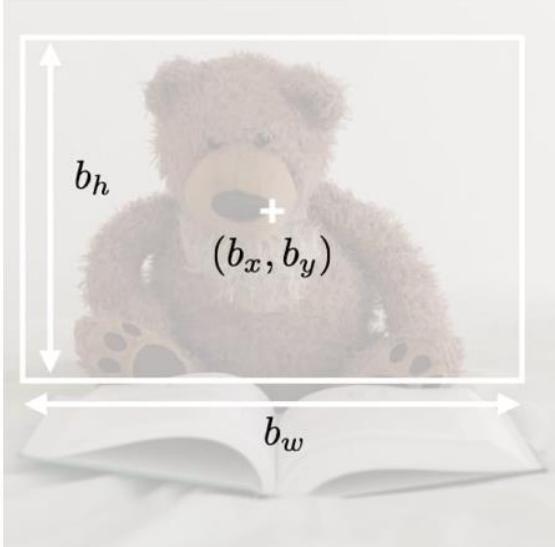
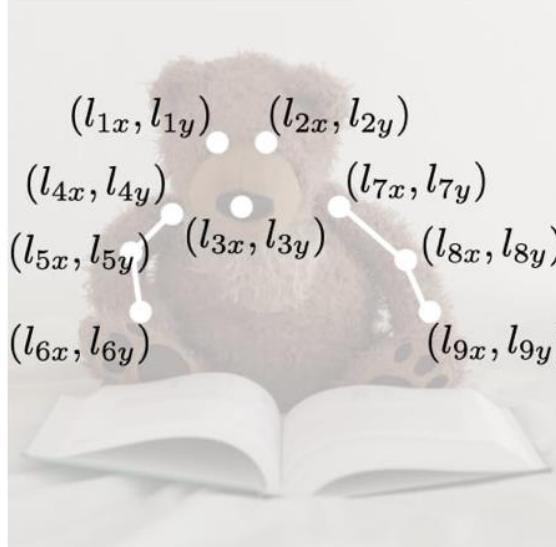
□ **Types of models** — There are 3 main types of object recognition algorithms, for which the nature of what is predicted is different. They are described in the table below:

Image classification	Classification w. localization	Detection
 Teddy bear	 Teddy bear	 Teddy bear Book
<ul style="list-style-type: none">• Classifies a picture• Predicts probability of object	<ul style="list-style-type: none">• Detects an object in a picture• Predicts probability of object and where it is located	<ul style="list-style-type: none">• Detects up to several objects in a picture• Predicts probabilities of objects and where they are located
Traditional CNN	Simplified YOLO, R-CNN ^{S. Ameen}	YOLO, R-CNN

Week 9: CNN – Object detection



- **Detection** — In the context of object detection, different methods are used depending on whether we just want to locate the object or detect a more complex shape in the image. The two main ones are summed up in the table below:

Bounding box detection	Landmark detection
<ul style="list-style-type: none">• Detects the part of the image where the object is located	<ul style="list-style-type: none">• Detects a shape or characteristics of an object (e.g. eyes)• More granular
 A photograph of a brown teddy bear sitting on an open book. A white bounding box is drawn around the bear. Inside the box, a center point is marked with a plus sign and labeled (b_x, b_y) . The vertical height of the box is labeled b_h and the horizontal width is labeled b_w .	 A photograph of a brown teddy bear sitting on an open book. Nine specific points on the bear's body are marked with white dots and labeled with coordinates: (l_{1x}, l_{1y}) , (l_{2x}, l_{2y}) , (l_{3x}, l_{3y}) , (l_{4x}, l_{4y}) , (l_{5x}, l_{5y}) , (l_{6x}, l_{6y}) , (l_{7x}, l_{7y}) , (l_{8x}, l_{8y}) , and (l_{9x}, l_{9y}) . Lines connect these points to form a polygonal shape.
Box of center (b_x, b_y) , height b_h and width b_w	Reference points $(l_{1x}, l_{1y}), \dots, (l_{nx}, l_{ny})$

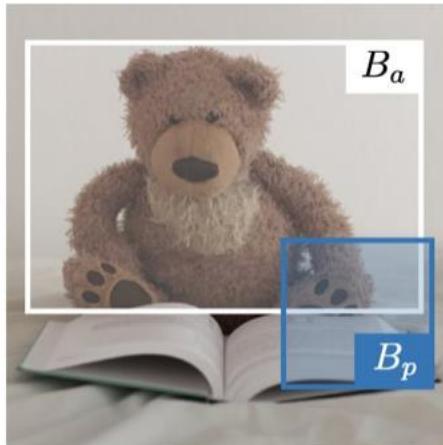
Week 9: CNN – Object detection



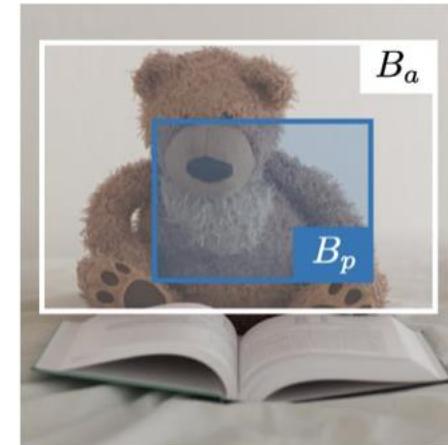
- **Intersection over Union** — Intersection over Union, also known as IoU, is a function that quantifies how correctly positioned a predicted bounding box B_p is over the actual bounding box B_a . It is defined as:

$$\text{IoU}(B_p, B_a) = \frac{B_p \cap B_a}{B_p \cup B_a}$$

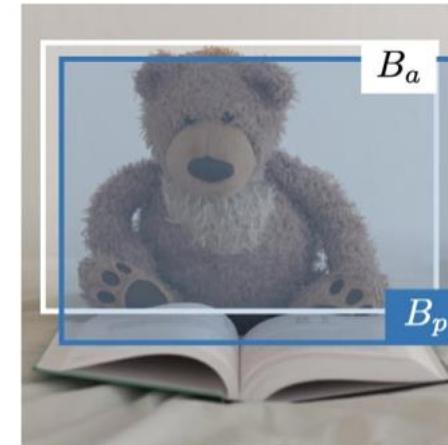
Remark: we always have $\text{IoU} \in [0, 1]$. By convention, a predicted bounding box B_p is considered as being reasonably good if $\text{IoU}(B_p, B_a) \geq 0.5$.



$$\text{IoU}(B_p, B_a) = 0.1$$



$$\text{IoU}(B_p, B_a) = 0.5$$



$$\text{IoU}(B_p, B_a) = 0.9$$

Week 9: CNN – Face verification and recognition



□ **Types of models** — Two main types of model are summed up in table below:

Face verification	Face recognition
<ul style="list-style-type: none">• Is this the correct person?• One-to-one lookup	<ul style="list-style-type: none">• Is this one of the K persons in the database?• One-to-many lookup
<p>Query</p> <p>The diagram illustrates the Face verification process. It shows two 'Query' images at the top: a brown teddy bear and a white teddy bear. Below each query is a green circle containing a checkmark or an 'X'. A green line connects the first query to its corresponding green circle, and a red line connects the second query to its red circle. At the bottom are two 'Reference' images: a brown teddy bear and a white teddy bear. The reference images are enclosed in boxes that match the colors of the circles above them (green for the brown bear, red for the white bear).</p> <p>Database</p> <p>The diagram illustrates the Face recognition process. It shows a single 'Query' image of a brown teddy bear at the top. A green line connects this query to three images in a 'Database' at the bottom: a brown teddy bear, a white teddy bear, and a dog. The database images are enclosed in boxes that match the color of the green line (green for the brown bear, red for the white bear and dog).</p>	

Week 9: CNN – Neural style transfer



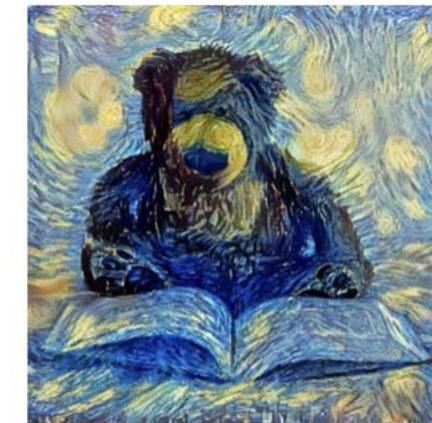
- **Motivation** — The goal of neural style transfer is to generate an image G based on a given content C and a given style S .



+



=



Content C

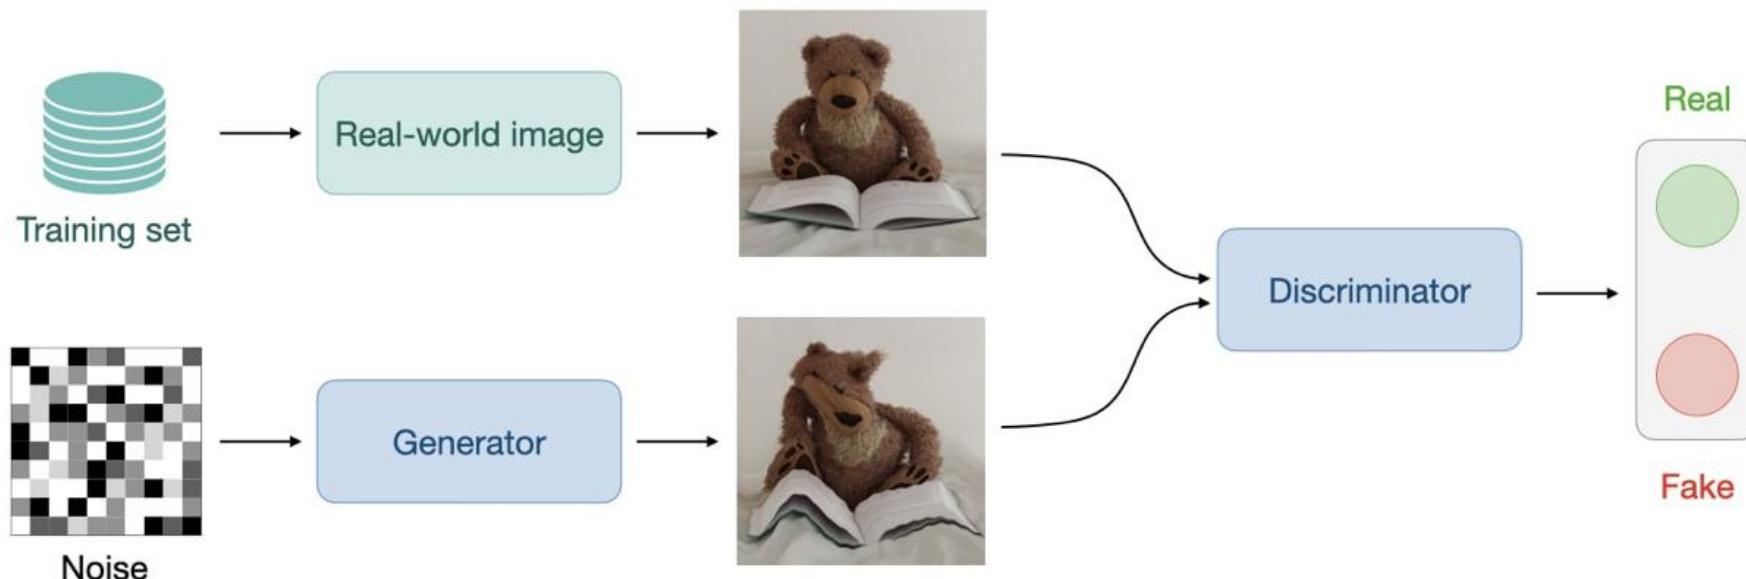
Style S

Generated image G

Week 9: CNN – GAN

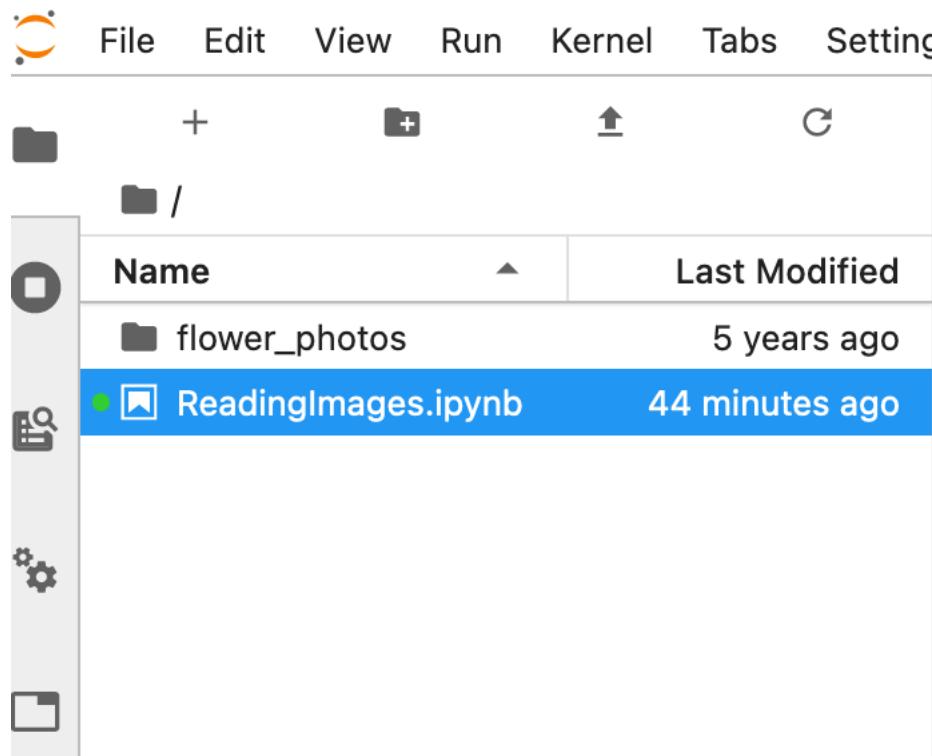


- **Generative Adversarial Network** — Generative adversarial networks, also known as GANs, are composed of a generative and a discriminative model, where the generative model aims at generating the most truthful output that will be fed into the discriminative which aims at differentiating the generated and true image.

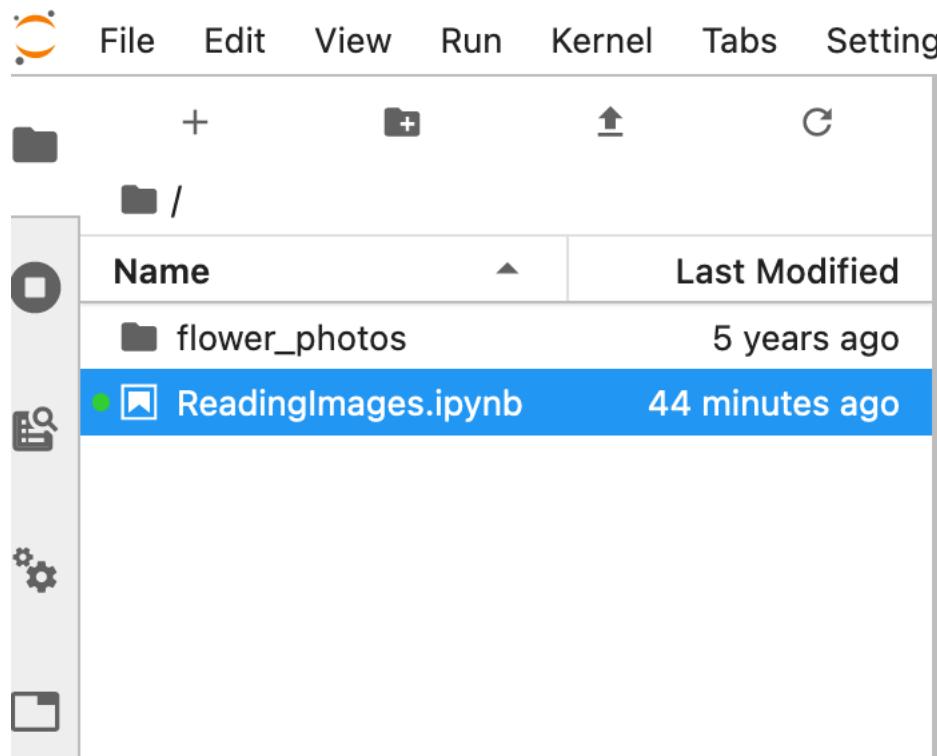


Remark: use cases using variants of GANs include text to image, music generation and synthesis.

Week 9 – Reading the Images



Week 9 – Reading the Images

A screenshot of a file explorer window showing the contents of a folder named 'flower_photos'. The folder path is displayed at the top: '/ flower_photos /'. The table lists the files and subfolders with their names and last modified dates.

Name	Last Modified
daisy	5 years ago
dandelion	5 years ago
roses	5 years ago
sunflowers	5 years ago
tulips	5 years ago
LICENSE.txt	5 years ago

Week 9 – Reading the Images



```
[1]: import numpy as np
      import os
      import PIL
      import PIL.Image
      import tensorflow as tf
      import tensorflow_datasets as tfds
```

```
[2]: data_dir = "flower_photos"
```

```
[4]: batch_size = 32
      img_height = 180
      img_width = 180
```

Week 9 – Reading the Images

Loading



```
: train_ds = tf.keras.preprocessing.image_dataset_from_directory(  
    data_dir,  
    validation_split=0.2,  
    subset="training",  
    seed=123,  
    image_size=(img_height, img_width),  
    batch_size=batch_size)
```

Found 3670 files belonging to 5 classes.
Using 2936 files for training.

```
: val_ds = tf.keras.preprocessing.image_dataset_from_directory(  
    data_dir,  
    validation_split=0.2,  
    subset="validation",  
    seed=123,  
    image_size=(img_height, img_width),  
    batch_size=batch_size)
```

Week 9 – Reading the Images

Loading



```
from tensorflow.keras import layers  
  
normalization_layer = tf.keras.layers.experimental.preprocessing.Rescaling(1./255)
```

Week 9 – Reading the Images

Loading



```
AUTOTUNE = tf.data.AUTOTUNE

train_ds = train_ds.cache().prefetch(buffer_size=AUTOTUNE)
val_ds = val_ds.cache().prefetch(buffer_size=AUTOTUNE)
```

Week 9 – Reading the Images Loading / Modelling



```
num_classes = 5

model = tf.keras.Sequential([
    layers.experimental.preprocessing.Rescaling(1./255),
    layers.Conv2D(32, 3, activation='relu'),
    layers.MaxPooling2D(),
    layers.Conv2D(32, 3, activation='relu'),
    layers.MaxPooling2D(),
    layers.Conv2D(32, 3, activation='relu'),
    layers.MaxPooling2D(),
    layers.Flatten(),
    layers.Dense(128, activation='relu'),
    layers.Dense(num_classes)
])
```

Week 9 – Reading the Images Loading / Modelling



```
model.compile(  
    optimizer='adam',  
    loss=tf.losses.SparseCategoricalCrossentropy(from_logits=True),  
    metrics=['accuracy'])
```

Week 9 – Reading the Images Loading / Modelling / Training



```
model.fit(  
    train_ds,  
    validation_data=val_ds,  
    epochs=1  
)  
92/92 [=====] - 47s 501ms/step - loss: 1.5055 - accuracy: 0.3228 - val_loss: 1.1538 - val_accuracy: 0.5368
```

Lecture - Week 9 : AI



- Next - RNN



Week 9 – : Deep Learning Modelling Sequences



Sentences:

$x_1 \quad x_2 \quad x_3 \quad x_4 \quad x_5 \quad x_6 \quad x_7 \quad x_8 \quad x_9 \quad x_{10} \quad x_{11} \quad x_{12}$

Paris Talks Set Stage for Action as Risks to the Climate Rise

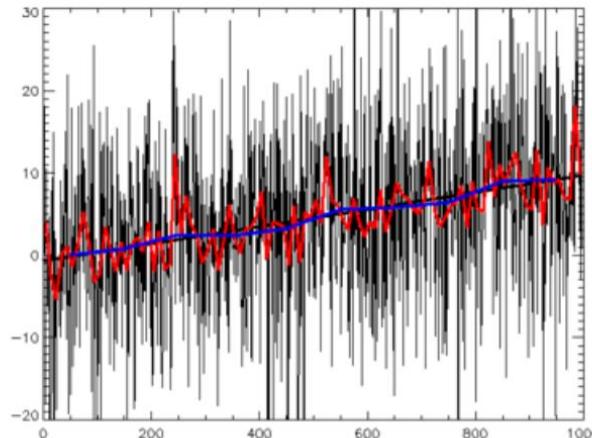
Week 9 – : Deep Learning Modelling Sequences



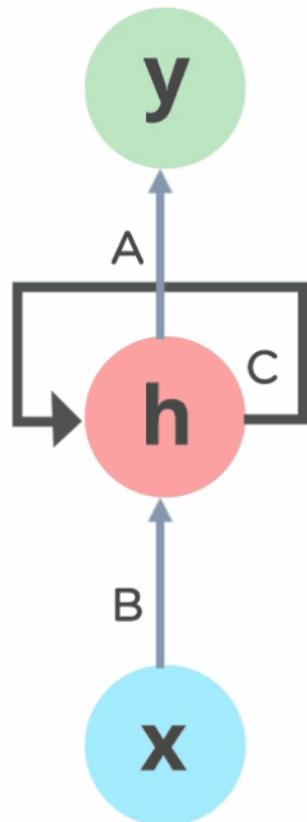
Sentences:

$x_1 \quad x_2 \quad x_3 \quad x_4 \quad x_5 \quad x_6 \quad x_7 \quad x_8 \quad x_9 \quad x_{10} \quad x_{11} \quad x_{12}$
Paris Talks Set Stage for Action as Risks to the Climate Rise

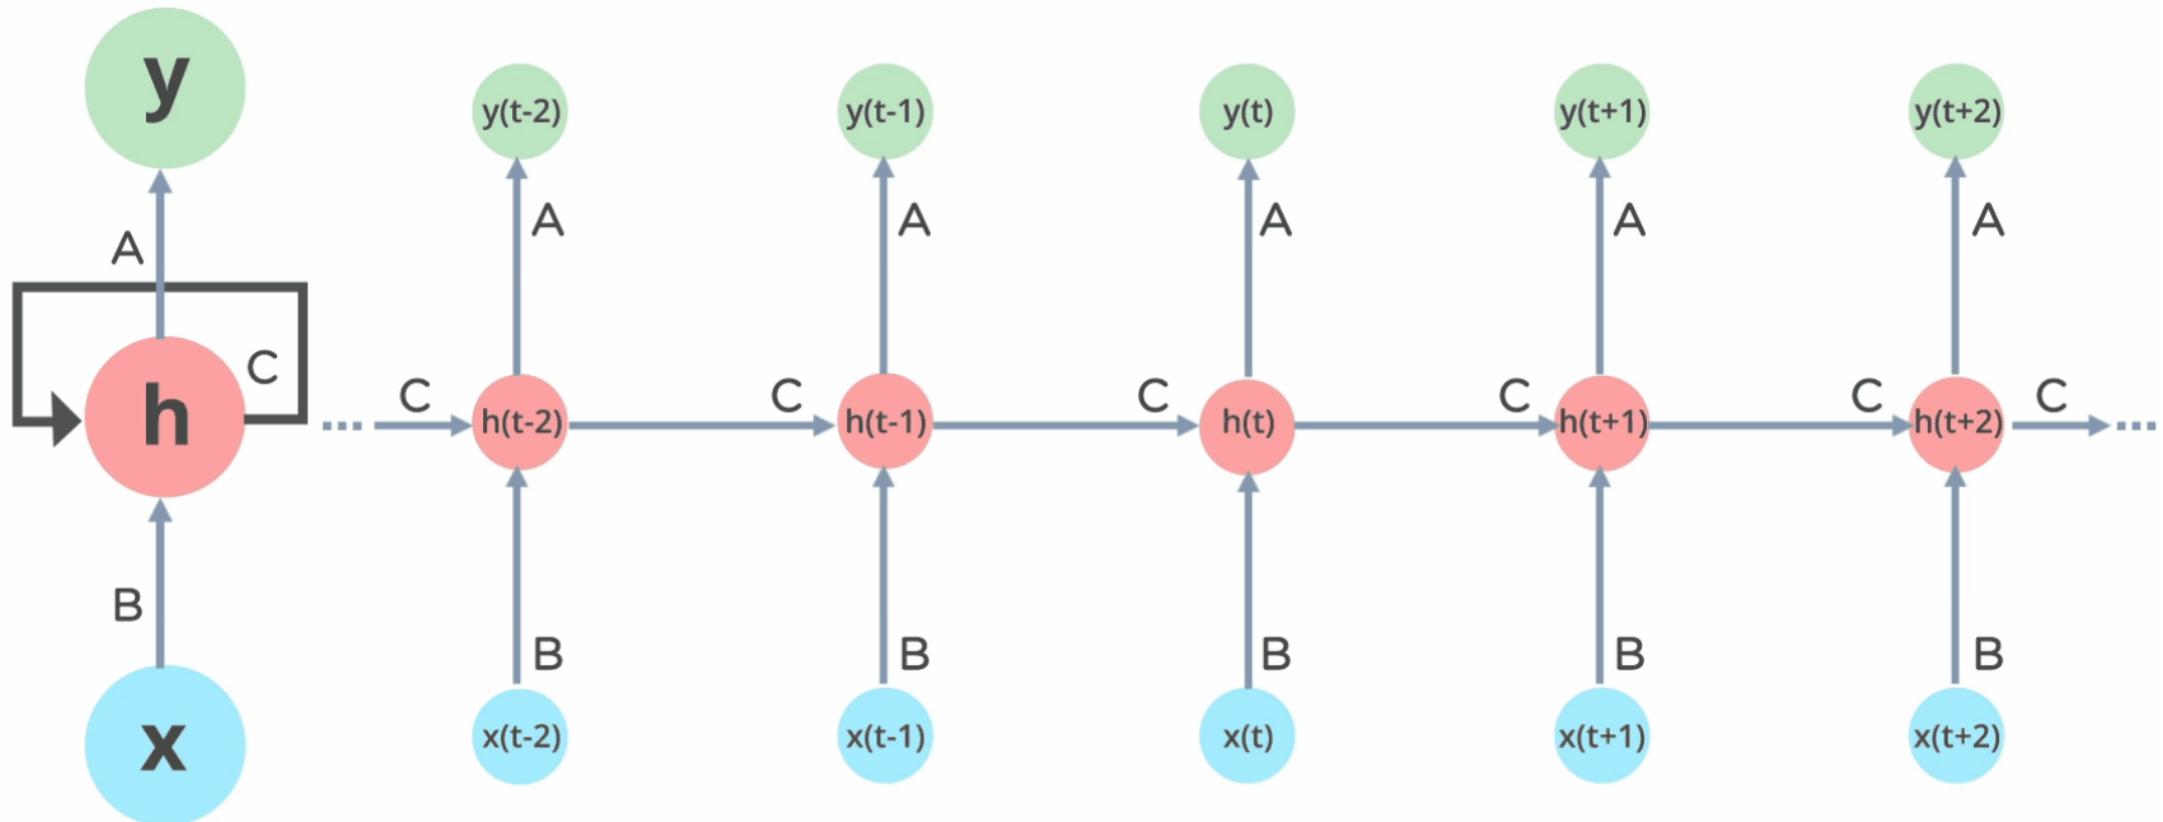
Time series:



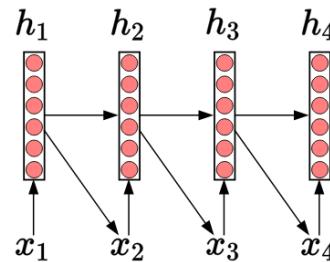
Week 9 – : Deep Learning Modelling Sequences



Week 9 – : Deep Learning Recurrent neural networks (RNN)



Week 9 – : Deep Learning RNN – Prediction



$h_1 = \text{Encode}(x_1)$

$x_2 \sim \text{Decode}(h_1)$

$h_2 = \text{Encode}(h_1, x_2)$

$x_3 \sim \text{Decode}(h_2)$

$h_3 = \text{Encode}(h_2, x_3)$

$x_4 \sim \text{Decode}(h_3)$

$h_4 = \text{Encode}(h_3, x_4)$

Update context vector:

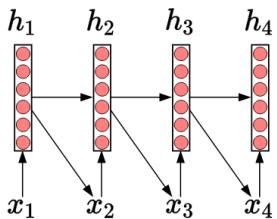
$h_t = \text{Encode}(h_{t-1}, x_t)$

Predict next character:

$x_{t+1} = \text{Decode}(h_t)$

context h_t compresses x_1, \dots, x_t

Week 9 – : Deep Learning RNN – Prediction



$$h_1 = \text{Encode}(x_1)$$

$$x_2 \sim \text{Decode}(h_1)$$

$$h_2 = \text{Encode}(h_1, x_2)$$

$$x_3 \sim \text{Decode}(h_2)$$

$$h_3 = \text{Encode}(h_2, x_3)$$

$$x_4 \sim \text{Decode}(h_3)$$

$$h_4 = \text{Encode}(h_3, x_4)$$

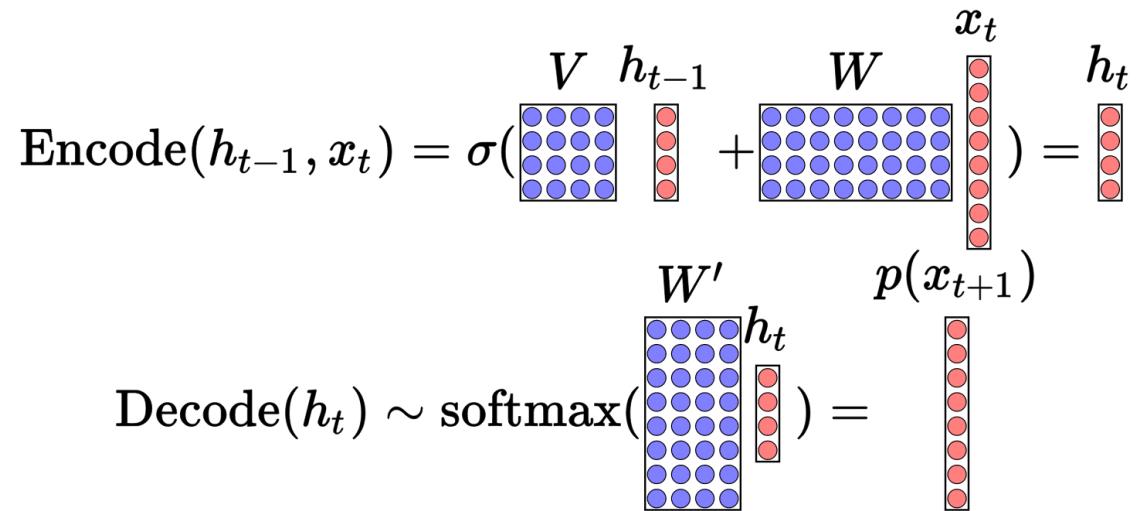
Update context vector:

$$h_t = \text{Encode}(h_{t-1}, x_t)$$

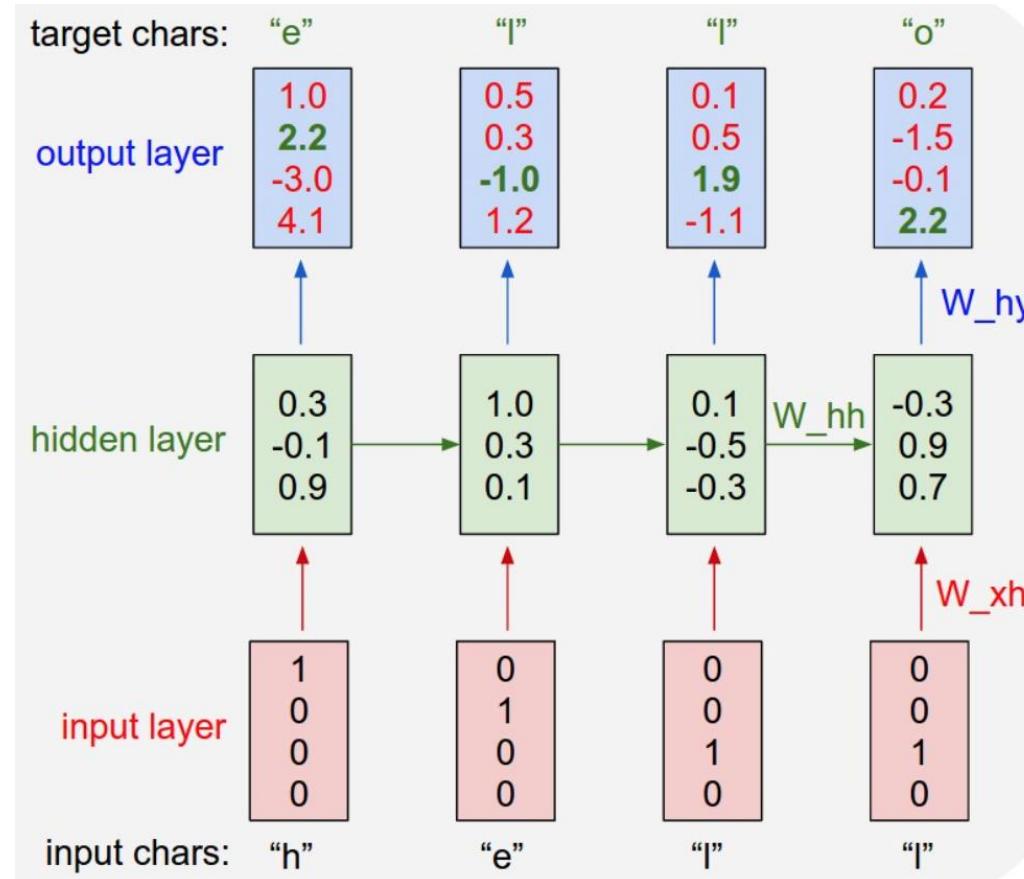
Predict next character:

$$x_{t+1} = \text{Decode}(h_t)$$

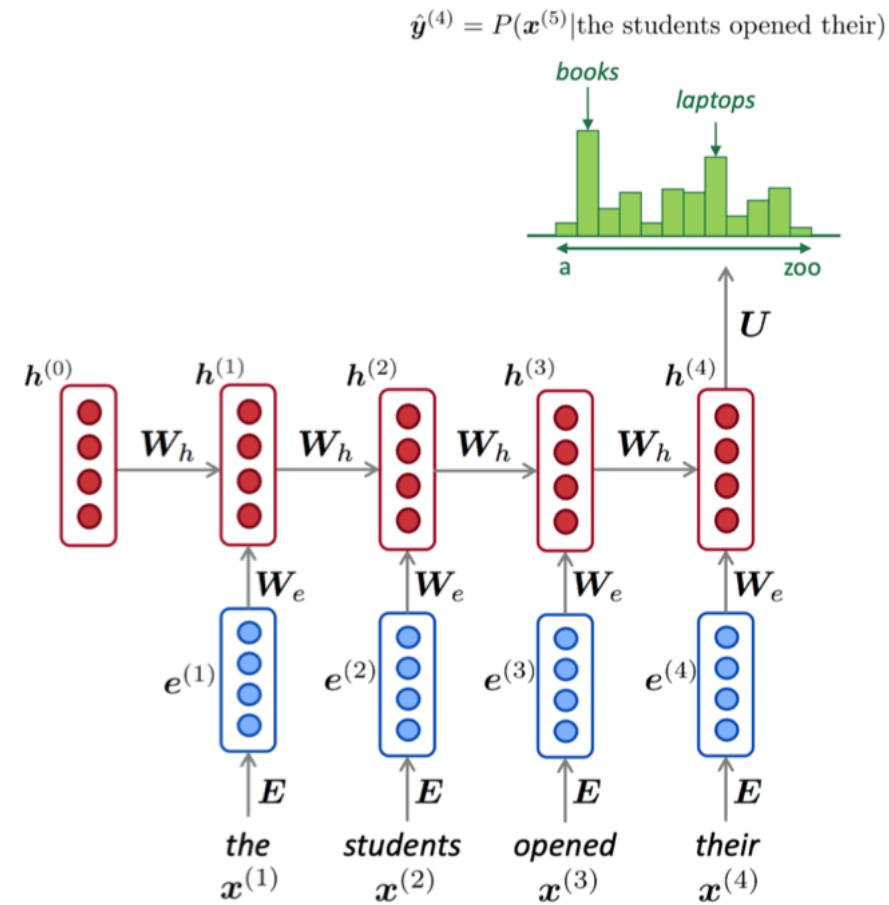
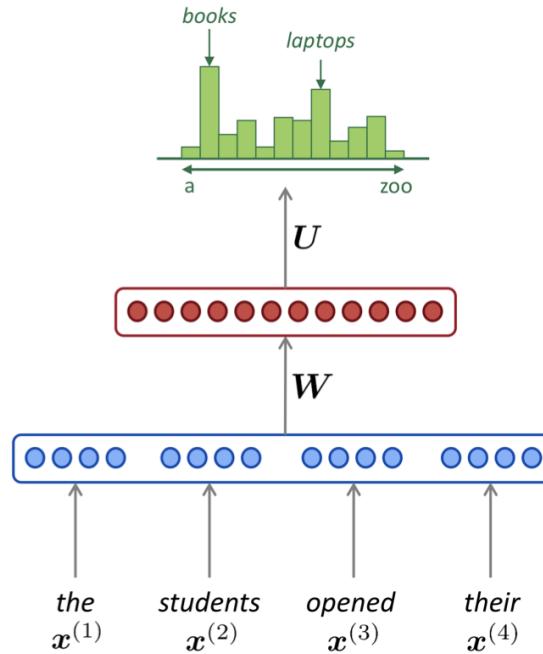
context h_t compresses x_1, \dots, x_t



Week 9 – : Deep Learning RNN – Prediction – Assignment



Week 9 – : Deep Learning Example NN VS RNN



Week 9 – : Deep Learning RNN – Coding



```
simple_rnn = tf.keras.layers.SimpleRNN(4)
```

Week 9 – : Deep Learning RNN – Coding



```
simple_rnn = tf.keras.layers.SimpleRNN(  
    4, return_sequences=True, return_state=True)
```

return_sequences: Boolean. Whether to return the last output in the output sequence, or the full sequence.
Default: **False**.

return_state: Boolean. Whether to return the last state in addition to the output. Default: **False**

Week 9 – : Deep Learning RNN – Coding



```
import numpy as np
import tensorflow as tf
from tensorflow import keras
from tensorflow.keras import layers
model = keras.Sequential()
from tensorflow.keras.layers import SimpleRNN, LSTM, Dense
```

Week 9 – : Deep Learning RNN – Coding



```
# define model
model = keras.Sequential()
model.add(SimpleRNN(100, input_shape=(5,1)))
model.add(Dense(1))
```

Week 9 – : Deep Learning RNN – Coding



```
# define model
model = keras.Sequential()
model.add(SimpleRNN(100, input_shape=(5,1)))
model.add(Dense(1))
```

```
model.summary()
```

```
Model: "sequential_7"
```

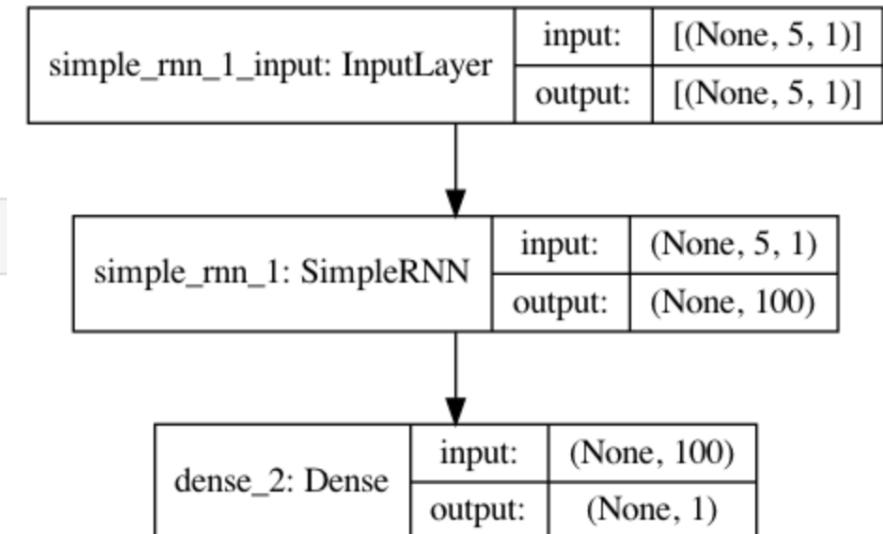
Layer (type)	Output Shape	Param #
simple_rnn_1 (SimpleRNN)	(None, 100)	10200
dense_2 (Dense)	(None, 1)	101

```
Total params: 10,301
```

```
Trainable params: 10,301
```

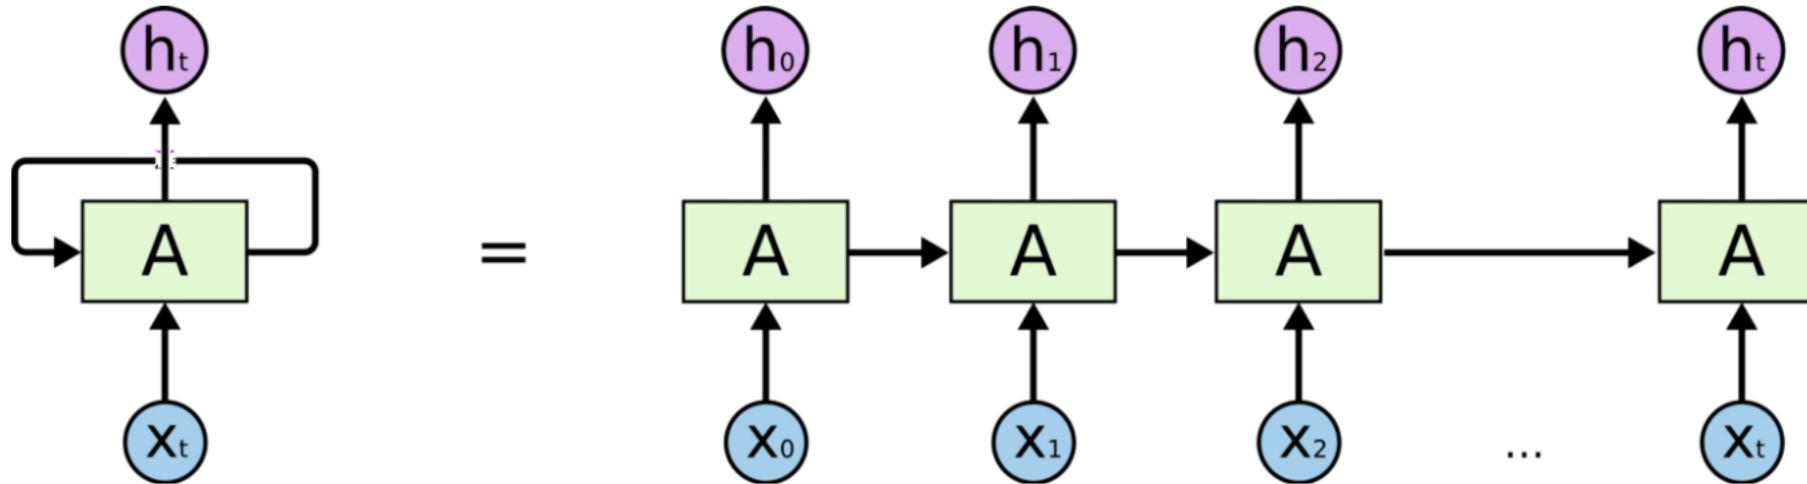
```
Non-trainable params: 0
```

```
from tensorflow.keras.utils import plot_model
plot_model(model, 'model.png', show_shapes=True)
```

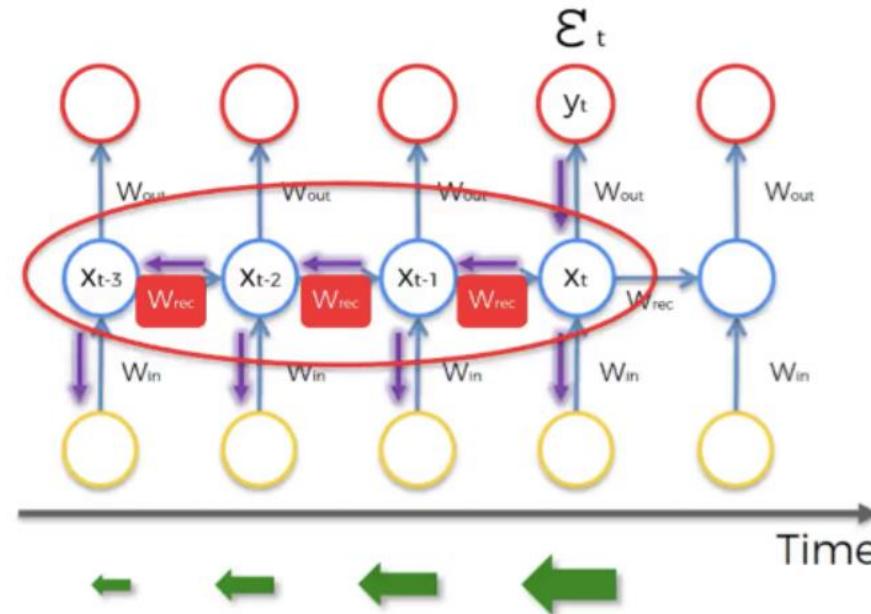


```
# compile the model
model.compile(-----)
# fit the model
model.fit(-----)
# evaluate the model
```

Week 9 – : Deep Learning RNN – Problems - Vanishing gradients



Week 9 – : Deep Learning RNN – Problems - Vanishing gradients



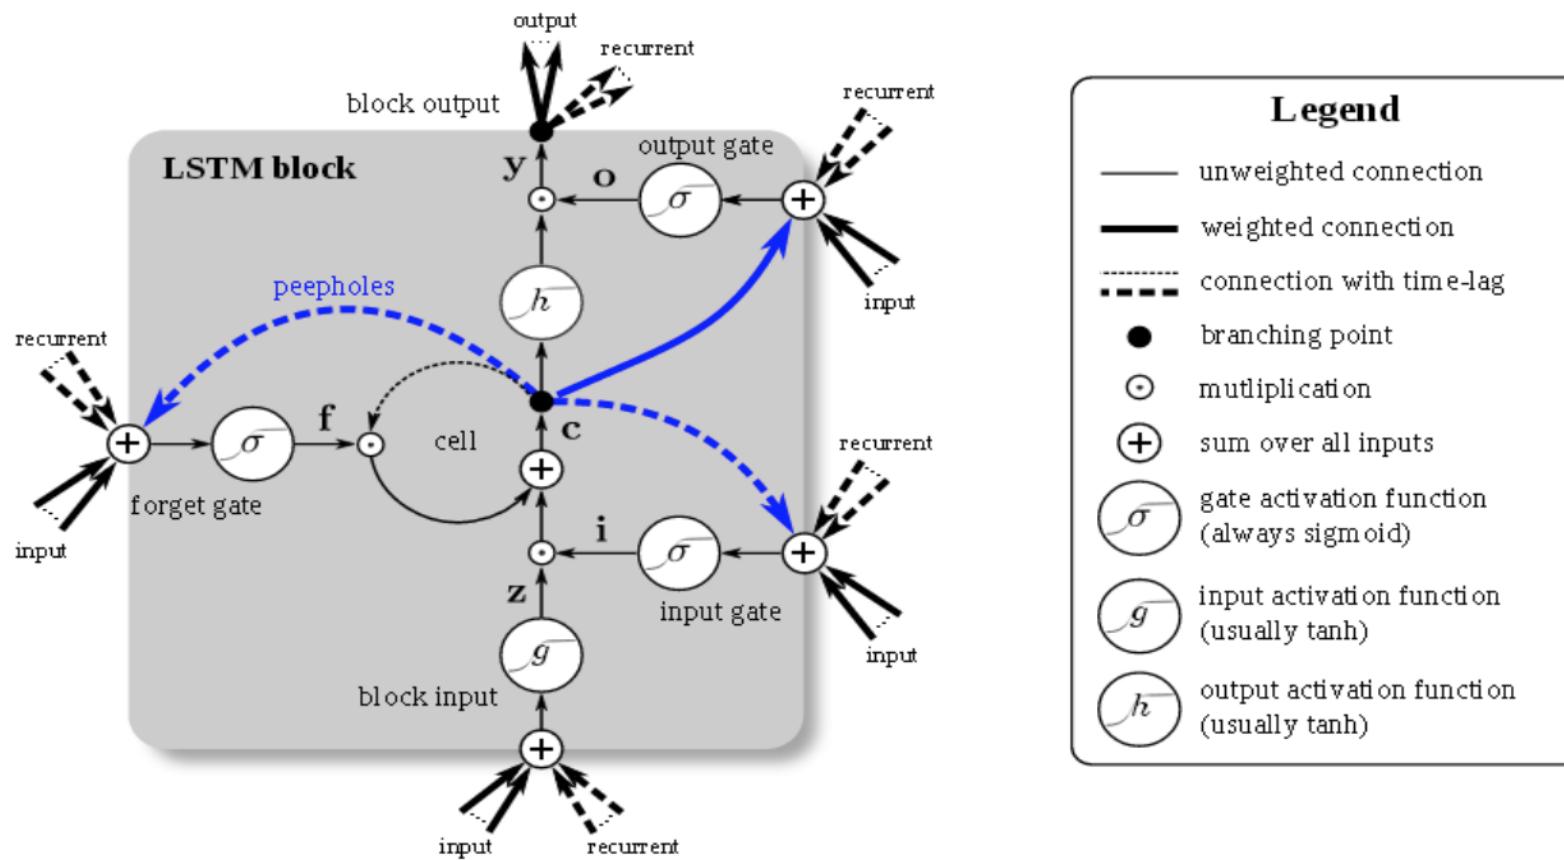
$$\frac{\partial \mathcal{E}}{\partial \theta} = \sum_{1 \leq t \leq T} \frac{\partial \mathcal{E}_t}{\partial \theta} \quad (3)$$

$$\frac{\partial \mathcal{E}_t}{\partial \theta} = \sum_{1 \leq k \leq t} \left(\frac{\partial \mathcal{E}_t}{\partial \mathbf{x}_t} \frac{\partial \mathbf{x}_t}{\partial \mathbf{x}_k} \frac{\partial^+ \mathbf{x}_k}{\partial \theta} \right) \quad (4)$$

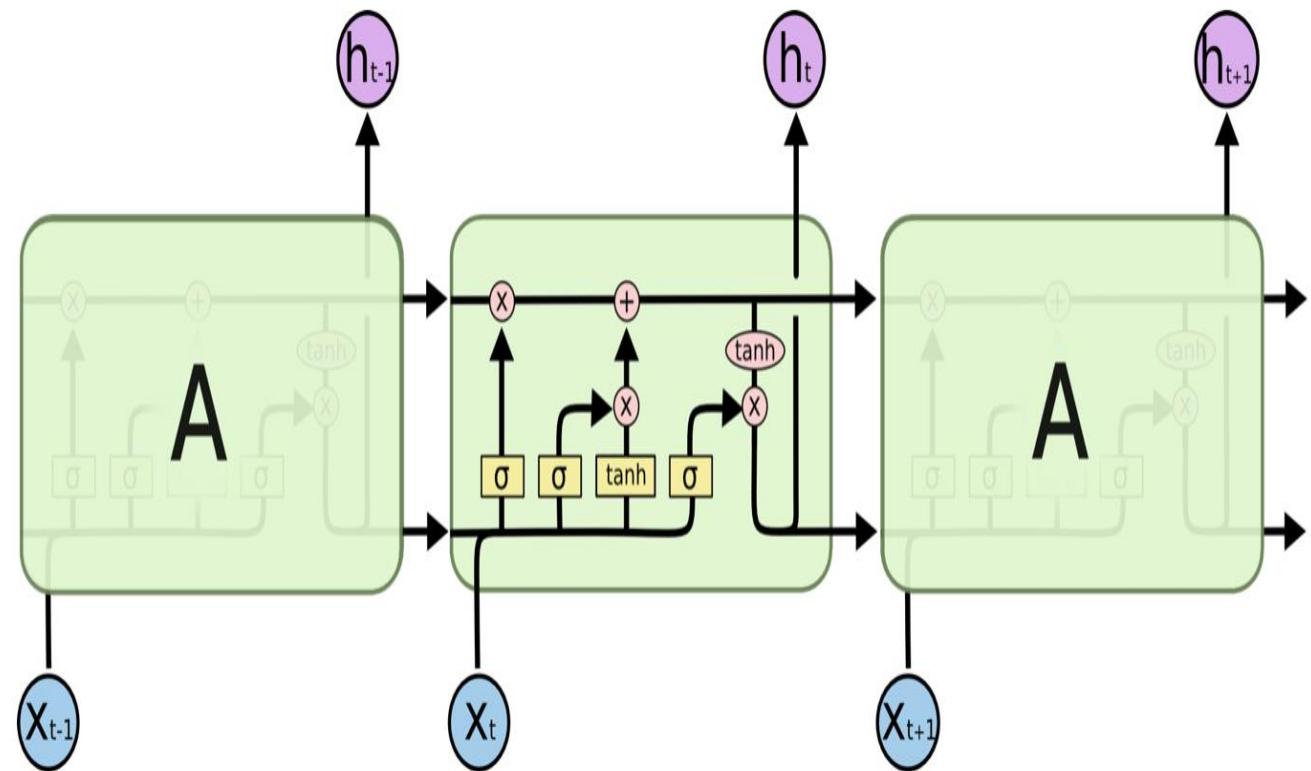
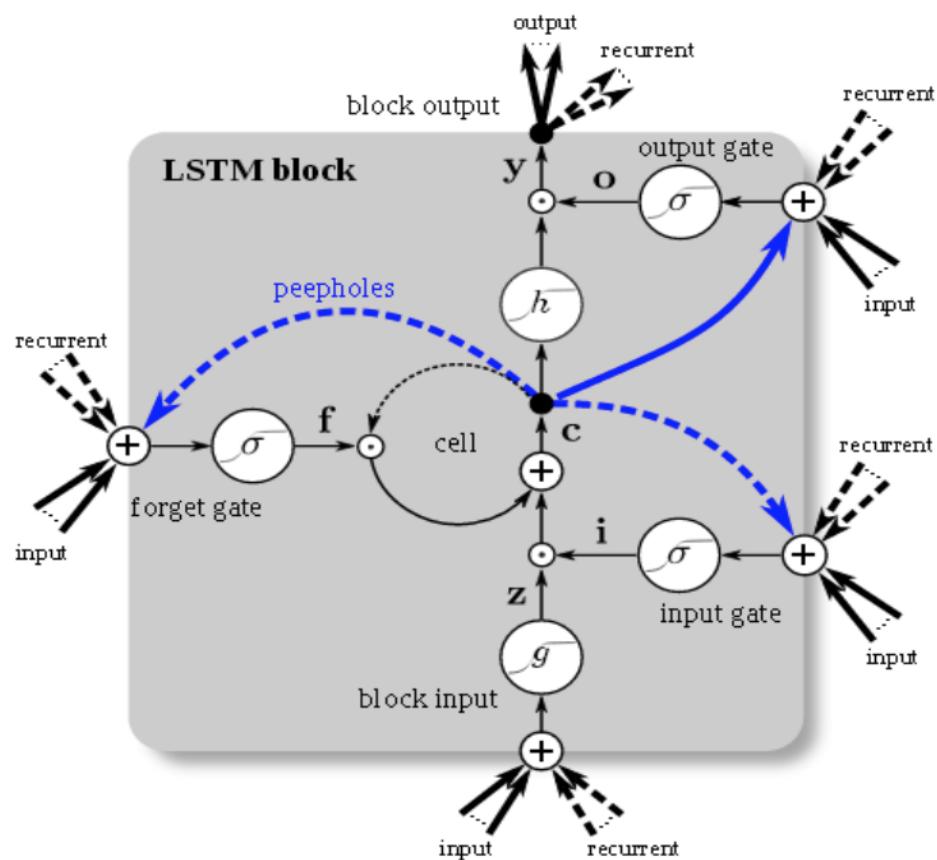
$$\frac{\partial \mathbf{x}_t}{\partial \mathbf{x}_k} = \prod_{t \geq i > k} \frac{\partial \mathbf{x}_i}{\partial \mathbf{x}_{i-1}} = \prod_{t \geq i > k} \mathbf{W}_{rec}^T diag(\sigma'(\mathbf{x}_{i-1})) \quad (5)$$

$\mathbf{W}_{rec} \sim \text{small}$ Vanishing
 $\mathbf{W}_{rec} \sim \text{large}$ Exploding

Week 9 – : Deep Learning Long Short-Term Memory (LSTM)



Week 9 – : Deep Learning Long Short-Term Memory (LSTM)



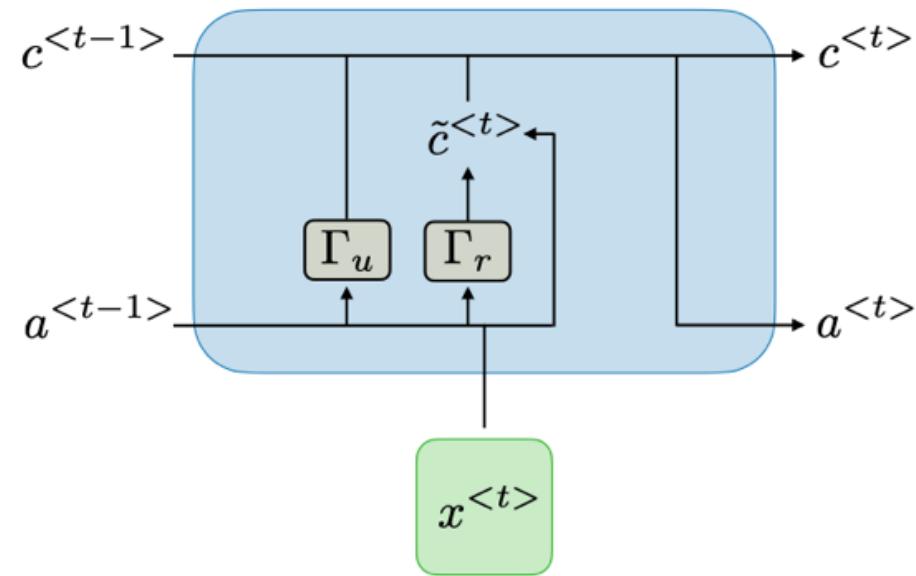
Week 9 – : Deep Learning LSTM -Coding



```
# define model
model = keras.Sequential()
model.add(SimpleRNN(100, input_shape=(5,1)))
model.add(Dense(1))
```

```
# define model
model = keras.Sequential()
model.add(LSTM(100, input_shape=(5,1)))
model.add(Dense(1))
```

Week 9 – : Deep Learning Gated Recurrent Unit (GRU)

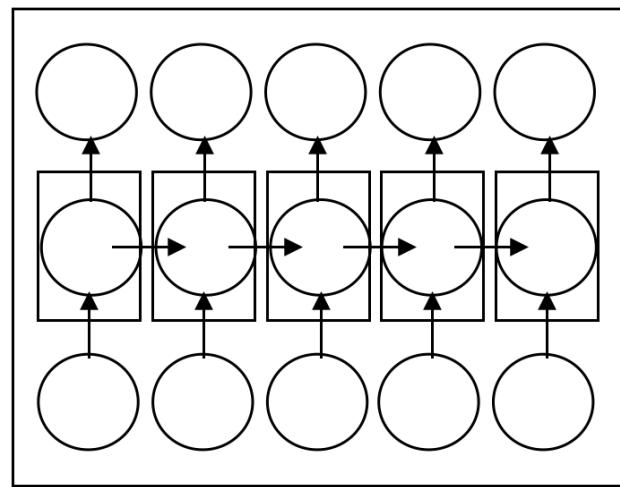


Week 9 – : Deep Learning Gated Recurrent Unit (GRU)

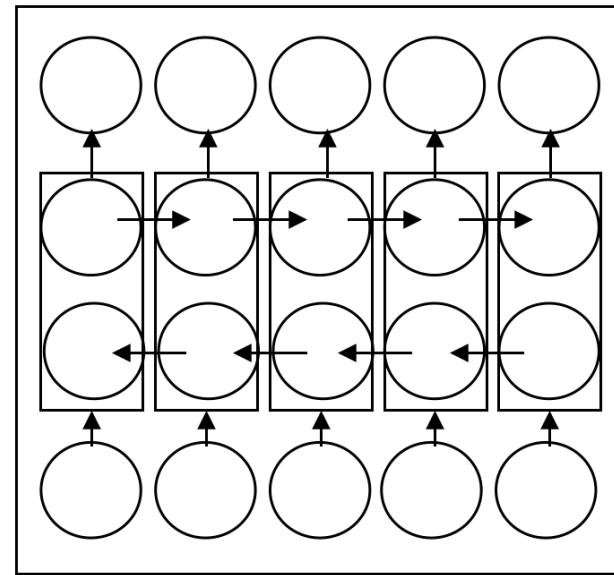


```
# define model
model = keras.Sequential()
model.add(GRU(100, input_shape=(5,1)))
model.add(Dense(1))
```

Week 9 – : Deep Learning Bidirectional



(a)



(b)

Structure overview

- (a) unidirectional RNN
- (b) bidirectional RNN

Week 9 – : Deep Learning Bidirectional LSTM - Code



```
model = Sequential()
```

```
model.add(Bidirectional(LSTM(10, return_sequences=True)
, backward_layer=LSTM(10, activation='relu', return_sequences=True, go_backwards=True)
, input_shape=(5, 10)))
```

```
model.add(Dense(5))
```

Week 9 – : Deep Learning Bidirectional LSTM - Code

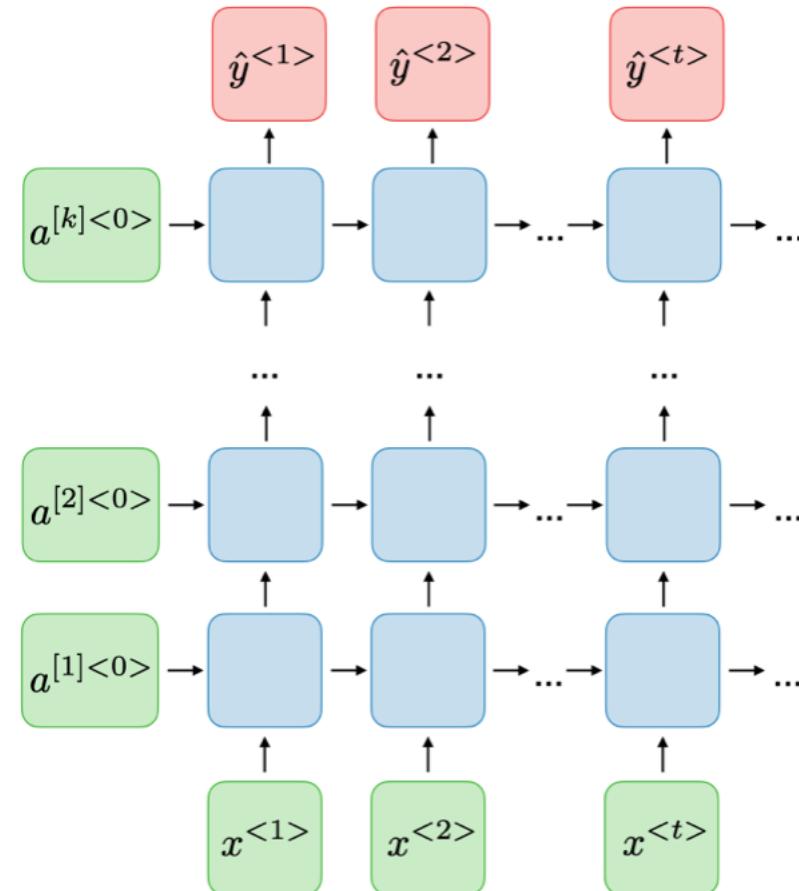


```
model = Sequential()
model.add(Bidirectional(LSTM(10, return_sequences=True), input_shape=(5, 10)))
model.add(Bidirectional(LSTM(10)))
model.add(Dense(5))
model.add(Activation('softmax'))
model.compile(loss='categorical_crossentropy', optimizer='rmsprop')

# With custom backward layer
model = Sequential()
forward_layer = LSTM(10, return_sequences=True)
backward_layer = LSTM(10, activation='relu', return_sequences=True,
                      go_backwards=True)
model.add(Bidirectional(forward_layer, backward_layer=backward_layer,
                       input_shape=(5, 10)))
model.add(Dense(5))
model.add(Activation('softmax'))
model.compile(loss='categorical_crossentropy', optimizer='rmsprop')
```

Week 9 – : Deep Learning

Deep RNN



Week 9 – : Deep Learning

Deep GRU (DGRU)



```
# define model
model = keras.Sequential()
model.add(GRU(100, input_shape=(5,1), return_sequences = True))
model.add(GRU(100, return_sequences = True)),
model.add(GRU(100)),
model.add(Dense(1))
```

Week 9 – : Deep Learning

Applications of RNNs:



RNN models are mostly used in the fields of natural language processing and speech recognition.

Type of RNN	Illustration	Example
One-to-one $T_x = T_y = 1$	A diagram illustrating a one-to-one Recurrent Neural Network (RNN). It shows three rectangular boxes connected vertically by arrows. The bottom box is green and labeled 'x'. An arrow points from 'x' up to a blue box labeled 'a<0>'. Another arrow points from 'a<0>' up to a red box labeled 'ŷ'.	Traditional neural network

Week 9 – : Deep Learning

Applications of RNNs:



RNN models are mostly used in the fields of natural language processing and speech recognition.

Type of RNN	Illustration	Example
One-to-many $T_x = 1, T_y > 1$	A diagram illustrating a One-to-many Recurrent Neural Network (RNN). It shows an input sequence x (green box) being processed by a hidden state $a^{<0>} \rightarrow$ (blue box). This leads to the first output $\hat{y}^{<1>} \rightarrow$ (red box). Subsequent outputs $\hat{y}^{<2>} \rightarrow \dots \rightarrow \hat{y}^{<T_y>}$ (red boxes) are generated sequentially, with each output $\hat{y}^{<t>}$ being produced by the hidden state $a^{<t>} \rightarrow$ (blue box) at time step t . The hidden states are connected by recurrent arrows, showing the temporal dependency between them.	Music generation

Week 9 – : Deep Learning

Applications of RNNs:



RNN models are mostly used in the fields of natural language processing and speech recognition.

Type of RNN	Illustration	Example
Many-to-one $T_x > 1, T_y = 1$	A diagram illustrating a Many-to-one Recurrent Neural Network (RNN). It shows an input sequence $x^{<1>} \rightarrow x^{<2>} \rightarrow \dots \rightarrow x^{<T_x>}$ entering from below, with arrows pointing upwards to hidden states $a^{<0>} \rightarrow a^{<1>} \rightarrow \dots \rightarrow a^{<T_x>}$. The final hidden state $a^{<T_x>}$ has an arrow pointing up to the output layer, which contains a single red box labeled \hat{y} .	Sentiment classification

Week 9 – : Deep Learning

Applications of RNNs:



RNN models are mostly used in the fields of natural language processing and speech recognition.

Type of RNN	Illustration	Example
Many-to-many $T_x = T_y$	A diagram illustrating a Many-to-many Recurrent Neural Network (RNN). It shows two parallel sequences of hidden states and outputs. The input sequence $x^{<1>} \rightarrow x^{<2>} \rightarrow \dots \rightarrow x^{<T_x>}$ feeds into a series of blue rectangular hidden state boxes. The output sequence $a^{<0>} \rightarrow \dots \rightarrow a^{<T_y>}$ is generated from these hidden states, with each a box having an upward arrow pointing to its corresponding \hat{y} box. The output boxes are labeled $\hat{y}^{<1>} \rightarrow \hat{y}^{<2>} \rightarrow \dots \rightarrow \hat{y}^{<T_y>}$.	Name entity recognition

Week 9 – : Deep Learning

Applications of RNNs:



RNN models are mostly used in the fields of natural language processing and speech recognition.

Type of RNN	Illustration	Example
Many-to-many $T_x \neq T_y$	A diagram illustrating a Many-to-many Recurrent Neural Network (RNN). It shows two sequences of hidden states. The input sequence $x^{<1>} \rightarrow \dots \rightarrow x^{<T_x>}$ feeds into a sequence of hidden states $a^{<0>} \rightarrow \dots \rightarrow a^{<T_y>}$. The final hidden state $a^{<T_y>}$ is passed through a layer of output nodes to produce the output sequence $\hat{y}^{<1>} \rightarrow \dots \rightarrow \hat{y}^{<T_y>}$.	Machine translation

Week 9 – : Deep Learning Applications of RNNs: Example



```
from numpy import sqrt
from numpy import asarray
from tensorflow.keras.layers import LSTM

# split a univariate sequence into samples
def split_sequence(sequence, n_steps):
    X, y = list(), list()
    for i in range(len(sequence)):
        # find the end of this pattern
        end_ix = i + n_steps
        # check if we are beyond the sequence
        if end_ix > len(sequence)-1:
            break
        # gather input and output parts of the pattern
        seq_x, seq_y = sequence[i:end_ix], sequence[end_ix]
        X.append(seq_x)
        y.append(seq_y)
    return asarray(X), asarray(y)
```

```
# load the dataset
path = 'https://raw.githubusercontent.com/jbrownlee/Datasets/master/monthly-car-sales.csv'
df = read_csv(path, header=0, index_col=0, squeeze=True)
df.tail()
```

Month	
1968-08	16722
1968-09	14385
1968-10	21342
1968-11	17180
1968-12	14577

Week 9 – : Deep Learning Applications of RNNs: Example



```
# retrieve the values
values = df.values.astype('float32')
# specify the window size
n_steps = 5
# split into samples
X, y = split_sequence(values, n_steps)
# reshape into [samples, timesteps, features]
X = X.reshape((X.shape[0], X.shape[1], 1))
# split into train/test
n_test = 12
X_train, X_test, y_train, y_test = X[:-n_test], X[-n_test:], y[:-n_test], y[-n_test:]
print(X_train.shape, X_test.shape, y_train.shape, y_test.shape)
```

(91, 5, 1) (12, 5, 1) (91,) (12,)

Week 9 – : Deep Learning Applications of RNNs: Example



```
# define model
model = Sequential()
model.add(LSTM(100, activation='relu', kernel_initializer='he_normal', input_shape=(n_steps,1)))
model.add(Dense(50, activation='relu', kernel_initializer='he_normal'))
model.add(Dense(50, activation='relu', kernel_initializer='he_normal'))
model.add(Dense(1))
# compile the model
model.compile(optimizer='adam', loss='mse', metrics=['mae'])
# fit the model
model.fit(X_train, y_train, epochs=350, batch_size=32, verbose=0, validation_data=(X_test, y_test))
# evaluate the model
mse, mae = model.evaluate(X_test, y_test, verbose=0)
print('MSE: %.3f, RMSE: %.3f, MAE: %.3f' % (mse, sqrt(mse), mae))
# make a prediction
row = asarray([18024.0, 16722.0, 14385.0, 21342.0, 17180.0]).reshape((1, n_steps, 1))
yhat = model.predict(row)
print('Predicted: %.3f' % (yhat))
```

Week 9 – : Deep Learning Applications of RNNs: Example



```
model.summary()
```

Model: "sequential_7"

Layer (type)	Output Shape	Param #
=====		
lstm_3 (LSTM)	(None, 100)	40800
dense_19 (Dense)	(None, 50)	5050
dense_20 (Dense)	(None, 50)	2550
dense_21 (Dense)	(None, 1)	51
=====		

Total params: 48,451

Trainable params: 48,451

Non-trainable params: 0

Week 9 : Deep Learning with NLP



weather	weather	play
Sunny	Hot	No
Overcast	Hot	Yes
Sunny	Mild	No

```
[ ] 1 # Assigning features and label variables
2 weather=['Sunny','Sunny','Overcast','Rainy','Rainy','Rainy','Overcast','Sunny','Sunny',
3 'Rainy','Sunny','Overcast','Overcast','Rainy']
4 temp=['Hot','Hot','Hot','Mild','Cool','Cool','Mild','Cool','Mild','Mild','Mild','Hot','Mild']
5 play=['No','No','Yes','Yes','Yes','No','Yes','No','Yes','Yes','Yes','Yes','Yes','Yes','No']
```

Week 9 : Deep Learning with NLP - One-hot encodings



```
: # Converting strings into numbers.
le = preprocessing.LabelEncoder()
weatherEncode = le.fit(weather)
weather_encoded = weatherEncode.transform(weather)
print('weather',weather_encoded)

# Converting string temp into numbers
leT = preprocessing.LabelEncoder()
WorkDayEncode = leT.fit(WorkDay)
WorkDay_encoded = WorkDayEncode.transform(WorkDay)
print ('WorkDay',WorkDay_encoded)

leTT = preprocessing.LabelEncoder()
TimeDayEncode = leTT.fit(WorkDay)
TimeDay_encoded = TimeDayEncode.transform(WorkDay)
print ('WorkDay',TimeDay_encoded)

# Converting string labels into numbers
leL = preprocessing.LabelEncoder()
Lebal_encoded=leL.fit(TrafficJam)
label=Lebal_encoded.transform(TrafficJam)
print ("TrafficJam:",label)

weather [0 0 0 0 0 0 1 1 1 1 1 1 2 2 2 2 2 2]
WorkDay [1 1 1 0 0 0 1 1 1 0 0 0 1 1 1 0 0 0]
WorkDay [1 1 1 0 0 0 1 1 1 0 0 0 1 1 1 0 0 0]
TrafficJam: [1 0 1 0 0 0 1 1 1 0 0 0 1 1 1 1 0 1]
```

Week 9 : Deep Learning with NLP - One-hot encodings



The cat sat on the mat

One-hot encoding

	cat	mat	on	sat	the
the =>	0	0	0	0	1
cat =>	1	0	0	0	0
sat =>	0	0	0	1	0

Week 9 : Deep Learning with NLP - One-hot encodings



Discrete Representation

motel[000000000010000]
AND
hotel [000000010000000] = 0

Week 9 : Deep Learning with NLP - Word embeddings



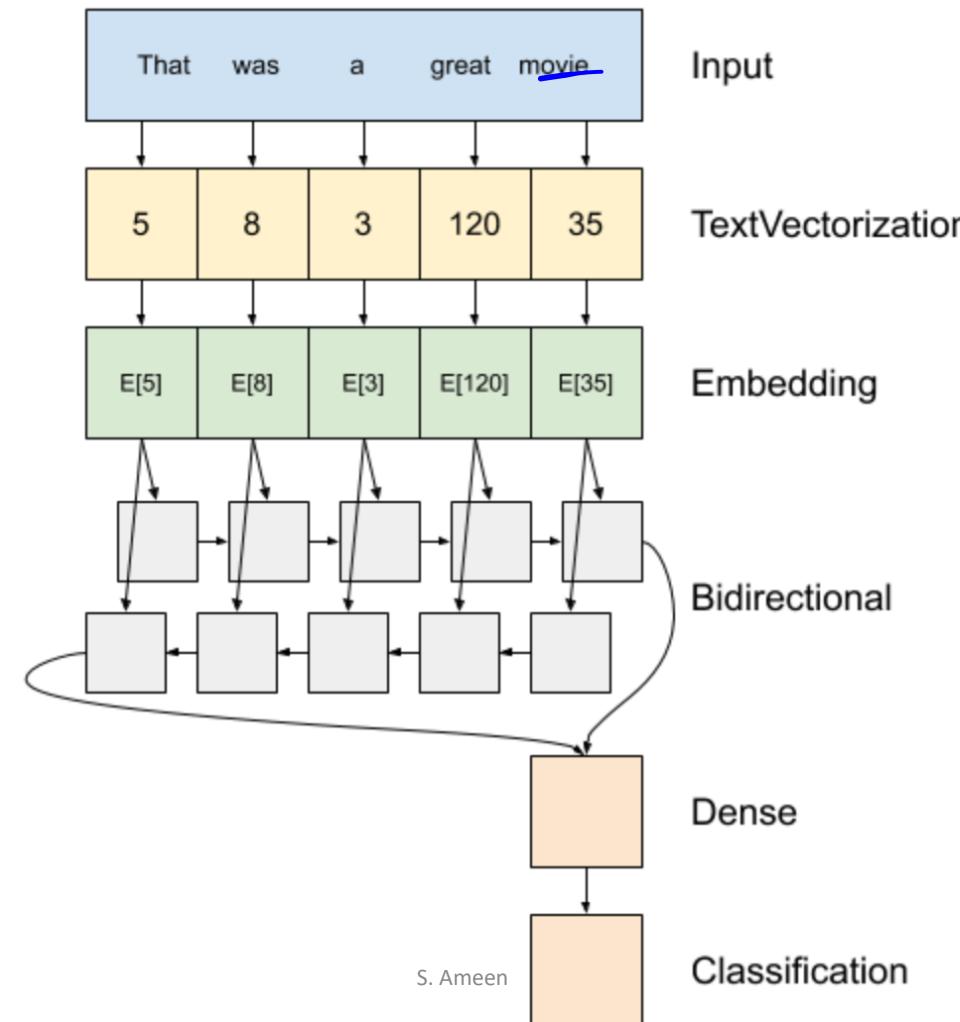
Word embeddings

Word embeddings give us a way to use an efficient, dense representation in which similar words have a similar encoding.

A 4-dimensional embedding

cat =>	1.2	-0.1	4.3	3.2
mat =>	0.4	2.5	-0.9	0.5
on =>	2.1	0.3	0.1	0.4

Week 9 : Deep Learning with NLP - Word embeddings



Week 9 : Deep Learning with NLP - Word embeddings



```
import csv
import tensorflow as tf
import numpy as np
from tensorflow.keras.preprocessing.text import Tokenizer
from tensorflow.keras.preprocessing.sequence import pad_sequences

vocab_size = 1000
embedding_dim = 16
max_length = 120
trunc_type='post'
padding_type='post'
oov_tok = "<OOV>"
training_portion = .8
```

A 4-dimensional embedding

cat =>	1.2	-0.1	4.3	3.2
mat =>	0.4	2.5	-0.9	0.5
on =>	2.1	0.3	0.1	0.4

```
>>> sequence = [[1], [2, 3], [4, 5, 6]]
>>> tf.keras.preprocessing.sequence.pad_sequences(sequence)
array([[0, 0, 1],
       [0, 2, 3],
       [4, 5, 6]], dtype=int32)
```

```
>>> tf.keras.preprocessing.sequence.pad_sequences(sequence, value=-1)
array([[-1, -1, 1],
       [-1, 2, 3],
       [4, 5, 6]], dtype=int32)
```

```
>>> tf.keras.preprocessing.sequence.pad_sequences(sequence, padding='post')
array([[1, 0, 0],
       [2, 3, 0],
       [4, 5, 6]], dtype=int32)
```

```
>>> tf.keras.preprocessing.sequence.pad_sequences(sequence, maxlen=2)
array([[0, 1],
       [2, 3],
       [5, 6]], dtype=int32)
```

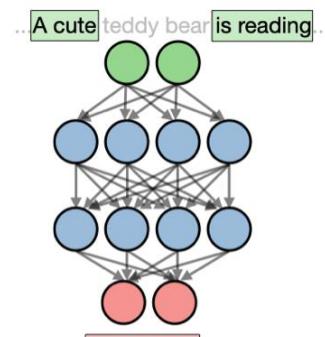
Week 9 : Deep Learning with NLP - word2Vec



Word2Vec: There are two models that are commonly used to train these embeddings:

The Skip-gram model takes the input as each word in the corpus, sends them to a hidden layer (embedding layer) and from there it predicts the context words. Once trained, the embedding for a particular word is obtained by feeding the word as input and taking the hidden layer value as the final embedding vector.

The CBOW (Continuous Bag of Words) model takes the input the context words for the given word, sends them to a hidden layer (embedding layer) and from there it predicts the original word. Once again, after training, the embedding for a particular word is obtained by feeding the word as input and taking the hidden layer value as the final embedding vector.



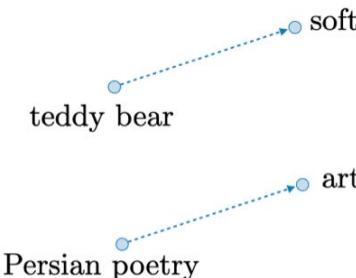
Train network on proxy task



Extract high-level representation



Compute word embeddings



Extract high-level representation

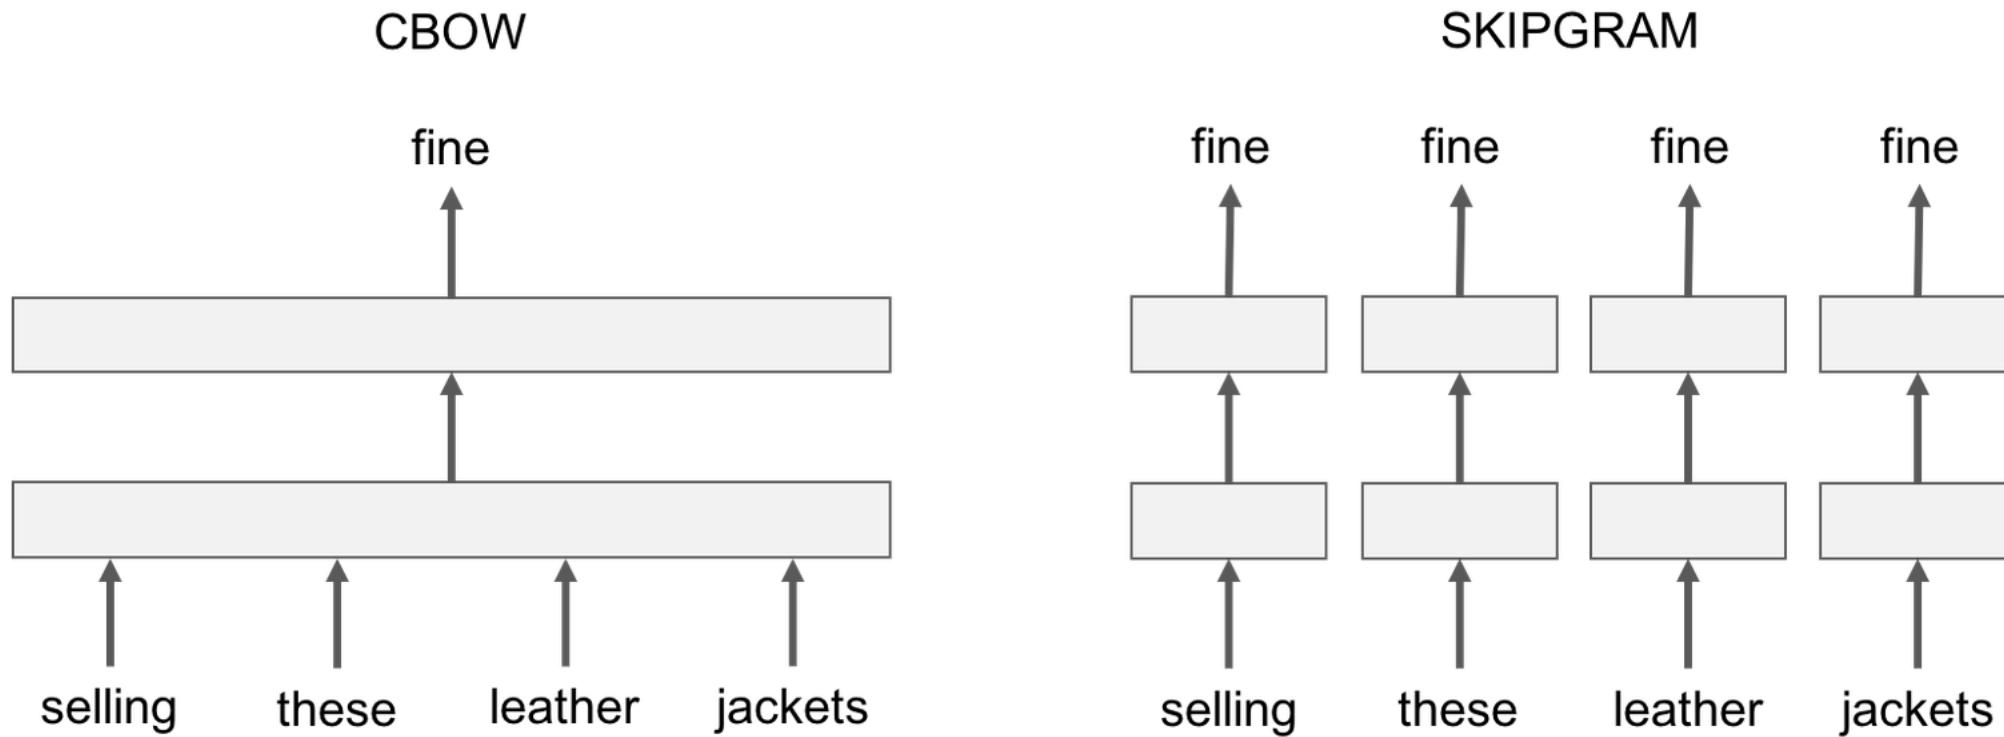


Compute word embeddings



Tomáš Mikolov

Week 9 : Deep Learning with NLP – CBOW vs Skipgram



Week 9 : Deep Learning with NLP - GloVe: Global Vectors for Word Representation



GloVe is an unsupervised learning algorithm for obtaining vector representations for words. Training is performed on aggregated global word-word co-occurrence statistics from a corpus, and the resulting representations showcase interesting linear substructures of the word vector space.

$$J(\theta) = \frac{1}{2} \sum_{i,j=1}^W f(P_{ij})(u_i^T v_j - \log P_{ij})^2$$



Week 9 : Deep Learning with NLP - GloVe:



Nearest words to
frog:

1. frogs
2. toad
3. litoria
4. leptodactylidae
5. rana
6. lizard
7. eleutherodactylus



litoria



rana



leptodactylidae

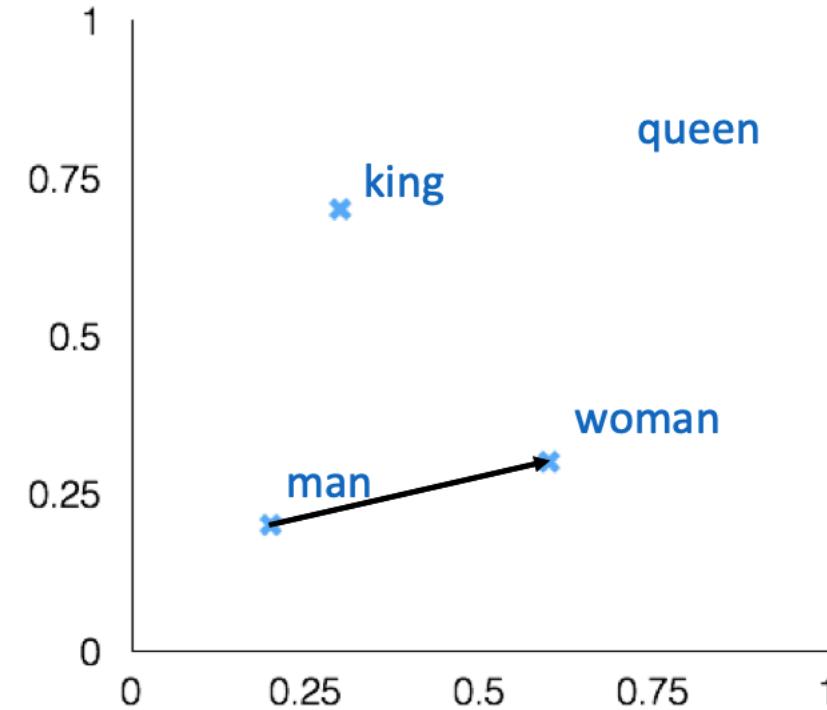


eleutherodactylus

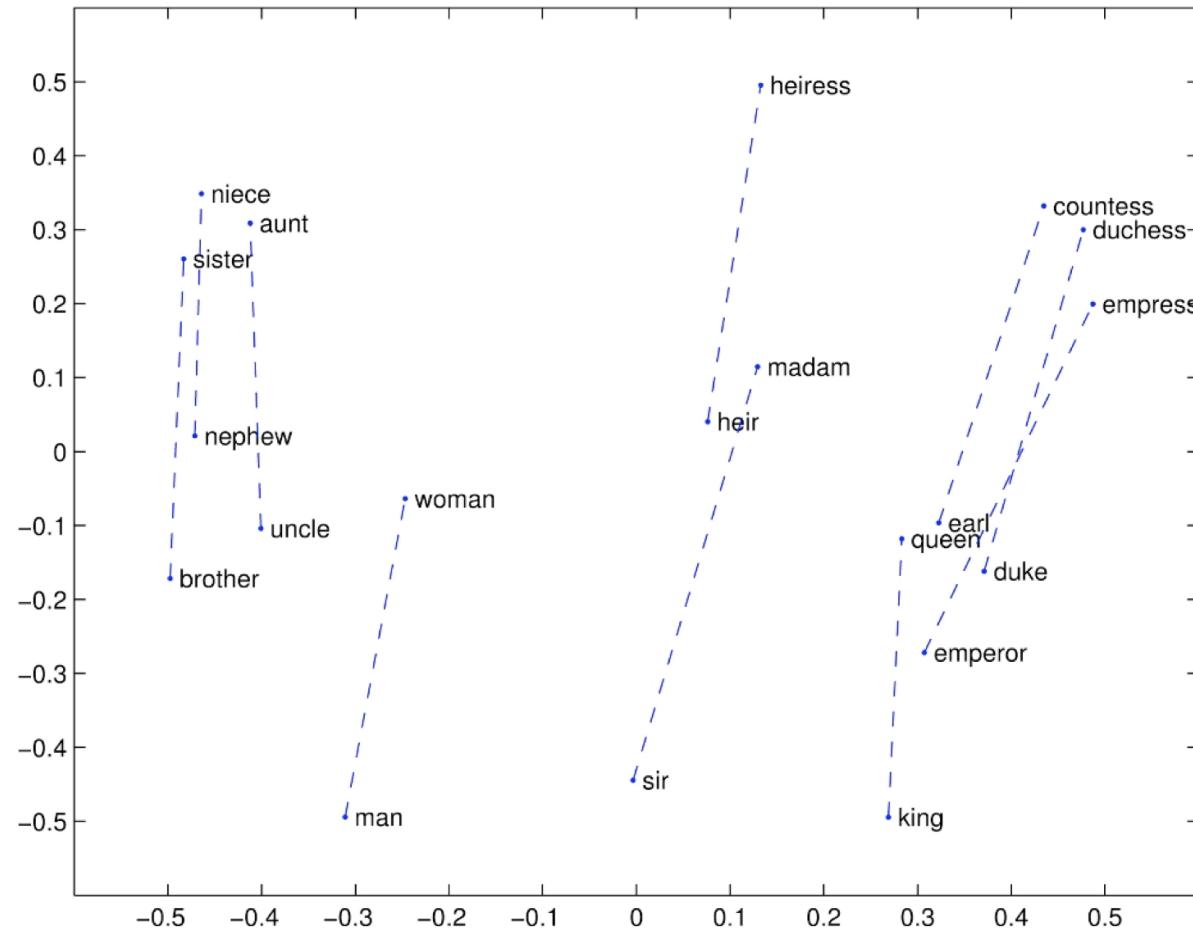
Week 9 : Deep Learning with NLP - GloVe:



man:woman :: king:?		
+	king	[0.30 0.70]
-	man	[0.20 0.20]
+	woman	[0.60 0.30]
<hr/>		queen [0.70 0.80]



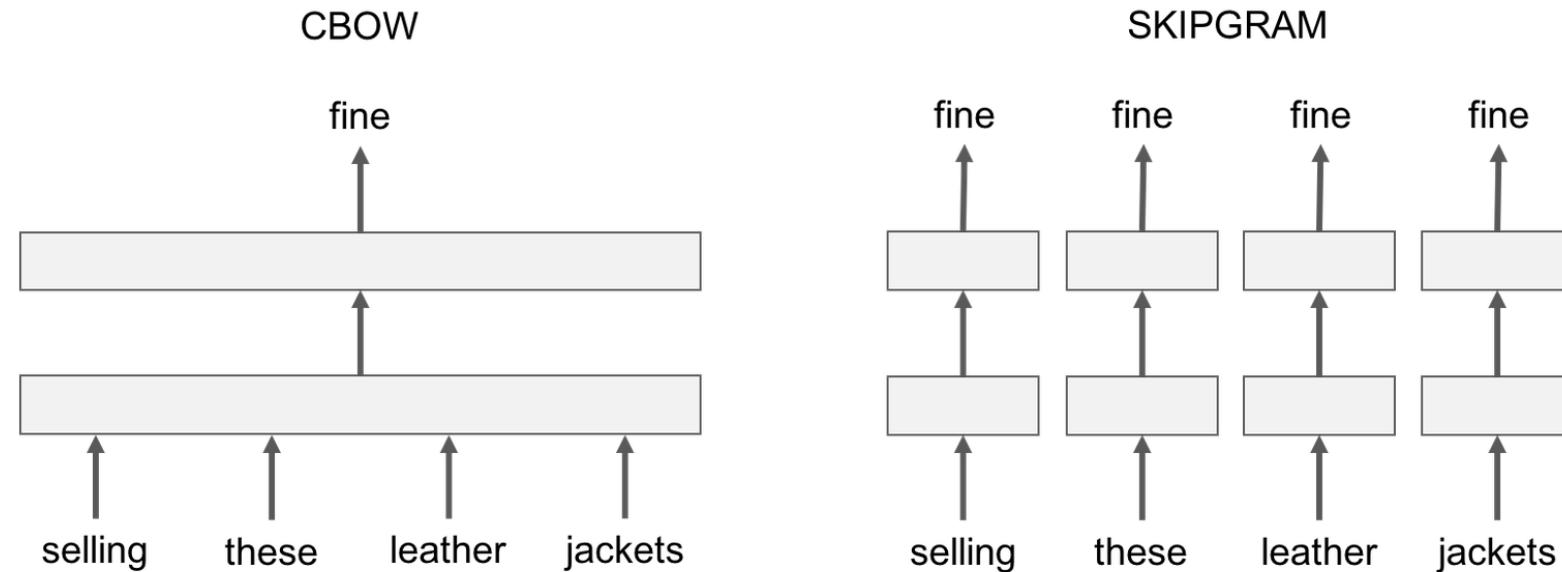
Week 9 : Deep Learning with NLP - GloVe:



Week 9 : Deep Learning with NLP - FastText:



FastText: is a library for efficient learning of word representations and sentence classification.



```
import fastText
```

<https://fasttext.cc/>

Week 9 : Deep Learning with NLP :

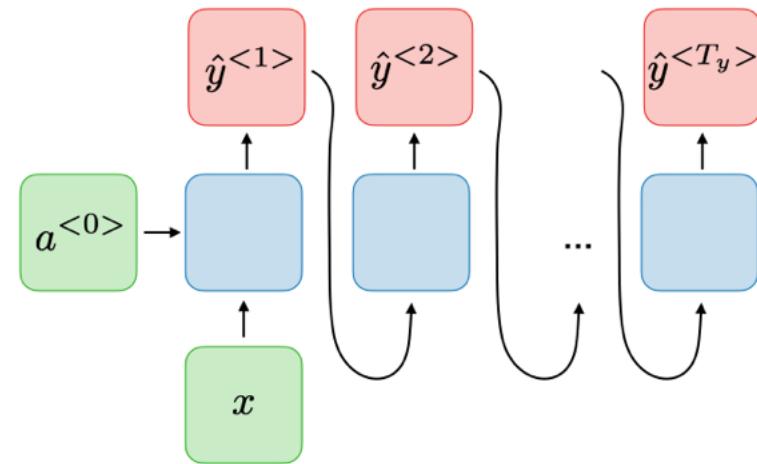


Demo

Week 9 : Deep Learning with NLP : Application :



One-to-many
 $T_x = 1, T_y > 1$



Music generation

Week 9 : Deep Learning with NLP : Application : - Shakespeare



PANDARUS:

Alas, I think he shall be come approached and the day
When little strain would be attain'd into being never fed,
And who is but a chain and subjects of his death,
I should not sleep.

Second Senator:

They are away this miseries, produced upon my soul,
Breaking and strongly should be buried, when I perish
The earth and thoughts of many states.

DUKE VINCENTIO:

Well, your wit is in the care of side and that.

Second Lord:

They would be ruled after this chamber, and
my fair nues begun out of the fact, to be conveyed,
Whose noble souls I'll have the heart of the wars.

Clown:

Come, sir, I will make did behold your worship.

VIOLA:

I'll drink it.

Week 9 : Deep Learning with NLP : Application :- Wikipedia



Naturalism and decision for the majority of Arab countries' capitalide was grounded by the Irish language by [[John Clair]], [[An Imperial Japanese Revolt]], associated with Guangzham's sovereignty. His generals were the powerful ruler of the Portugal in the [[Protestant Immineners]], which could be said to be directly in Cantonese Communication, which followed a ceremony and set inspired prison, training. The emperor travelled back to [[Antioch, Perth, October 25|21]] to note, the Kingdom of Costa Rica, unsuccessful fashioned the [[Thrales]], [[Cynth's Dajoard]], known in western [[Scotland]], near Italy to the conquest of India with the conflict. Copyright was the succession of independence in the slop of Syrian influence that was a famous German movement based on a more popular servicious, non-doctrinal and sexual power post. Many governments recognize the military housing of the [[Civil Liberalization and Infantry Resolution 265 National Party in Hungary]], that is sympathetic to be to the [[Punjab Resolution]] (PJS)[<http://www.humah.yahoo.com/guardian.cfm/7754800786d17551963s89.htm> Official economics Adjoint for the Nazism, Montgomery was swear to advance to the resources for those Socialism's rule, was starting to signing a major tripad of aid exile.]]

Week 9 : Deep Learning with NLP : Application : - Algebraic Geometry (Latex)



Proof. Omitted. \square

Lemma 0.1. Let \mathcal{C} be a set of the construction.

Let \mathcal{C} be a gerber covering. Let \mathcal{F} be a quasi-coherent sheaves of \mathcal{O} -modules. We have to show that

$$\mathcal{O}_{\mathcal{O}_X} = \mathcal{O}_X(\mathcal{L})$$

Proof. This is an algebraic space with the composition of sheaves \mathcal{F} on $X_{\text{étale}}$ we have

$$\mathcal{O}_X(\mathcal{F}) = \{\text{morph}_1 \times_{\mathcal{O}_X} (\mathcal{G}, \mathcal{F})\}$$

where \mathcal{G} defines an isomorphism $\mathcal{F} \rightarrow \mathcal{G}$ of \mathcal{O} -modules. \square

Lemma 0.2. This is an integer Z is injective.

Proof. See Spaces, Lemma ??.

Lemma 0.3. Let S be a scheme. Let X be a scheme and X is an affine open covering. Let $\mathcal{U} \subset \mathcal{X}$ be a canonical and locally of finite type. Let X be a scheme. Let X be a scheme which is equal to the formal complex.

The following to the construction of the lemma follows.

Let X be a scheme. Let X be a scheme covering. Let

$$b : X \rightarrow Y' \rightarrow Y \rightarrow Y' \times_X Y \rightarrow X.$$

be a morphism of algebraic spaces over S and Y .

Proof. Let X be a nonzero scheme of X . Let X be an algebraic space. Let \mathcal{F} be a quasi-coherent sheaf of \mathcal{O}_X -modules. The following are equivalent

- (1) \mathcal{F} is an algebraic space over S .
- (2) If X is an affine open covering.

Consider a common structure on X and X the functor $\mathcal{O}_X(U)$ which is locally of finite type. \square

This since $\mathcal{F} \in \mathcal{F}$ and $x \in \mathcal{G}$ the diagram

$$\begin{array}{ccccc}
 S & \longrightarrow & & & \\
 \downarrow & & & & \\
 \xi & \longrightarrow & \mathcal{O}_{X'} & \searrow & \\
 \text{gor}_s & & \uparrow & & \\
 & & =\alpha' & \longrightarrow & \\
 & & \downarrow & & \\
 \text{Spec}(K_\psi) & & & & \text{Mor}_{\text{Sets}} \quad d(\mathcal{O}_{X/k}, \mathcal{G}) \\
 & & & & \downarrow \\
 & & & & X
 \end{array}$$

is a limit. Then \mathcal{G} is a finite type and assume S is a flat and \mathcal{F} and \mathcal{G} is a finite type f_* . This is of finite type diagrams, and

- the composition of \mathcal{G} is a regular sequence,
- $\mathcal{O}_{X'}$ is a sheaf of rings.

Proof. We have see that $X = \text{Spec}(R)$ and \mathcal{F} is a finite type representable by algebraic space. The property \mathcal{F} is a finite morphism of algebraic stacks. Then the cohomology of X is an open neighbourhood of U . \square

Proof. This is clear that \mathcal{G} is a finite presentation, see Lemmas ??.

A reduced above we conclude that U is an open covering of C . The functor \mathcal{F} is a “field”

$$\mathcal{O}_{X,x} \longrightarrow \mathcal{F}_{\overline{x}} \dashv^{-1} (\mathcal{O}_{X_{\text{étale}}}) \longrightarrow \mathcal{O}_{X_k}^{-1} \mathcal{O}_{X_\lambda} (\mathcal{O}_{X_k}^{\overline{x}})$$

is an isomorphism of covering of \mathcal{O}_{X_k} . If \mathcal{F} is the unique element of \mathcal{F} such that X is an isomorphism.

The property \mathcal{F} is a disjoint union of Proposition ?? and we can filtered set of presentations of a scheme \mathcal{O}_X -algebra with \mathcal{F} are opens of finite type over S . If \mathcal{F} is a scheme theoretic image points. \square

If \mathcal{F} is a finite direct sum \mathcal{O}_{X_λ} is a closed immersion, see Lemma ?? . This is a sequence of \mathcal{F} is a similar morphism.

Week 9 : Deep Learning with NLP : Application :- Linux Source Code



```
static int indicate_policy(void)
{
    int error;
    if (fd == MARN_EPT) {
        /*
         * The kernel blank will coeld it to userspace.
         */
        if (ss->segment < mem_total)
            unblock_graph_and_set_blocked();
        else
            ret = 1;
        goto bail;
    }
    segaddr = in_SB(in.addr);
    selector = seg / 16;
    setup_works = true;
    for (i = 0; i < blocks; i++) {
        seq = buf[i++];
        bpf = bd->bd.next + i * search;
        if (fd) {
            current = blocked;
        }
    }
    rw->name = "Getjbbregs";
    bprm_self_clearl(&iv->version);
    regs->new = blocks[(BPF_STATS << info->historidac)] | PFMR_CLOBATHINC_SECONDS << 12;
    return segtable;
}
```

Week 9 : Deep Learning with NLP : Application : - Generating Baby Names



*Rudi Levette Berice Lussa Hany Mareanne Chrestina Carissy Marylen Hammie Janye Marlise Jacacie
Hendred Romand Charienna Nenotto Ette Dorane Wallen Marly Darine Salina Elvyn Ersia Maralena Minoria Ellia
Charmin Antley Nerille Chelon Walmor Evena Jeryly Stachon Charisa Allisa Anatha Cathanie Geetra Alexie Jerin
Cassen Herbett Cossie Velen Daurenge Robester Shermond Terisa Licia Roselen Ferine Jayn Lusine Charyanne
Sales Sanny Resa Wallon Martine Merus Jelen Candica Wallin Tel Rachene Tarine Ozila Ketia Shanne Arnande
Karella Roselina Alessia Chasty Deland Berther Geamar Jackein Mellisand Sagdy Nenc Lessie Rasemy Guen
Gavi Milea Anneda Margoris Janin Rodelin Zeanna Elyne Janah Ferzina Susta Pey Castina*

Week 9 : Deep Learning with NLP : Application : - Poem



These are some random sampling generated by the network

ما أنا من أبي والشوق يُسْرِجُ بالأمْرِ الأَخِيرِ وَلَقَدْ أَرْدَتُ بِهِ الْعَرَاقَ وَمِنْ أَنْ يَتَادَمَ الْأَرْضُ وَبَيْنَ الْمُحَالِ وَالْمُعَالِي
وَالْأَمَانِي وَالْعُلُومُ وَالْحَيَاةِ فِي الْقُلْبِ أَنْ لَمْ تُسْتَرِّي

ثَمَنَّيْتُ الْحَرَى وَلَمْ أَبْقِ بِهَا مَلَاءِ مِنْ مَاءِ الْمُحَالِ الْمُسْتَرِّيِّ أَنْتَ
مَا أَرْدَاكَ أَنْ يَتَقَرَّبَا وَمَلَاءِ يَخْسِبُ بَعْدَ الْعَدْلِ الْأَنْفُسِ وَالْعَالَمِ الْحُبُّ عَلَى الْمَاءِ الْغُرَابِ وَالْعَيْنُ يَعْلُوْهُ عَلَى الْعَدْلِ
مُسْتَنْهَلُ الْحُبُّ يَسْتَبَقُّ يِهِ الْعَدْلِ فَالْهُوَيْ وَالْعَلَى وَالْحَيْزُ يَعْرِفُ الْأَحْدَادُ وَالْعَيْنُ يُسْمِعُ مِنْ سَمَاءِ الْجُنُبِ مَا أَوْدَعَتِنَّ
حَزِيبِهِ فِي الْمُسْتَبَاحِ وَجَادُهُمَا بِكَ الْمُسْتَسِلُمُ الْمُسْتَسِلُمُ الْمُسْتَسِلُمُ الْمُسْتَسِلُمُ
وَالْمُعَالِي وَالْعَمَرُ أَوْ لِمَاءِ مُحْتَمِلُ وَالرَّمْسُ يَسْمُوْهِ الْفَدَاؤُ وَسِرِّي بِالْقَرَائِنِ وَالْمُعَالِي وَاجْأَرُ الْحَسْبِ مَا
بَاتَ

وَالصَّمَتُ الَّذِي يُرجِي الرَّجَاءَ بِهِ الْجَوَادُ الرَّاقِي وَالشَّمْسُ يَعْرِفُ أَنَّهُ الْأَرْضُ الَّذِي يَحْرُقُ الْأَمْرُ الَّذِي أَخْرَجَ
الْخُلُودَ وَأَعْطَى الْأَرْضَ مِنْ بَعْضِ الْأَرْضِ وَلَمْ تَرِدْ إِلَى الْأَرْضِ مِنْ أَقْبَالِهِ وَتَرَامَى وَلَمْ أَدِرِ مِنْ أَنْجَاهِهِ وَالْعَدَائِمُ وَأَشْرَقَ مِنْ
عَيْنِ وَلَمْ يَعْرِفِ الْعَدَائِمُ حَزَانِ ما يَشْرِي الْكُرُومُ الْمُوَاعِدُولُمْ تَرْجَعُ الْأَرْضُ الْمُنَايَا وَأَعْنَرَتْهُ بِهِ مِنْ لَوَاءِ الْخَيْرِ وَالْمَاءِ
مَائِدٌ وَبَلْلُغُ فِي الْأَرْضِ الْحُبَابِ وَالْحَبْرُ وَأَصْبَحَتْ أَحَبِّي لَيْسَ يَعْرِفُ أَنَّنِي لَعْمَرُكَ أَرْضَى مِنْ بَنِي مَاءِ مَائِدٍ وَإِنْ كُنْتُ مِنْ
أَعْنَاقِهِ وَبَلِيلَةَ تَرَى كَيْفَ يَكِيمُ لَوْ تَرَى الْعَارِ مِنْ عَدَا وَأَعْرَضَ مِنْ حَالٍ وَأَرْضِي بِأَبْرِقَوْلَمْ أَكُ مِنْ مَاءِ وَلَا أَوْ تَرَى
بِهِ وَأَعْطِي بِهِ الْأَرْضَ الْحُبَابِ بِأَبْصَرِ الْعَصْرِ الْإِسْلَامِي <> مروانِ ابْنِ أَبِي حَفْصَةِ <> أَتَانِي كَأَنَّ الْعَدَاءَ الْحَبِيبُ رَقْم
الْقَصِيدَةِ : 27920
يَا رَبَّ أَسْرَى لَيْسَ يَحْسُبُهُ الْأَمْرُ رَقْمَ الْقَصِيدَةِ : 17967
أَصَابَ الْحَيَاةُ وَالْخَلْقُ مُرْتَجِعٌ



Salem Ameen

Sessional Lecturer at The University of Salford

1 article

Inspired by the work of [Andrej Karpathy blog](#) for training multi-layer Recurrent Neural Network (RNN, LSTM, and GRU) and then sampling from the model. I trained **char-rnn** to many of Arabic poems and surprisingly the network can match Arabic diacritics which is difficult to match to many people but the network learns to generate them. I trained the file with size of nearly 93MB. LSTM with 3 layers, the size of the network 700, 0.5 dropout while number of parameters in the model: 10341950

Week 9 : Deep Learning with NLP : Application : - Poem



28

Solo

A musical score for a solo instrument. The staff begins with a sixteenth-note chord followed by a sixteenth-note grace note before a eighth-note. This pattern repeats. The key signature changes to one sharp at the end of the measure. Measure number 28 is indicated above the staff.

29

Solo

A musical score for a solo instrument. The staff begins with a sixteenth-note chord followed by a sixteenth-note grace note before a eighth-note. This pattern repeats. The key signature changes to one sharp at the beginning of the measure. Measure number 29 is indicated above the staff.

Week 9 : AI



- Thank You

